

Lecture 8: Light, color and shadows

DHBW, Computer Graphics

Lovro Bosnar

22.2.2023.

Syllabus

- 3D scene
 - Object
 - Camera
 - Light
 - Rendering
 - Image and display
- Light
 - Light and color
 - Real world light sources
 - Light models
 - Physical lights
 - Non-physical lights
 - Environment illumination
 - Shadows





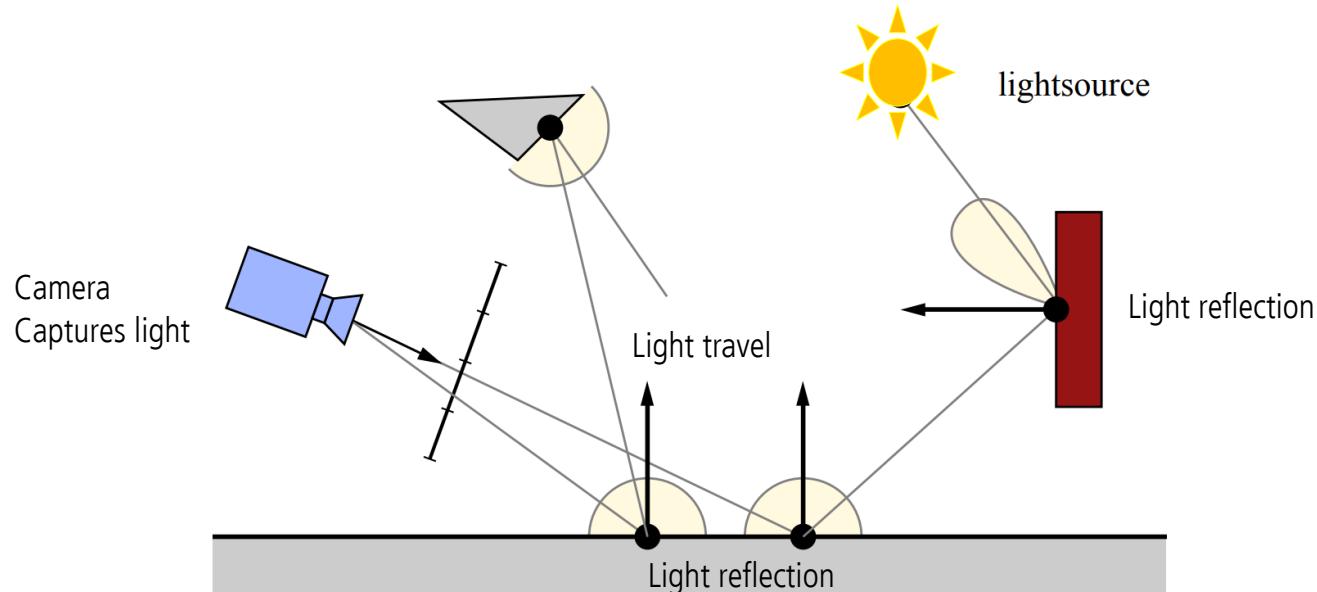
Light is important because without light...

Light sources with size and shape would be visible if we would look towards them.

Light and color

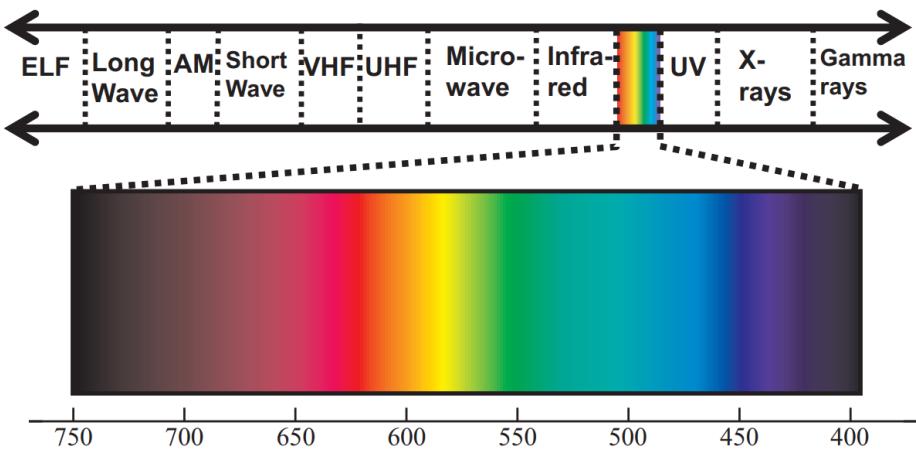
Light and image formation

- Light is emitted from light sources, travels along rays and interacts with object in the scene (e.g., surface reflection)
- Image created when light falls on the sensor (e.g., camera film)
- This process is described with **rendering equation**.

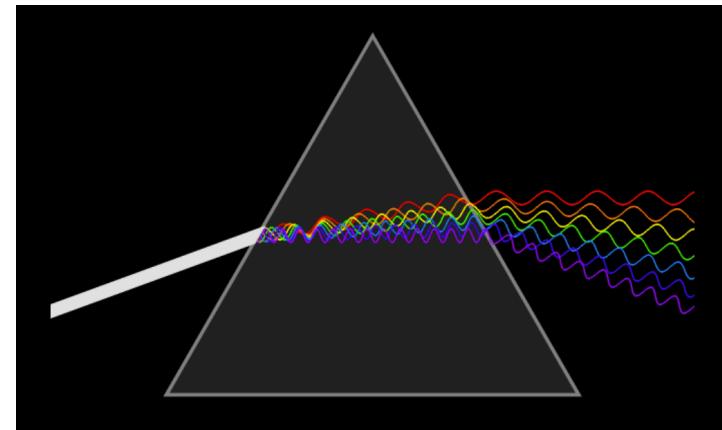


Light frequency and color

- Light is electromagnetic wave (or particle) visible by eye
 - Visible spectrum: wavelength 400 – 700 nm
 - Different wavelengths (frequencies) are perceived as different **color**
 - color is **perceptual (psychophysical)** phenomena rather than **physical** one



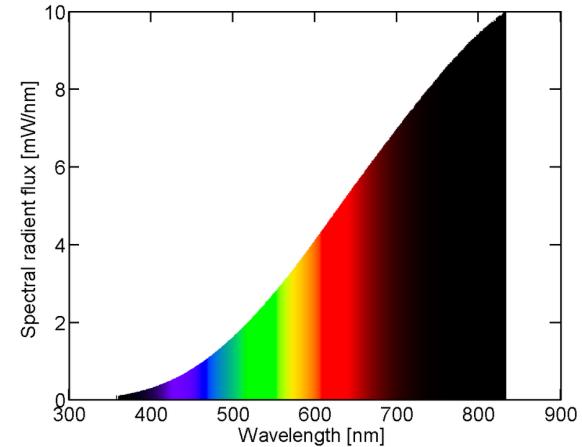
https://blog.selfshadow.com/publications/s2013-shading-course/#course_content



White light is made of all visible colors – adding all colors from light spectrum gives white light (color additivity, e.g. display)
https://en.wikipedia.org/wiki/Dispersive_prism

Spectral power distribution

- Light is generally characterized by **spectral power distribution (SPD)**
 - range and intensity of colors it produces at each wavelength in the visible spectrum
- Sun is reference for natural light
 - Standard: **daylight illuminant - D65** by International Commission on Illumination (midday sun Western/Northern Europe)
- **Color rendering index**
 - Ability of light source to reproduce color of object seen under D65, standard by CIE (Commission Internationale de l'Eclairage)
- White color is composed of all colors in equal amount: constant SPD
 - Doesn't exist in real world: CIE; E illuminant

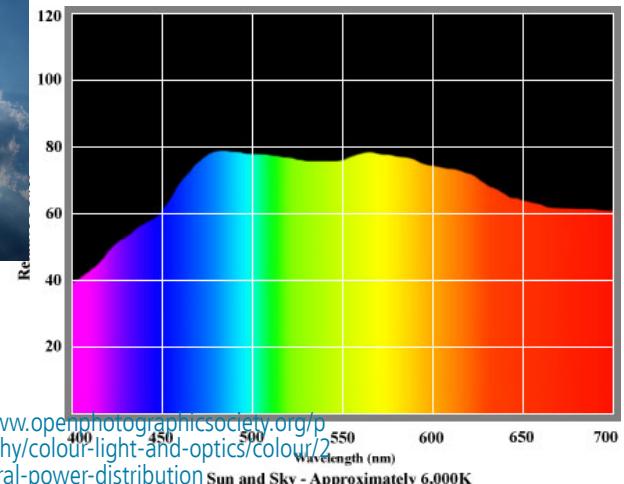


Perceive light as yellow since eye is less sensitive to red

https://en.wikipedia.org/wiki/Incandescent_light_bulb



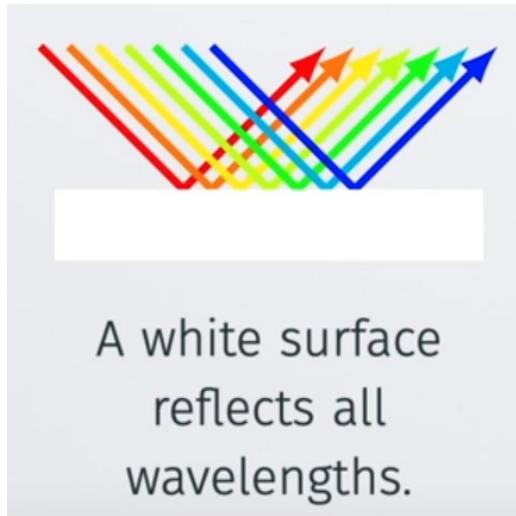
<https://commons.wikimedia.org/wiki/File:Sun-in-the-sky.jpg>



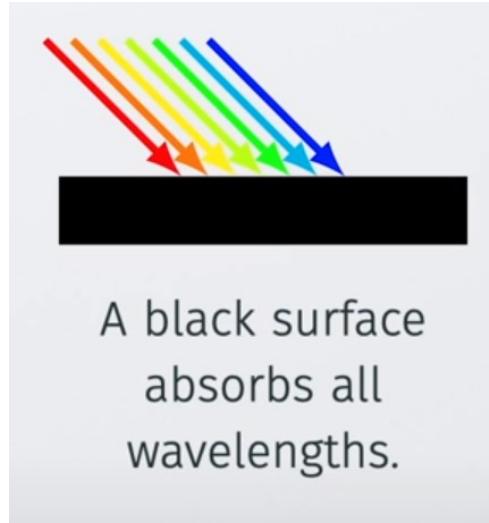
http://www.openphotographicsociety.org/b_photography/colour-light-and-optics/colour/2_03-spectral-power-distribution Sun and Sky - Approximately 6,000K

Material reflectance

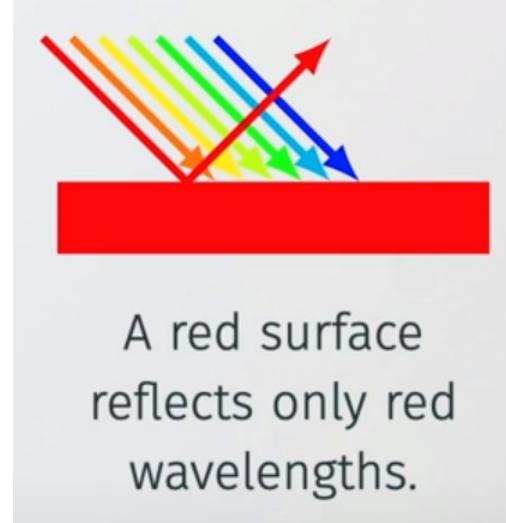
- Light falling on opaque surface will reflect and absorb
- Different materials are characterized with different **spectral reflectance curves**
 - Reflected light is perceived as color of object surface.



A white surface
reflects all
wavelengths.



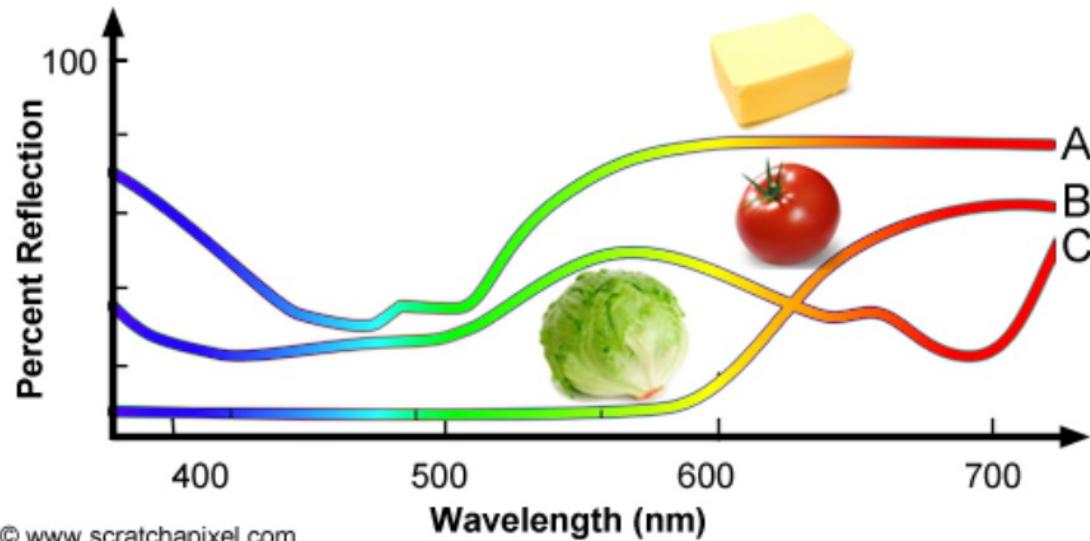
A black surface
absorbs all
wavelengths.



A red surface
reflects only red
wavelengths.

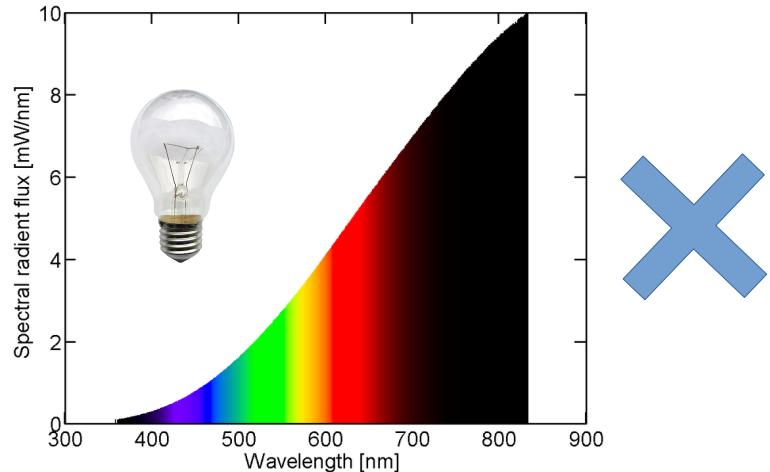
Spectral reflectance curve

- SPD of light source defines its color and intensity
- Spectral reflectance curve: SPD describing light reflected from objects
 - Fraction of the light reflected by the surface of the object over the amount of white light illuminating the object.

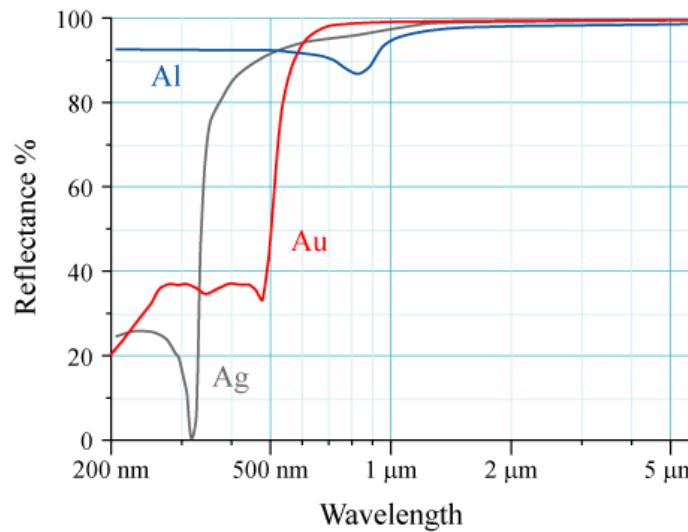


Material reflectance

- Reflected energy is calculated using reflectance R or absorbtion $A = 1 - R$, $A + R = 1$
 - R and A are spectral values, but in graphics we use RGB values

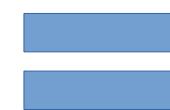


Power spectral distribution of light → incoming light energy (E)



Spectral reflectance curve → Material reflectance (R)

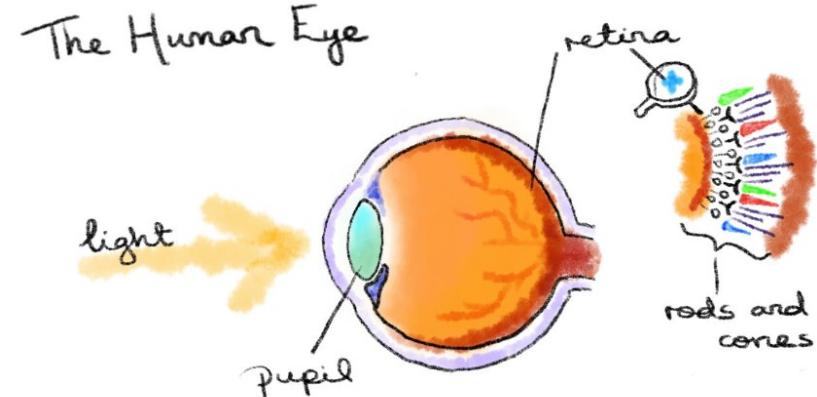
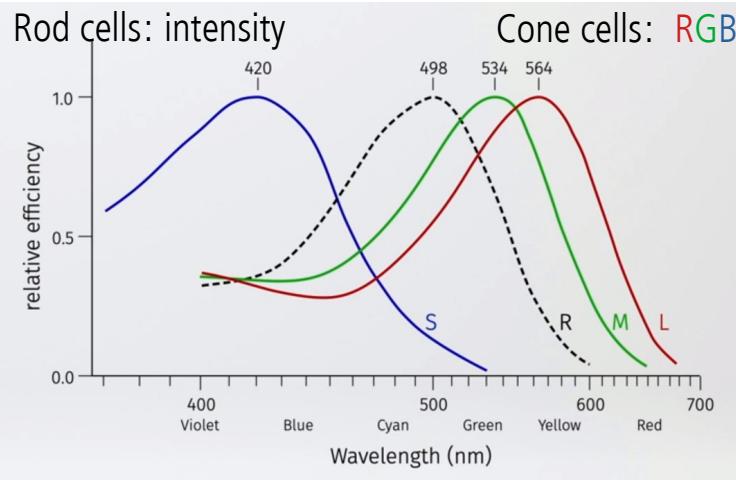
<https://en.wikipedia.org/wiki/Reflectance>



$$\text{Reflected} = E * R$$

Light and human visual system

- Human visual system; back of eye (the retina) is covered with light sensitive receptors – photo-receptive cells:
 - Cone cells – responsible for **trichromatic** color vision
 - “Active” in well lit environments (e.g., daylight)
 - Three types sensitive to: **red, green and blue** light
 - Rod cells – sensitive to smaller amount of light intensity but not reproducing colors
 - “Active” in low lit environments (e.g., night)

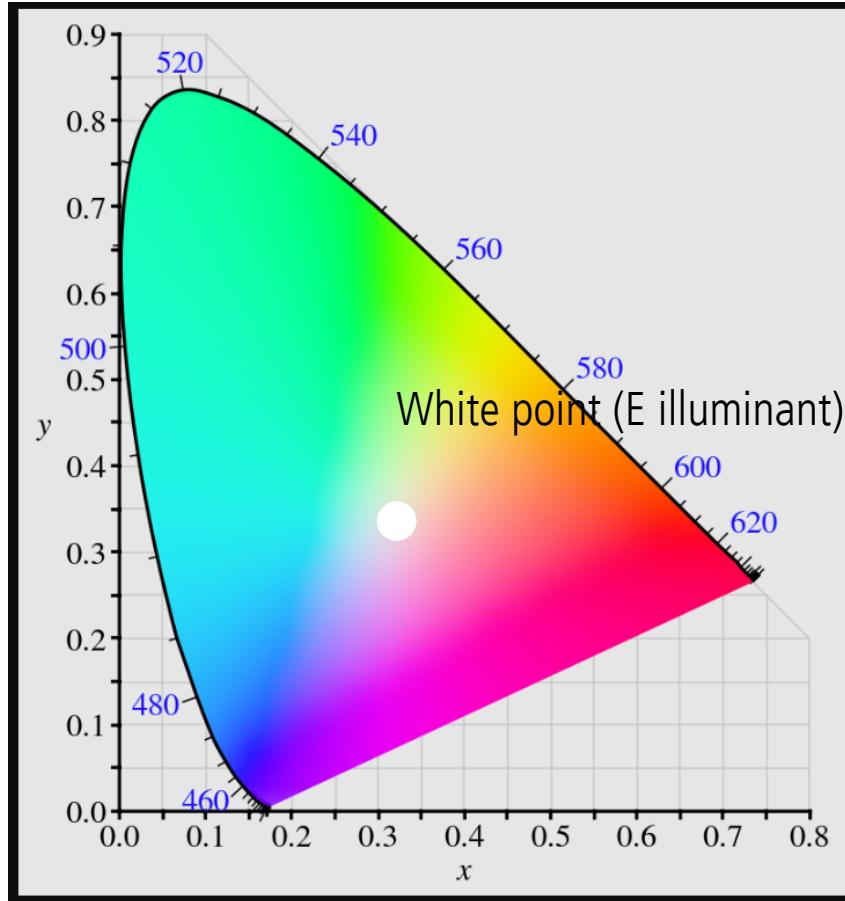


From SPD to color spaces

- Light from **light source** is characterized with spectral power distribution (SPD)
- Light reflected from **object** is defined with PSD; spectral reflectance curve
- Human visual system can not directly perceive spectral values, that is, PSDs
 - Specifying colors as wavelengths is impractical and in most cases not needed due to trichromatic human vision
 - Wave effects such as diffraction, interference and polarization will be ignored
- PSDs must be transformed in color models or color spaces suited for trichromatic vision
 - PSD is transformed in color hue and color brightness values which we can perceive
 - Geometric optics assumption: light is described with rays with intensity can color

Colorspace and gamut

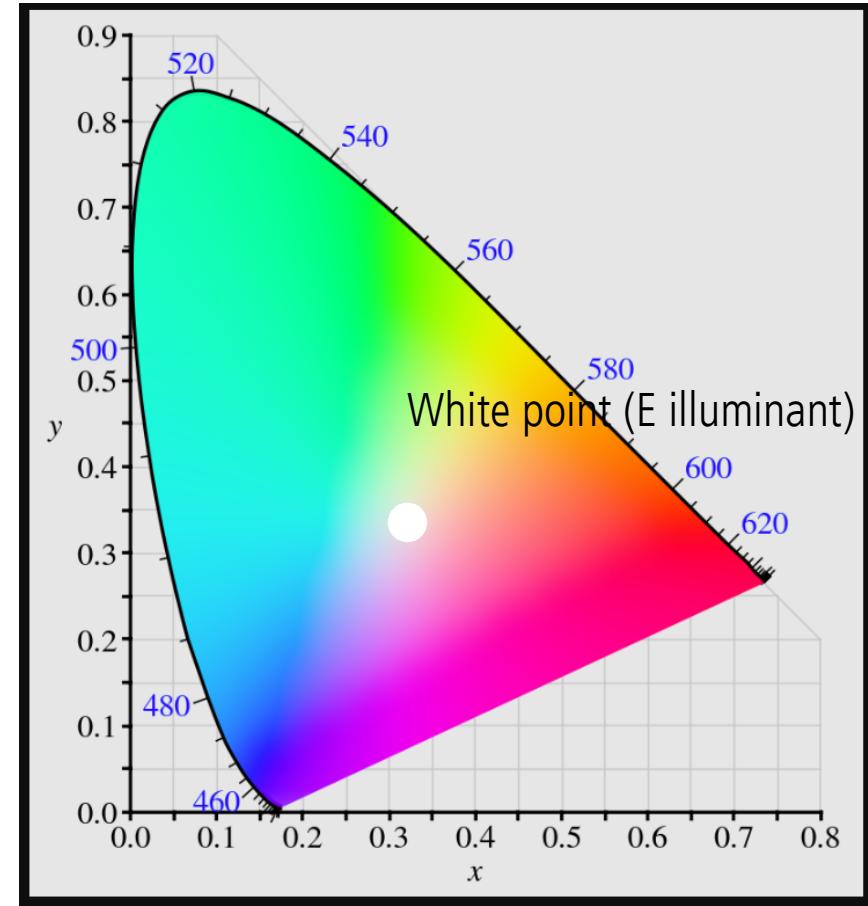
- **Colorspace**
 - Model to represent colors perceived by human visual system → reference for interpretation
 - Way of representing colors through numerical values
 - Specify color of pixels in an image
- **CIE XYZ** is foundation for all color models
 - RGB model is most important for computer graphics
- **Gamut**
 - Range of possible colors that can be represented by particular colorspace



Gamut of XYZ color space:
https://en.wikipedia.org/wiki/CIE_1931_color_space

XYZ colorspace

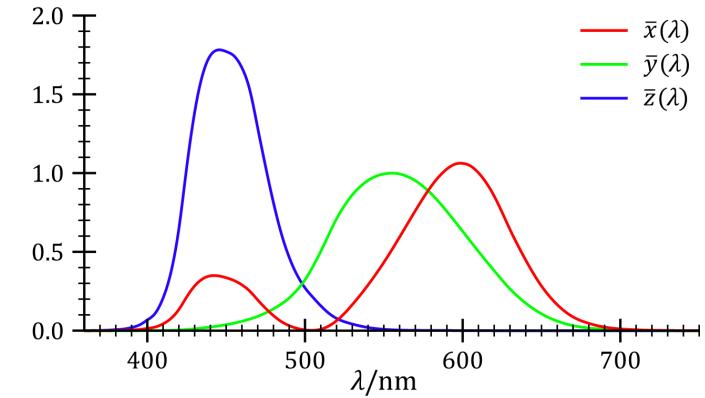
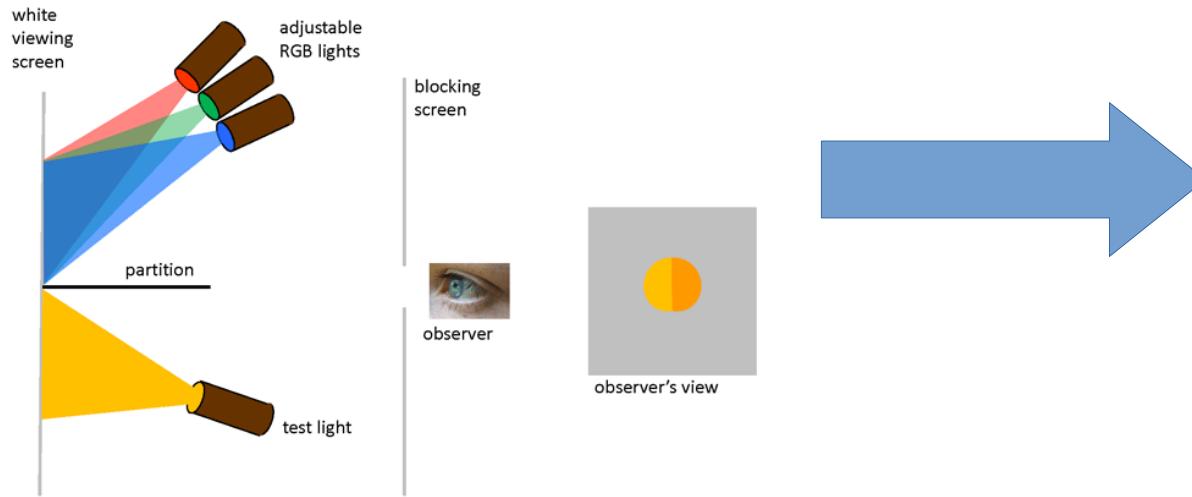
- Assumption: any color (in an additive color model) can be represented by a combination of red, green and blue light
 - Based on cone cells – trichromatic system
- Three curves: **CIE standard observer color matching functions**
 - Y curve is close to luminosity function → expressing the brightness
- Using color matching curve, PSD can be converted to X, Y and Z values
 - Artificial additive primary colors (X, Y, Z)



Gamut of XYZ color space:
https://en.wikipedia.org/wiki/CIE_1931_color_space

Color matching experiment

- Human visual system, cone cells can distinguish 3 signals: RGB → **trichromatic theory** of color vision
- Therefore, any color perceived by human visual system can be achieved by combining R, G and B signals.
- Colormatching experiment used that idea to form **color matching curves**.



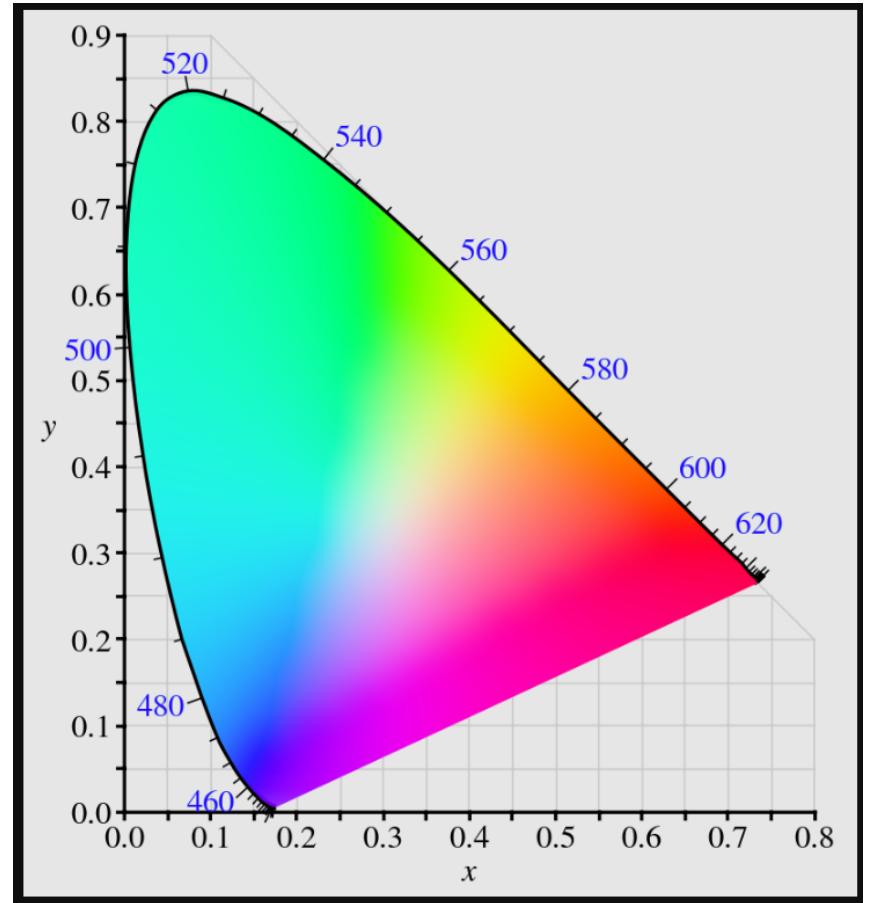
Color matching curves

XYZ colorspace

- $X + Y + Z = 1$, gamut can be visualized using:

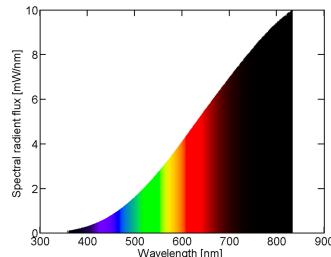
$$x := \frac{X}{X+Y+Z}, \quad y := \frac{Y}{X+Y+Z}$$

- Gamut of **XYZ color space** - gamut of human vision
 - Full range of possible colors visible by human visual system
- XYZ colorspace:
 - International, device-independent standard for color specification
 - System capable representing all colors human visual system can perceive
 - Gold standard for color storage
 - XYZ have to be transformed to RGB space to be visualized

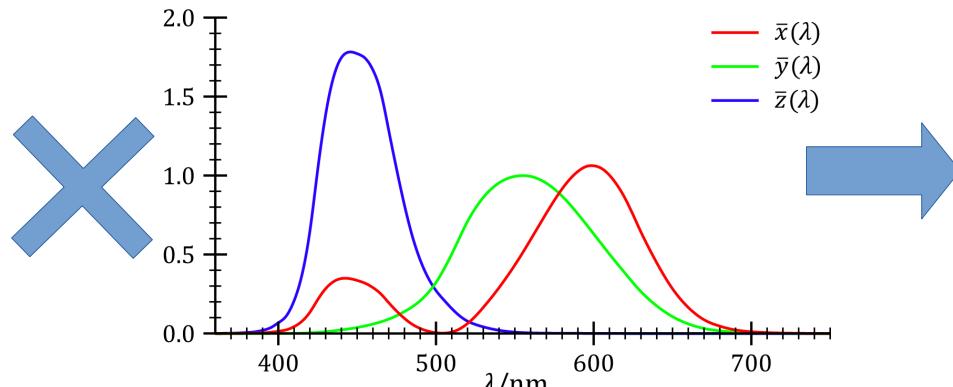


From PSD of light source to XYZ

- XYZ colorspace is obtained from PSD of light source using color matching curves
 - spectrum and the color-matching functions are discrete in a computer

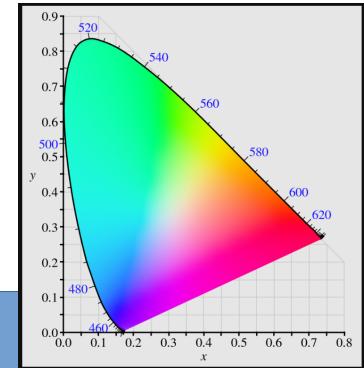


PSD of light source



https://en.wikipedia.org/wiki/CIE_1931_color_space

Color matching curves



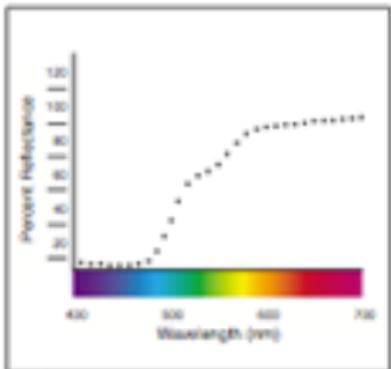
$$X = v_1$$

$$Y = v_2$$

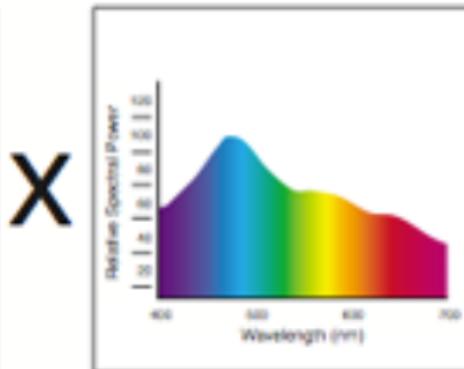
$$Z = v_3$$

From PSD of reflected light to XYZ

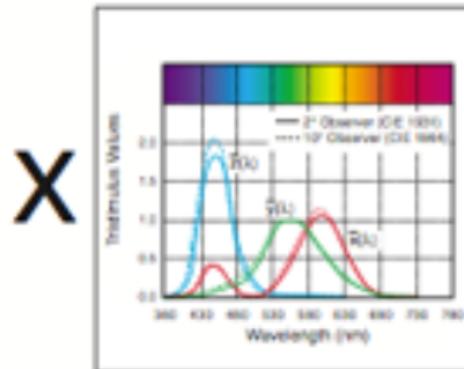
- Light reflected from object is represented as spectral reflectance curve
 - Therefore it must be multiplied by white light which caused the reflection and the standard is D65.



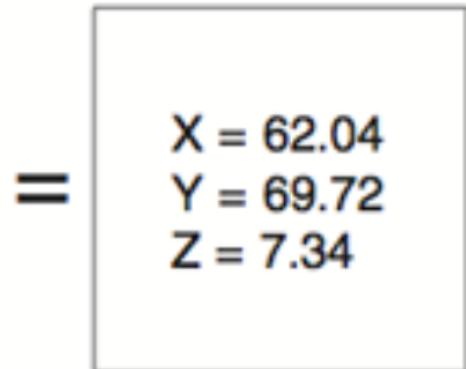
**Spectral
Reflectance
Curves**



D65 Illuminant



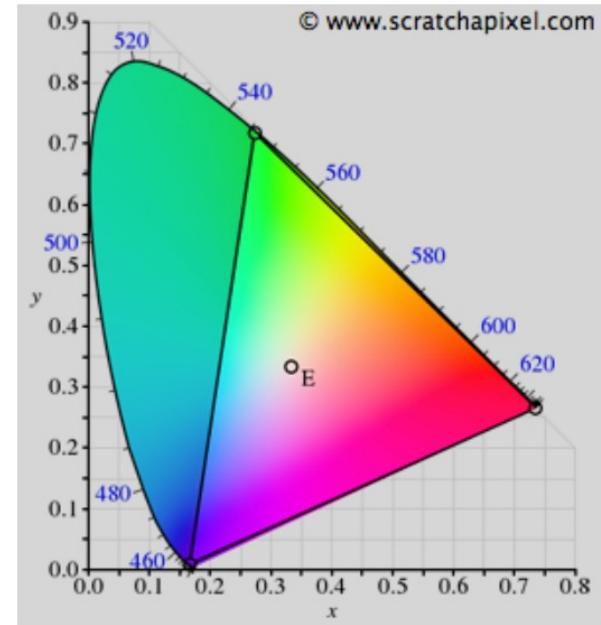
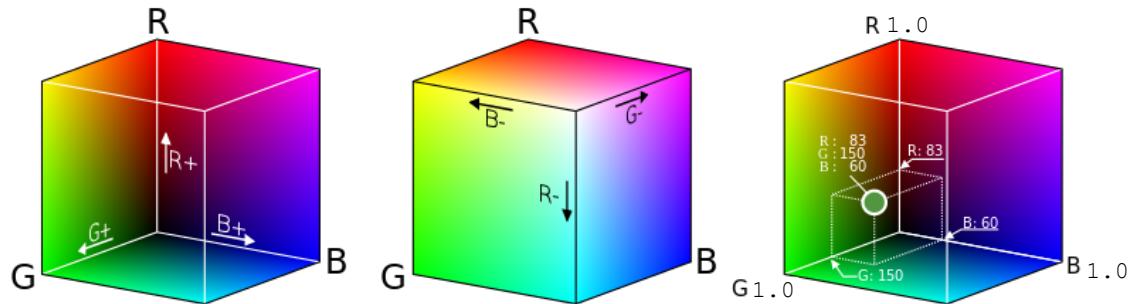
**CIE Color
Matching Functions**



**Tristimulus
Values**

RGB colorspace

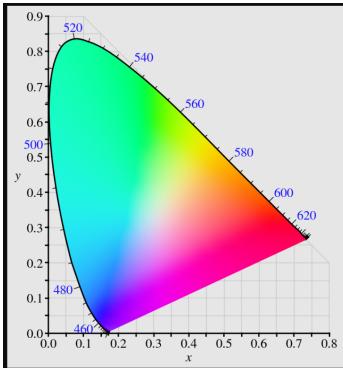
- Similar as human visual system: represent colors as simple sum of primary colors (R, G, B).
- Visualized as cube where R, G and B are placed on axis
- R,G,B intensities [0,1]



XYZ vs RGB gamut

XYZ to RGB

- XYZ and RGB colorspaces are defined by three coordinates
 - They define a 3D space, which can be plotted (cube for RGB)
- Conversion of color from XYZ to RGB can be seen as moving a point from one 3D space to another.
 - Linear transformation performed using 3x3 matrix
 - Color is not changed, only expressed in different colorspace



XYZ

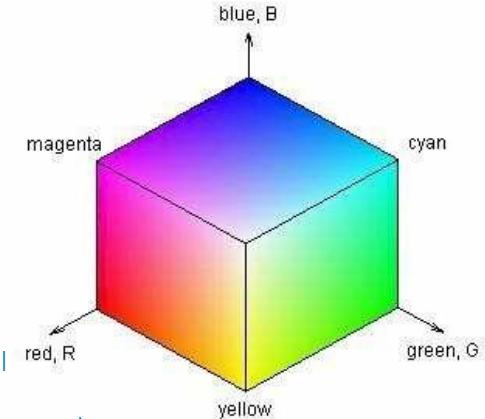
$$\begin{bmatrix} X \\ Y \\ Z \end{bmatrix} = \begin{bmatrix} 0.4124564 & 0.3575761 & 0.1804375 \\ 0.2126729 & 0.7151522 & 0.0721750 \\ 0.0193339 & 0.1191920 & 0.9503041 \end{bmatrix} \begin{bmatrix} R \\ G \\ B \end{bmatrix}$$

http://www.brucelindbloom.com/index.html?Eqn_RGB_XYZ_Matrix.html

<http://www.russellcottrell.com/photo/matrixCalculator.htm>

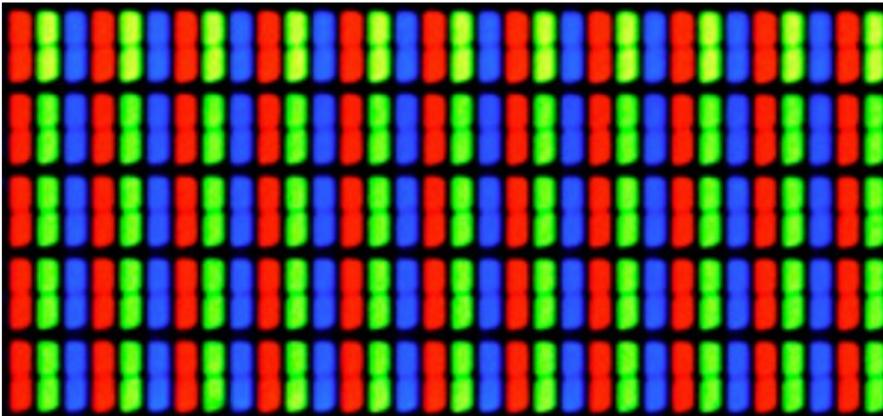
<https://www.oceanopticsbook.info/view/photometry-and-visibility/from-xyz-to-rgb>

<https://terathon.com/blog/rgb-xyz-conversion-matrix-accuracy/>



RGB colorspace

- Most **computer screens technology** is based on system similar to RGB color model.
 - Summing all RGB values gives white color
- RGB and XYZ are **linear** colorspace
- For rendering, RGB color space is often used
 - PSD is used in **spectral rendering** - more correct rendering, but expensive: computing XYZ from PSD for calculation takes time

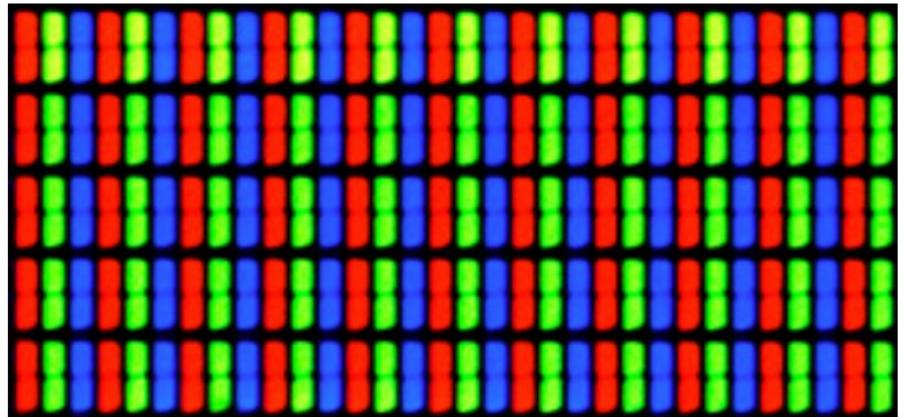
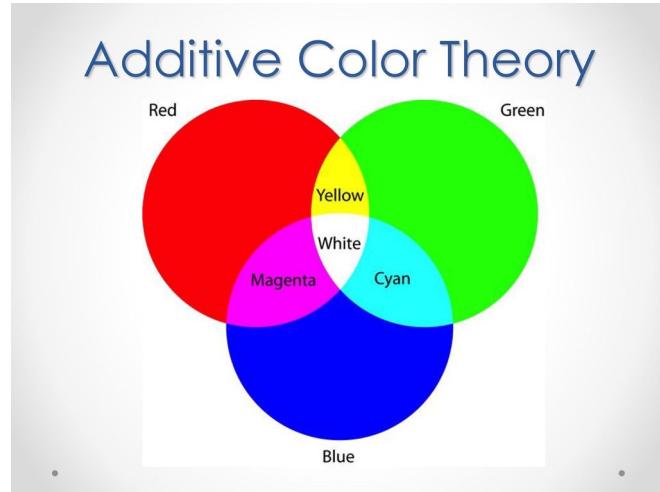


Close-up of LCD screen:

<https://jakubmarian.com/the-illusion-of-rgb-screens/>

Additive colorspaces

- Main colors: red, green, blue
- Any color is created by adding red, green or blue
- White color is made of red green and blue
- Usage: display devices, computer science, lighting technology → colors are created with light emission
- Display device (e.g., monitor) can further increase or decrease the overall global image intensity (brightness)



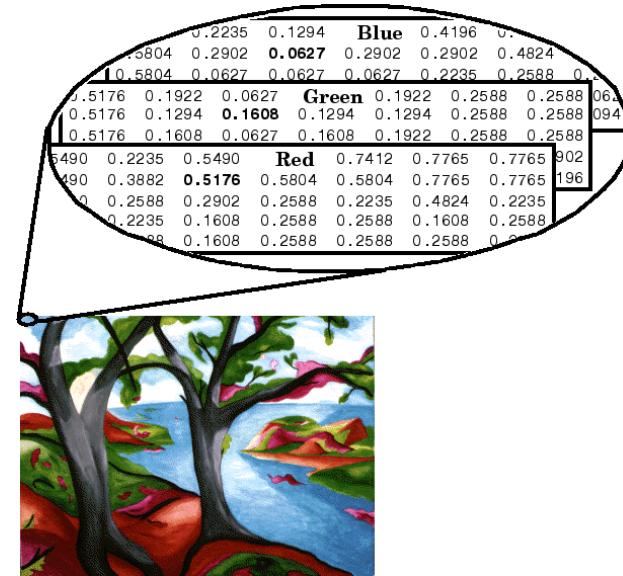
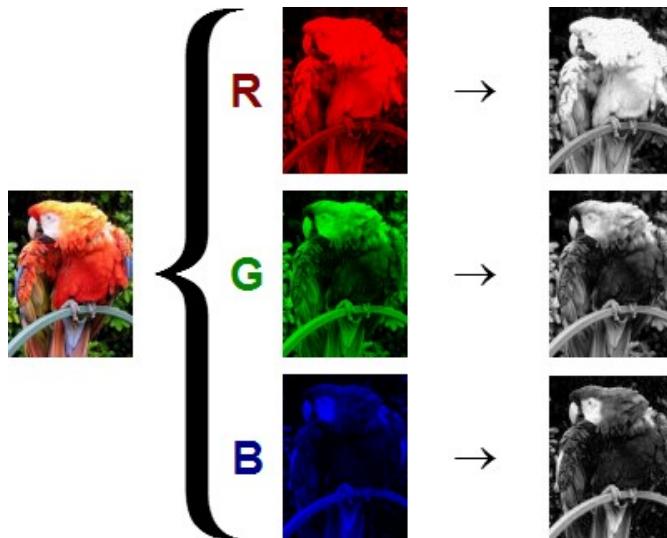
Close-up of LCD screen:

<https://jakubmarian.com/the-illusion-of-rgb-screens/>

<https://carsonparkdesign.com/what-makes-white-on-the-screen/>

Additive colorspaces

- Typical image formats store three intensities in three channels: R, G, B.
 - e.g., standard images: 8 bit per channel [0,255] range
 - $2^{(8 \times 3)} = 16.7$ million colors (term: true color)
 - e.g., HDR images: 10 or 12 bit per channel



https://www.mathworks.com/help/matlab/creating_plots/image-types.html

Subtractive colorspace

- Example: CMY colorspace
- Main colors: **cyan, magenta, yellow** - white light is made by subtracting
- Usage: art, painting, printing
- Color from canvas is not a result of light emitted from canvas, but light reflected from canvas
 - When white light hits canvas, what is not absorbed is reflected

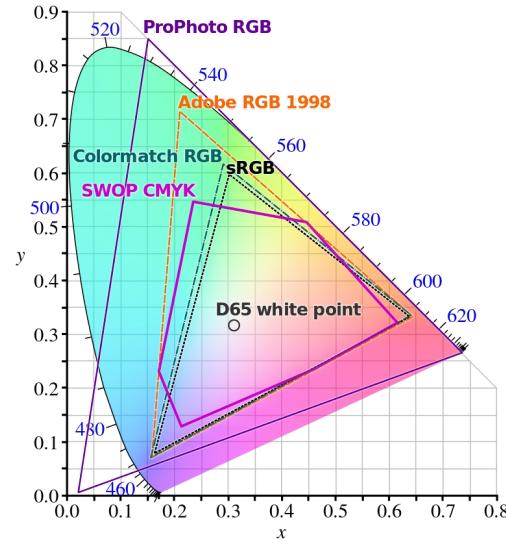


<https://www.greenleafblueberry.com/blogs/news/modern-primary-colors>

Colorspaces and transformations

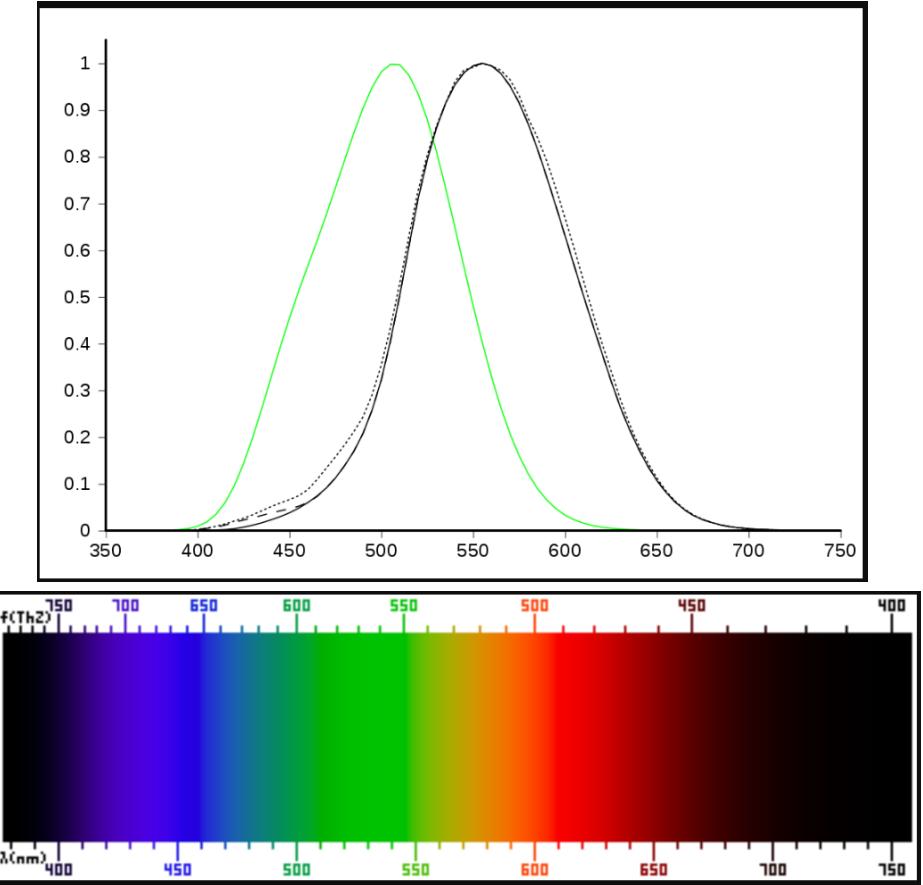
- RGB and XYZ are linear colorspaces and can be transformed using 3x3 matrix
- There are other colorspaces such as sRGB, HSV, HSL are non-linear
- 3x3 matrices can not be used to convert between non-linear and linear colorspace (and vice versa)
 - Any non-linear colorspace must be converted to linear colorspace before 3x3 matrix transformation.
- Different colorspace represent different gamut of colors
- RGB and CMY have small gamut! → device limitations

https://en.wikipedia.org/wiki/Color_space



Color brightness

- 555 to 560 nm (green towards yellow) are perceived by the human eye as being the brightest
 - Blue color: the dimmest
 - Green color: the brightest
- How bright color appear to human eye: **luminosity function**
 - Average sensitivity of human eye to different wavelengths for well lit conditions - photopic (cone rods).
- Brightness or lightness is present in **HSV** color model
- **Photometry**: science of measuring light taking in account human visual system
 - CIE photopic spectral luminous curve

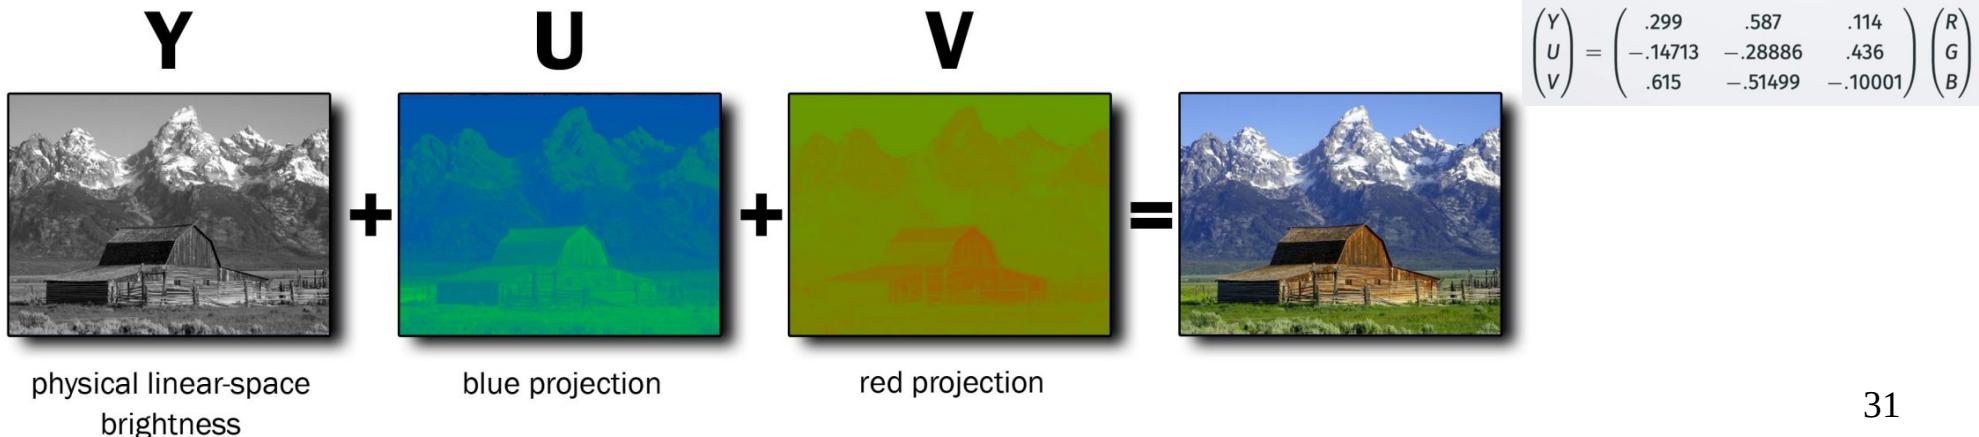


Luminosity function for photopic – green - (well lit) and scotopic (low lit) – black.

https://en.wikipedia.org/wiki/Luminous_efficiency_function

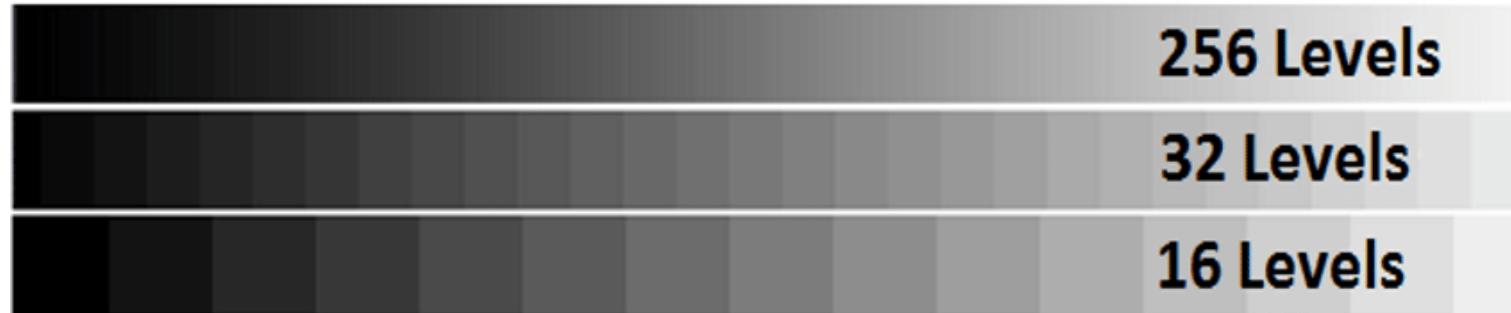
Brightness (luminance)

- Human eye is much more sensitive to spatial variations in brightness (gray scale) than to spatial variations in color
- Therefore, YUV color space can be particularly useful to highlight the importance of gray scale brightness
 - Y – brightness (luminance)
 - U and V – color (chroma)
- YUV can be obtained from RGB using linear transformations
- YUV is useful for compression: Y channel should get the most resolution (e.g., JPEG and H.264/H.265 formats)



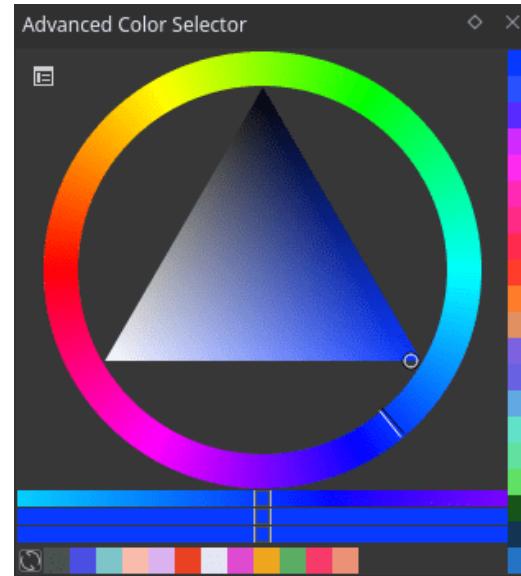
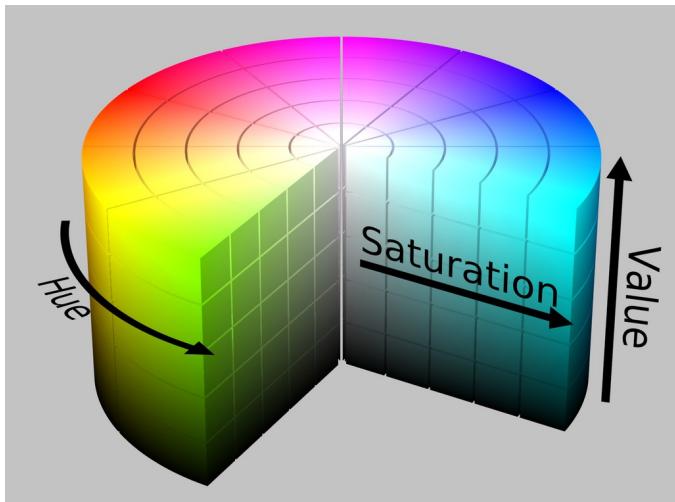
Brightness (luminance)

- Human eye can very well see small changes in brightness
- Many levels for brightness is needed to describe it without visual discontinuities
- **Dynamic range:** ratio between image with brightest white level and image with darkest black level
 - Very high ratios are possible in real world → image and display limitations!



HSV colorspace

- RGB colorspace is not always intuitive for “color picking”
- HSV is transformed from RGB:
 - H – Hue - color (rainbow of colors), wheel, given as angle
 - S – Saturation - how white the color is, how much color is mixed with white
 - V – Value – brightness of color
- Non-additive colorspace

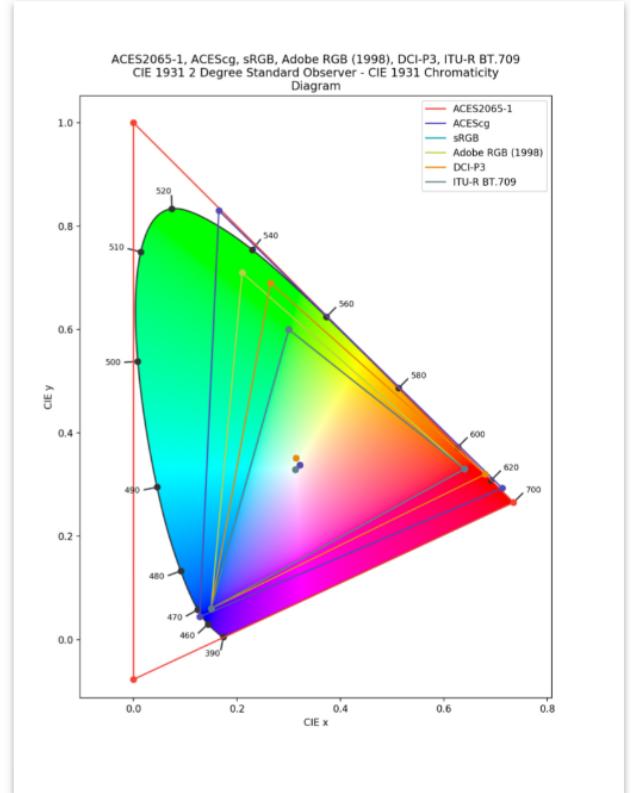


HSV is often used for colorpicking in modeling environments or painting programs:

https://docs.krita.org/en/reference_manual/dockers/advanced_color_selector.html

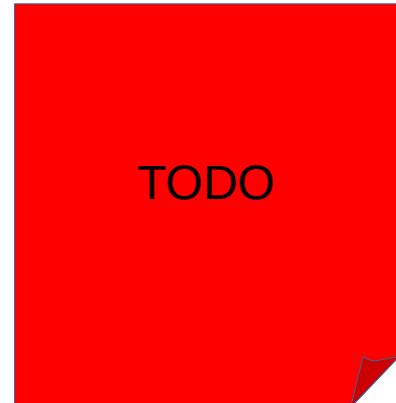
Note: ACES colorspace

- Academy Color Encoding Specification (ACES) is developed by Academy of Motion Picture Arts & Science technology committee
- Designed to cover the full gamut
- Goal: images should always look the same regardless of used camera and display device.
- Images in ACES color space can't be directly displayed to the screen. Conversion steps:
 - Reference rendering transform
 - Output device transform: depends on display device (computer monitor, digital projector, etc.)
- Becoming defacto standard for professional production



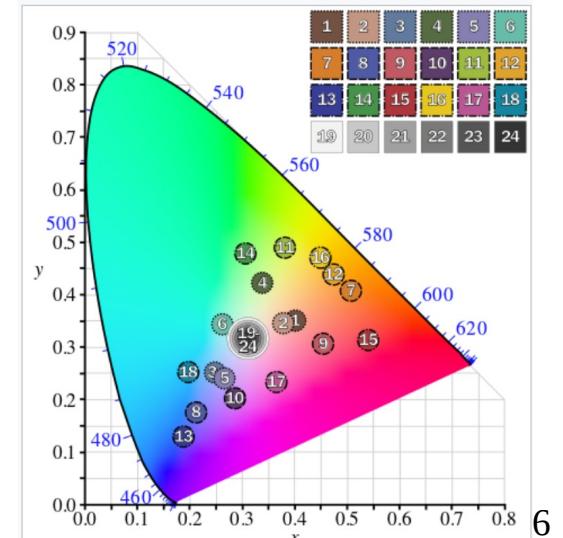
Color in production

- Digital drawing and image manipulation:
 - Krita: <https://krita.org/en/item/krita-features/>
 - Gimp
- 3D modeling:
 - Blender
- Game engines:
 - Godot



Note: Macbeth ColorChecker

- Colors have spectral reflectances mimicking materials which are often photographed.
 - Human skin, sky, flowers, etc.
- Goal: consistent color appearance under varying lighting conditions
 - Reference for tweaking exposition and color settings of camera



Light characterization

- Quantifying light:
 - **Radiometry** – measurement of radiation, physical transmission of light
 - **Photometry** – measurement related to human visual system
 - Summary measures: computed from radiometric quantities.
 - **Colorimetry** – Color perception

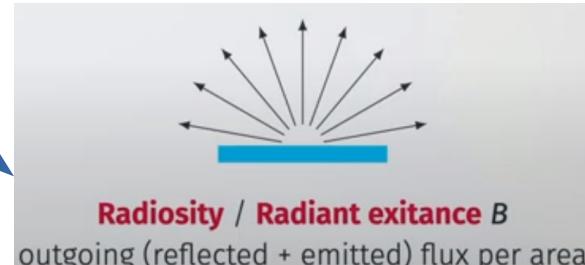
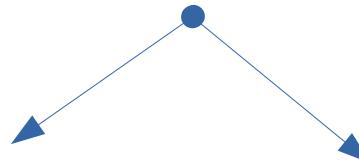
Radiometric quantities

- Radiant energy (Q) – **energy** carried by light (photons) [J]
- Radiant flux (Φ) – power carried by light - **energy over time** [$W=Js^{-1}$]

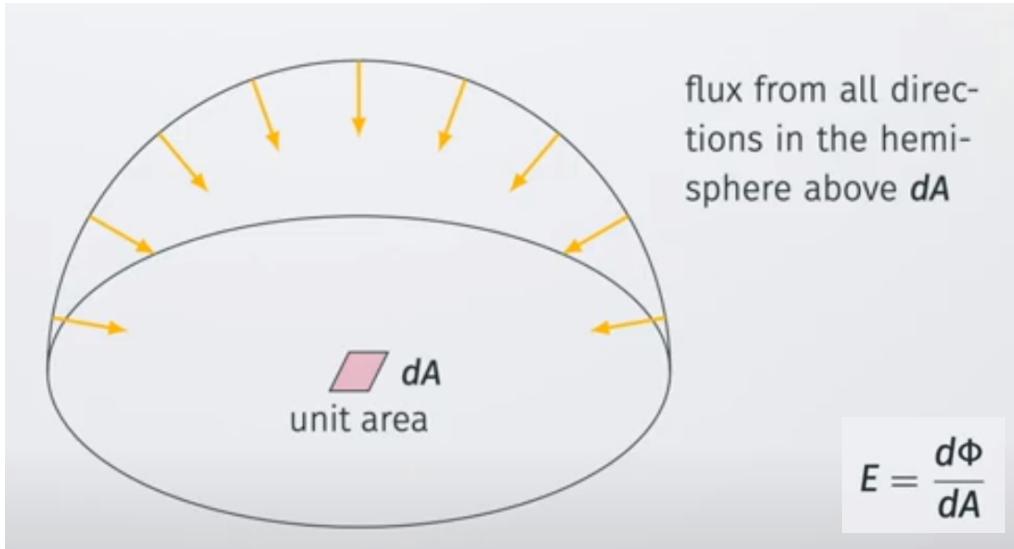
$$\Phi = \frac{dQ}{dt}$$

- Flux density (E) – **radiant flux over area** [Wm^{-2}]

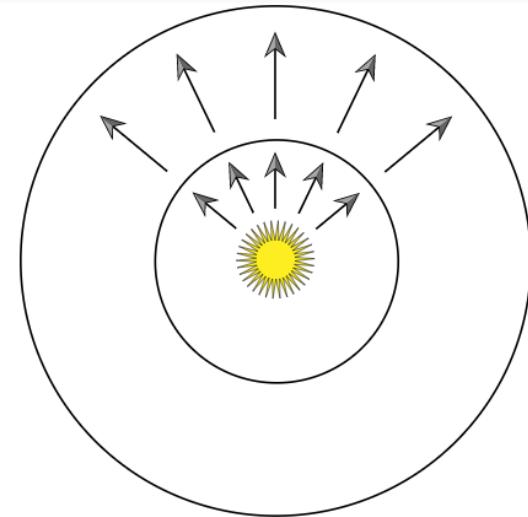
$$E = \frac{d\Phi}{dA}$$



Radiometric quantities



Irradiance – overall radiant flux (light flow) incoming onto surface.

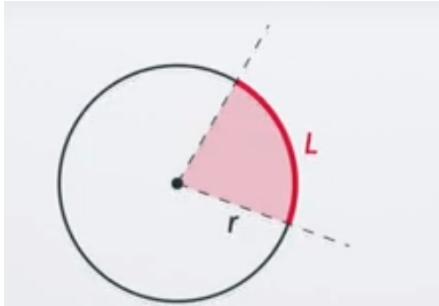


Radiant exitance – radiant flux (light flow) leaving a light source.

Amount of energy received from a light falls off with the squared distance from the light.

https://www.pbr-book.org/3ed-2018/Color_and_Radiometry/Radiometry

Radiometric quantities

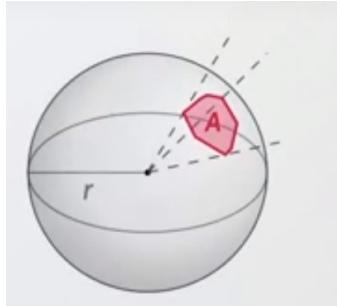


Planar angle: angle of segment of length L on a circle of radius r

$$\alpha = \frac{L}{r}.$$

Angle of full circle:

$$\alpha = \frac{2\pi r}{r} = 2\pi \text{ [rad = radians].}$$

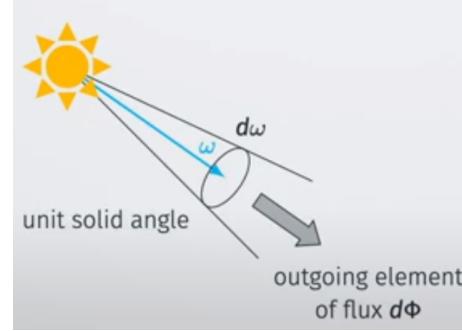


Solid angle: area of size A on a sphere of radius r (3D extension of the concept of plane angle)

$$\Omega = \frac{A}{r^2}.$$

Angle of full sphere:

$$\Omega = \frac{4\pi r^2}{r^2} = 4\pi \text{ [sr = steradians].}$$



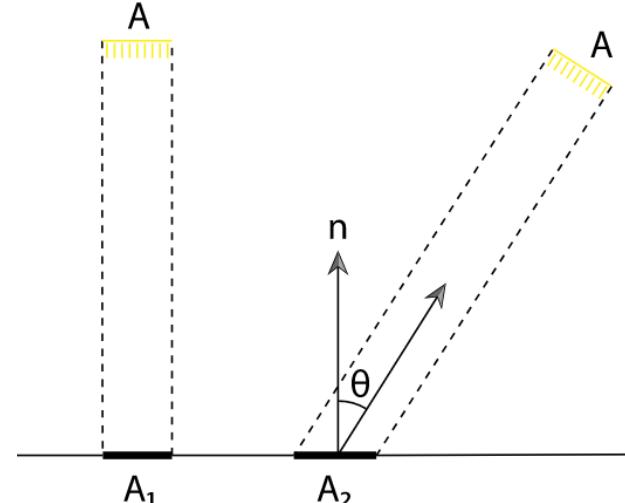
Radiant intensity (I) – radiant flux with respect to **solid angle** [W / sr]

- Radiant flux of light source in specific direction.
- Radiant flux reaching certain point from some point.

Radiometric quantities

- **Radiance (L)** – radiant flux per unit solid angle, per unit projected area [$\text{W} / (\text{m}^2 \text{ sr})$]
 - Incident on surface, emerging from surface or passing through surface in certain direction
- $L(p, \omega)$ – radiance of small area at p along small solid angle $d\omega$ in direction ω
 - Elementary description of light
 - Electromagnetic radiation in single ray

$$L = \frac{d^2\Phi}{d\omega dA_p} = \frac{d^2\Phi}{d\omega dA \cos \theta}$$

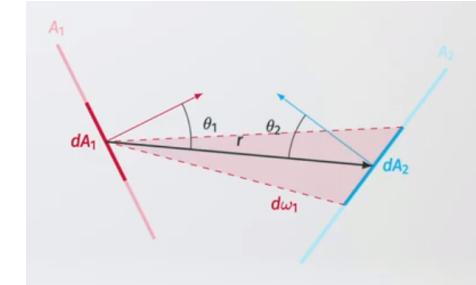
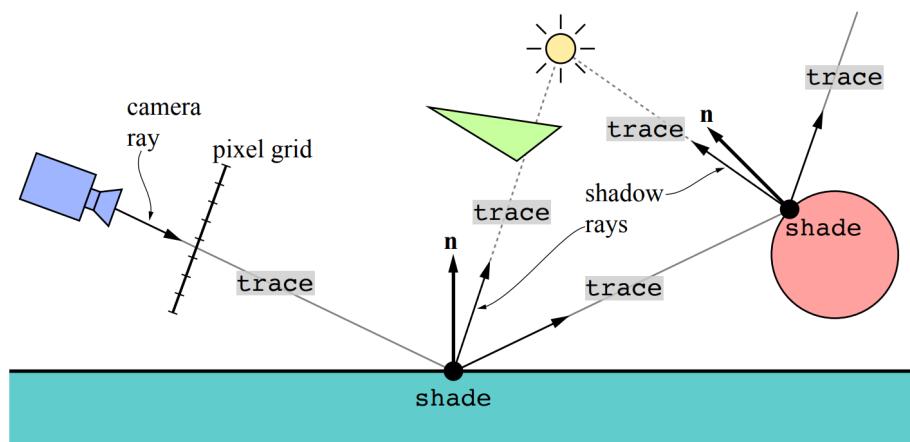


Radiance area is measured in a plane perpendicular to ray, if surface has orientation then cosine correction factor must be used.

* Image-based rendering uses this concept, namely light field technique.

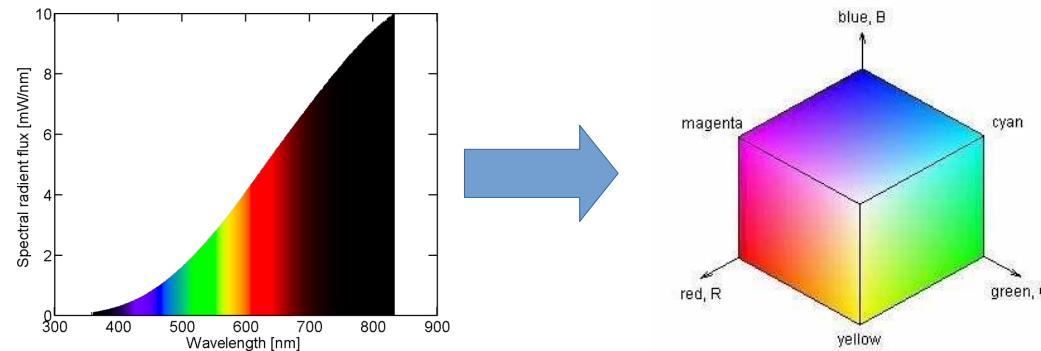
Radiometric quantities

- Radiance is what cameras measure – prime importance for rendering
 - Rendering can be seen as evaluating radiance for each camera ray
- The purpose of shading is to calculate radiance along a camera ray
 - Radiance is simplified with color and intensity
- Conservation of radiance along rays
 - If two points on two surfaces can “see each other”, then radiance outgoing from one surface and incoming to another surface is same and vice versa → light can be traced between surfaces until it reaches the sensor
 - Radiance is not affected by distance if light is traveling in a vacuum
 - Intensity is lower for surfaces which are more distant since light covers less area



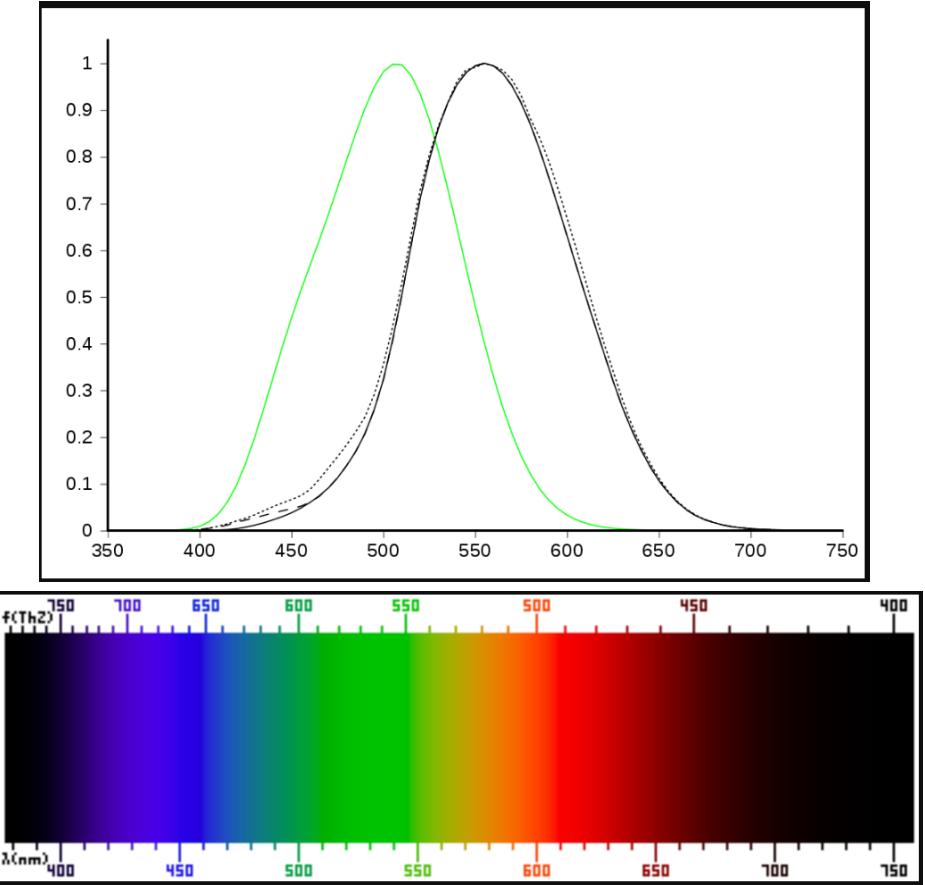
Radiometric quantities and PSD

- Most light waves contain mixture of many different wavelengths – **polychromatic** (light with one wavelength is called **monochromatic**).
 - Visualized using spectral power distribution (SPD) – plot showing how light's energy is distributed across wavelengths.
- All radiometric quantities have spectral distributions
 - Using full SPDs for rendering is complex*, often radiometric quantities are represented as RGB triplets.



Photometry

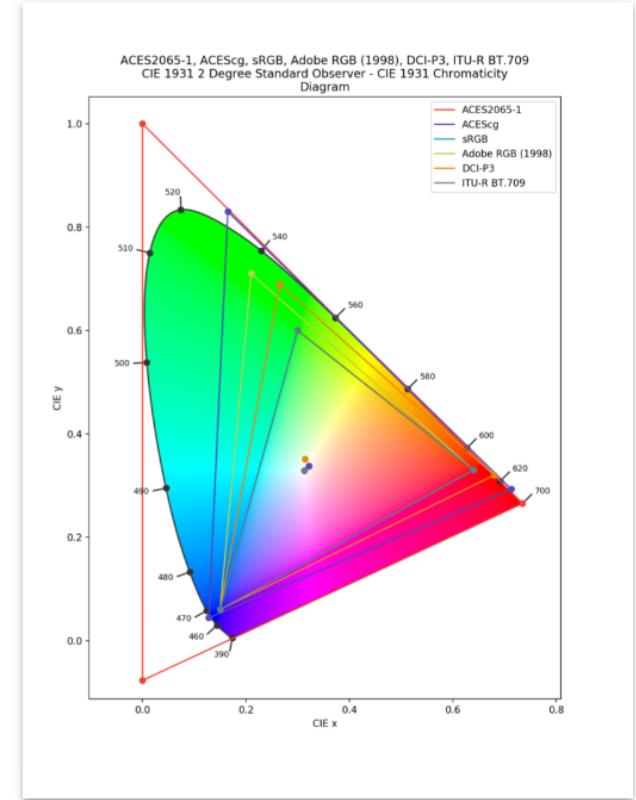
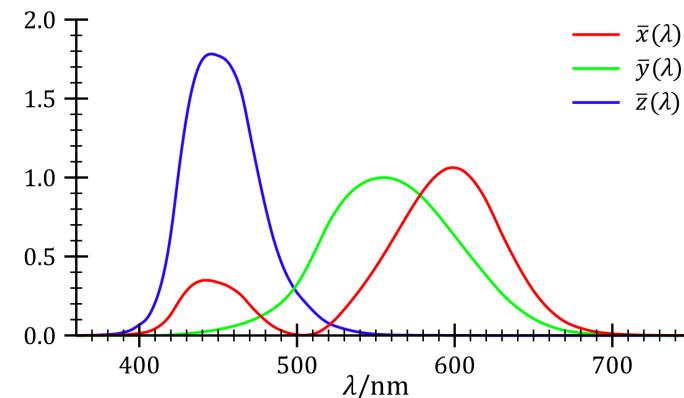
- Photometry is field like radiometry but it takes visual system and perception in account.
- Each radiometric quantity has equivalent photometric quantity
 - Radiant flux → Luminous flux (lumen - lm)
 - Irradiance → Illuminance (lux - lx)
 - Radiant intensity → Luminous intensity (candela - cd)
 - Radiance → Luminance (cd / m^2)
- Radiometric quantities are converted to photometric by using **CIE photopic spectral luminous curve**.



Luminosity function for photopic – green - (well lit) and
scotopic (low lit) – black.
https://en.wikipedia.org/wiki/Luminous_efficiency_function

Colorimetry

- Colorimetry: relationship between PSD and perception of color
- CIE standard observer color matching functions
 - PSD → XYZ colorspace
- Define colorspaces and their gamut for visualization
- RGB color spaces are most interesting for rendering:
 - SRGB (linear colorspace with sRGB primaries and white point)



Rendering with RGB

- RGB is perceptual rather than physical quantity
- RGB in Physically based rendering is a category error
 - Correct would be to use spectral quantities for rendering and RGB conversion in the end for visualization
- For interactive and appearance-based applications (e.g., film) RGB rendering is acceptable.
 - Object surface reflection is defined using RGB rather than spectral reflectance curve
 - Light sources are described with RGB rather than PSD
- For predictive simulations rendering spectral quantities must be used.



Tristimulus vs spectral rendering:
<https://hal.inria.fr/hal-03331619/document>

Recap

- Light has physical but also perceptual properties.
- Discussion on radiometry, photometry and colorimetry was needed to understand light and its perception:
 - Color (hue)
 - Intensity
- Light in computer graphics is represented using:
 - RGB values in $[0, 1]$
 - Intensity, multiplier of RGB values

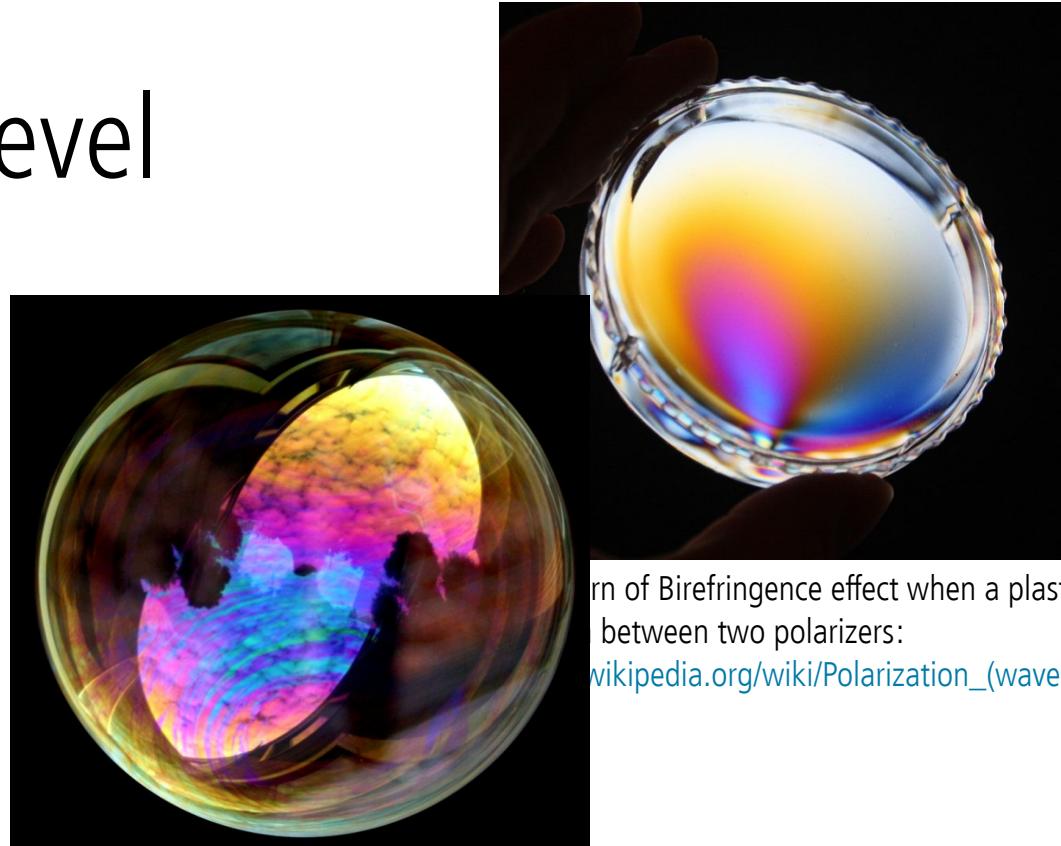
Real world light and sources of light

Light description

- Light can be described both on:
 - Microscopic scale: wave optics
 - Macroscopic scale: **geometric optics**

Light: microscopic level

- Light is quantized – it comes in individual and indivisible packets called **photons**
- At same time light is **wave-like**
- Some effects of light-matter interaction can be only explained (and thus modeled*) using **wave optics**:
 - Interference
 - polarization
 - Diffraction
 - Iridescence
 - Pleochroism and birefringent
 - Etc.
- Also, knowing microscopic background often helps in getting intuition for some shading and light transport methods**



Learn of Birefringence effect when a plastic box between two polarizers:
[wikipedia.org/wiki/Polarization_\(waves\)](https://en.wikipedia.org/wiki/Polarization_(waves))

White light interference in a soap bubble – iridescence phenomena (thin-film interference) -
https://en.wikipedia.org/wiki/Wave_interference

* It is important to note that complexity of model level depends on application. Often, in computer graphics wave effects can be “faked” to achieve desired appearance:

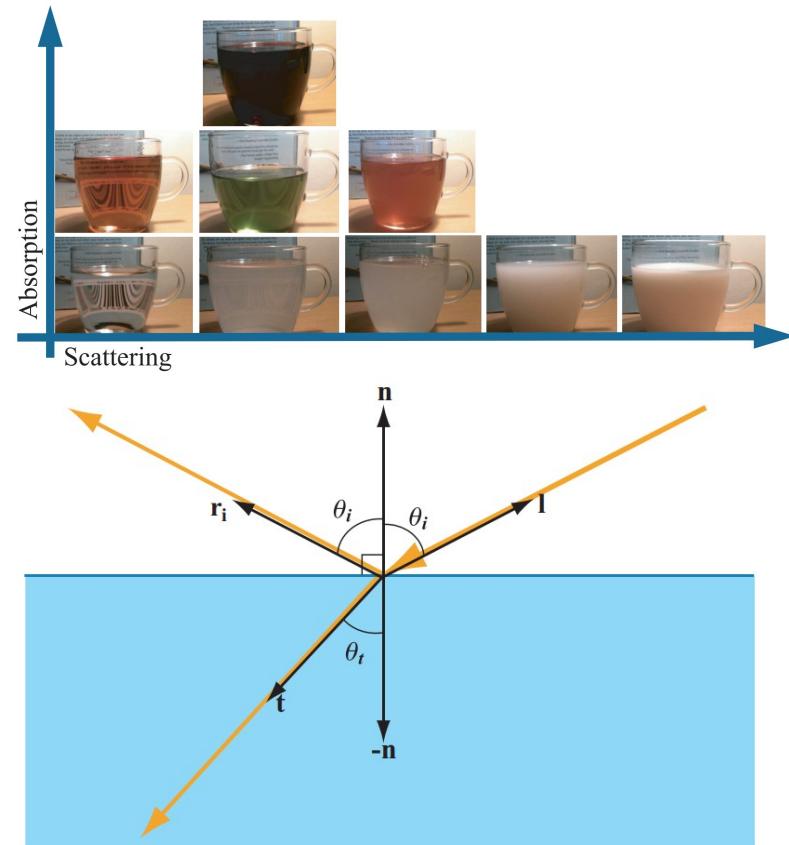
<https://developer.nvidia.com/gpugems/gpugems/part-i-natural-effects/chapter-8-simulating-diffraction>. On the other hand, physically based methods for achieving wave-effects are actively researched area

https://sites.cs.ucsb.edu/~lingqi/publications/202203_practical_plt_paper_lowres.pdf

** For example photon mapping: http://web.cs.wpi.edu/~emmanuel/courses/cs563/write_ups/zackw/photon_mapping/PhotonMapping.html

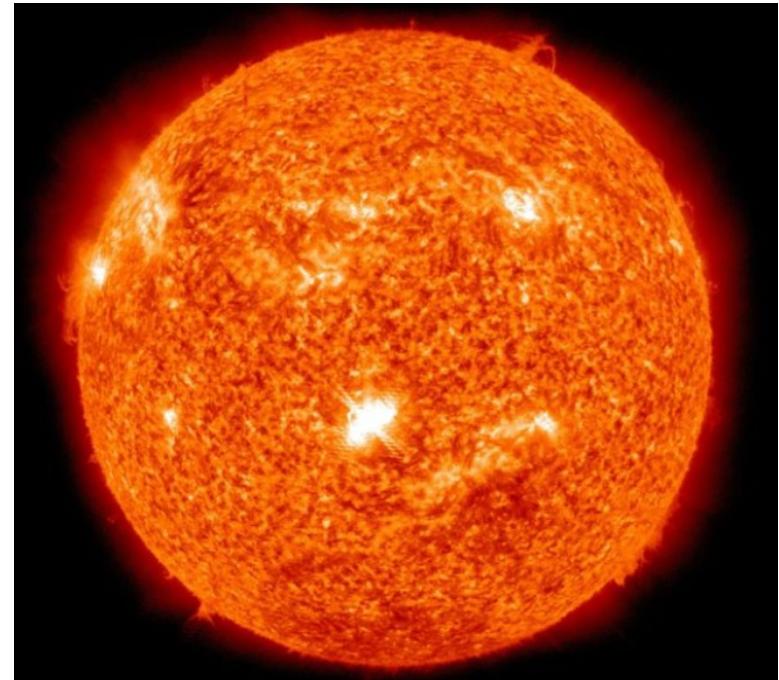
Light: macroscopic level

- Light in computer graphics is often modeled at macroscopic level – **geometric optics**.
 - Assumption: features with which light interacts are relative to light wavelength (or larger)
 - Light is described with rays.
 - Microscopic light behavior is described with **index of refraction (IOR)**
- In medium of constant IOR, light energy travels in a straight line
 - It only can get **absorbed**
- Light passing from one to another medium with different IOR changes speed
 - This is described with light **scattering**
 - Interface between media is represented by **surface**
- Scattering of light on surface is described with reflection and refraction described with Snell law (light direction) and Fresnel equations (light energy).



Sources of light

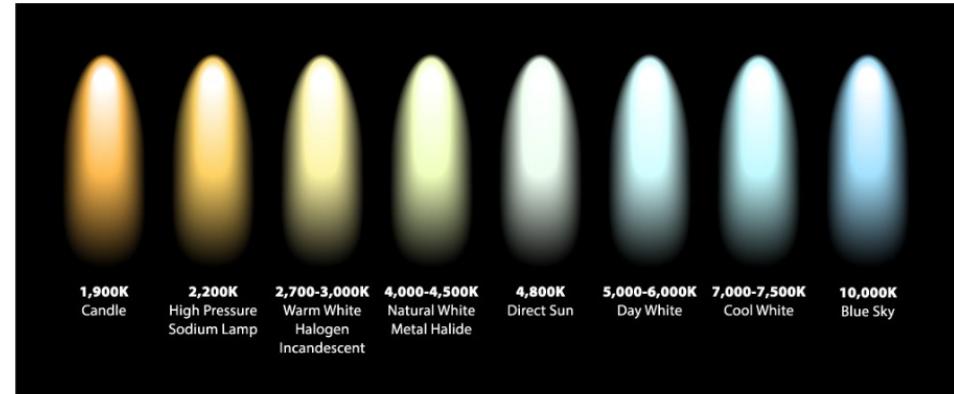
- In a real world, **every light source has a physical shape and size**, e.g., Sun or very hot objects.
- Motion of atomic particles that hold electrical charge causes objects to emit electromagnetic radiation over a range of wavelengths (Maxwell's equations)
- Different way how energy is converted into electromagnetic radiation:
 - Incandescent (tungsten) lamps
 - Halogen lamps
 - Gas-discharge lamps
 - LED lights



<https://education.nationalgeographic.org/resource/sun>

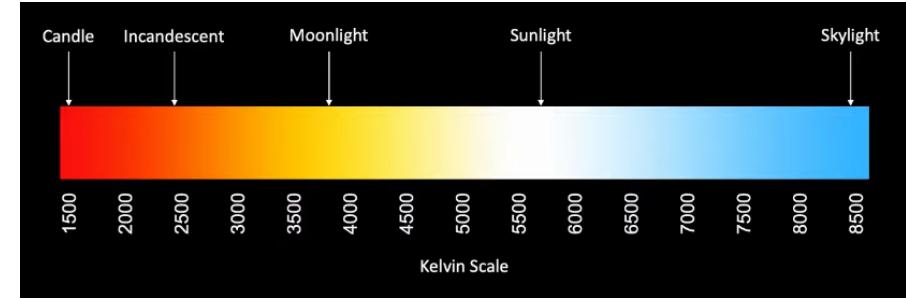
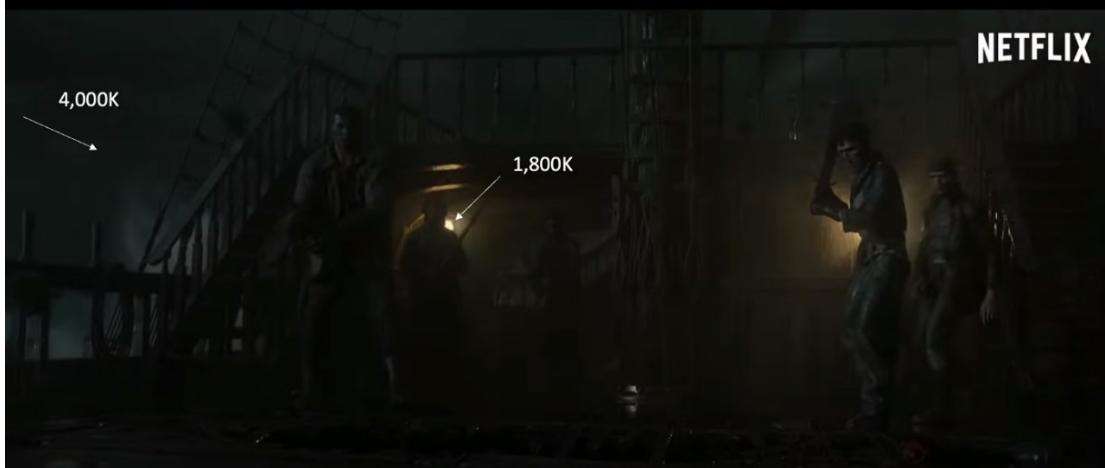
Color temperature

- Object which is heated to certain temperature emits light
- Color of emitted light depends on temperature
- Black body radiation



https://www.inlineelectric.com/color_temperature

Color temperature

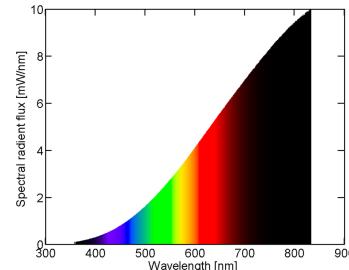


Secrets of photorealism:

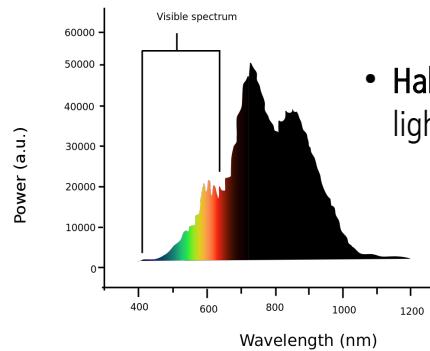
https://www.youtube.com/watch?v=Z8AAX-ENWvQ&t=2s&ab_channel=Blender



Sources of light

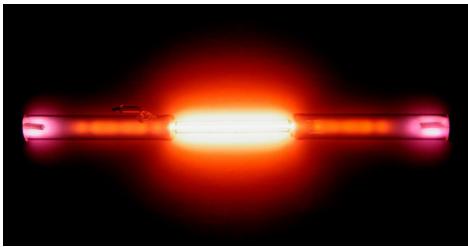


- **Incandescent (tungsten) lamps:**
 - flow of electricity through tungsten filament heats it up and causes it to emit electromagnetic radiation with distribution of wavelengths depending on filament temperature.
 - A frosted glass enclosure is often present to absorb some of the wavelengths.



- **Halogen lamps:** tungsten filament is enclosed in halogen gas. Part of the filament in an incandescent light evaporates when it's heated. Halogen causes evaporated tungsten to return to filament

Sources of light



Helium: https://en.wikipedia.org/wiki/Gas-discharge_lamp



LED lights:
https://en.wikipedia.org/wiki/Light-emitting_diode

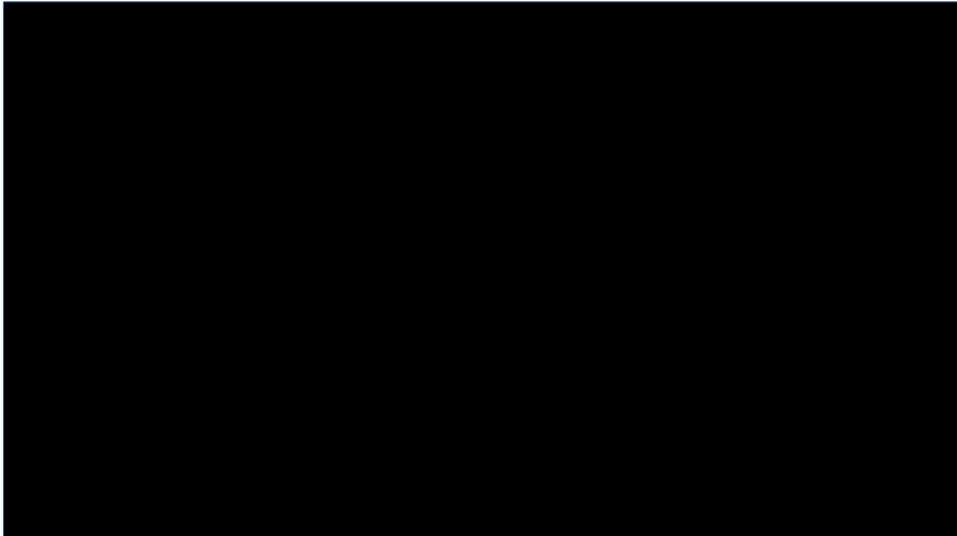
- **Gas-discharge lamps:** passing electrical current through hydrogen, neon, argon, or vaporized metal gas, causes light to be emitted at specific wavelengths that depend on the particular atom in the gas.
- Fluorescent coating on the bulb's interior is often used to transform the emitted frequencies to a wider range.

- **LED lights** are based on electroluminescence: they use materials that emit photons due to electrical current passing through them.

Light model

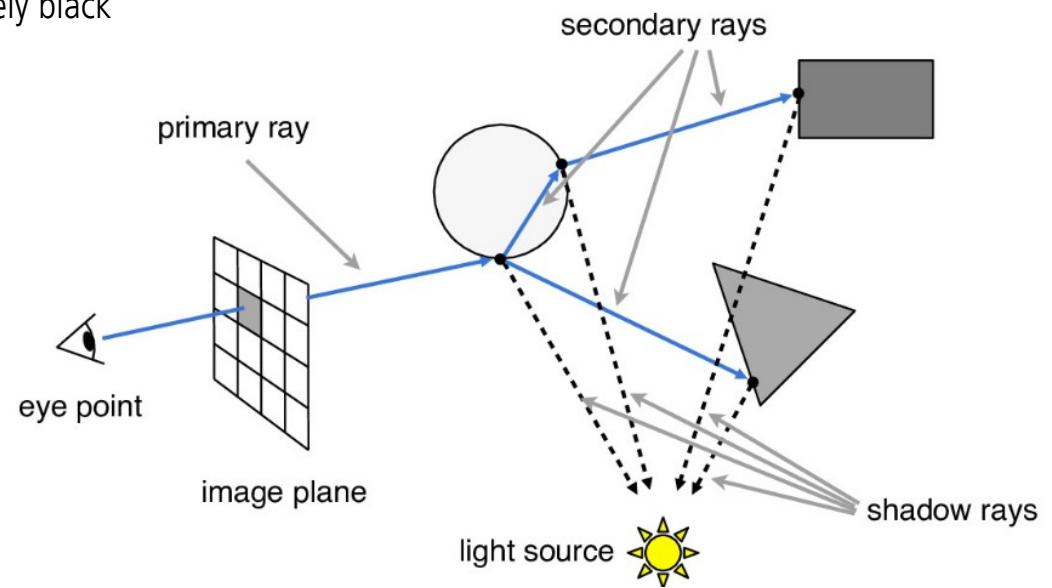
Rendering and light

- Light is emitted from a light source,
- Light propagates through vacuum (ignore participating media)
- Light interacts with objects (e.g., reflection) which depends on object material
- Very small portion of light enters camera, falls on film producing image



Rendering and light

- Rendering: calculating colors of pixels on virtual image plane
 - For each pixel, camera ray is generated and tested for intersection with objects in 3D scene → **visibility solving**
 - Color of intersection point (shading point) is calculated in **shading step**
 - Shading takes in account 3D object shape (i.e., normal), object material, viewing direction and incoming light
 - Without light resulting image would be completely black



Rendering and light

- Rendering equation: color (radiance) of shading point p in (view) direction ω

$$L_o(p, \omega_o) = L_e(p, \omega_o) + \int_{\Omega} f(p, \omega_o, \omega_i) L_i(p, \omega_i) (\omega_i \cdot n) d\omega_i$$

Emission Reflection

BRDF Incoming light attenuation

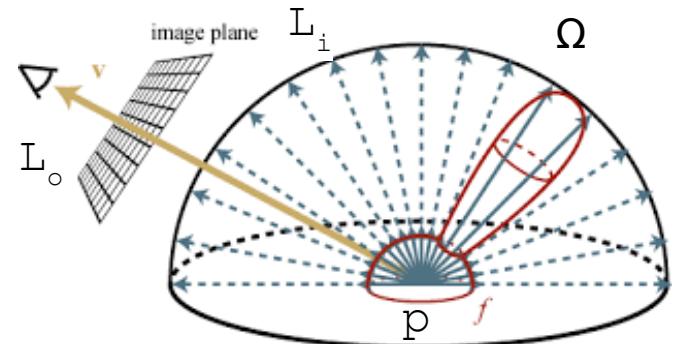
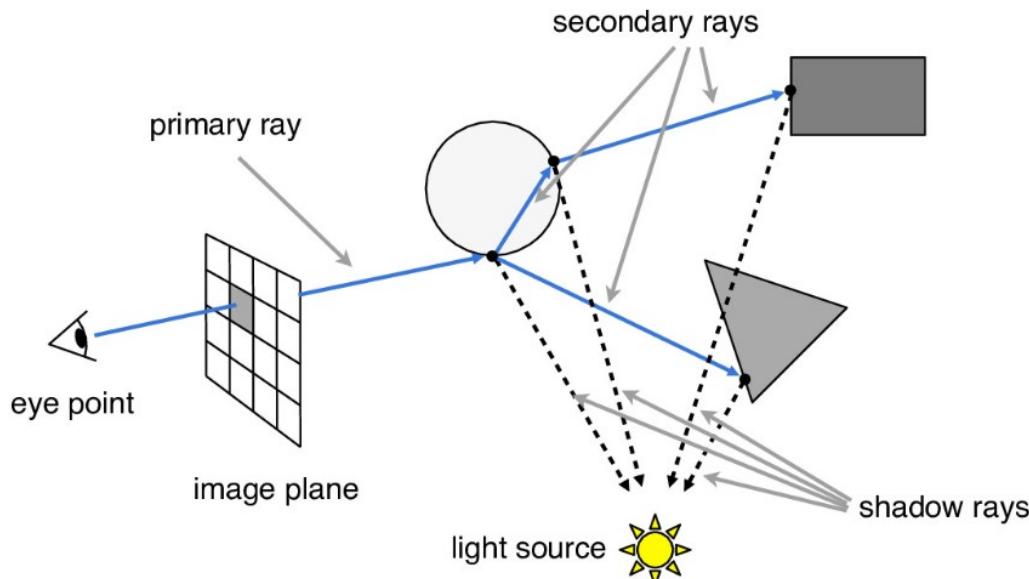
- Foundation of physically-based rendering
- Describes global illumination phenomena – light coming from all directions above object surface
- Recursive equation - not possible to solve analytically: approximation methods are used

Global illumination

- Global illumination: light can fall on surface from any direction:

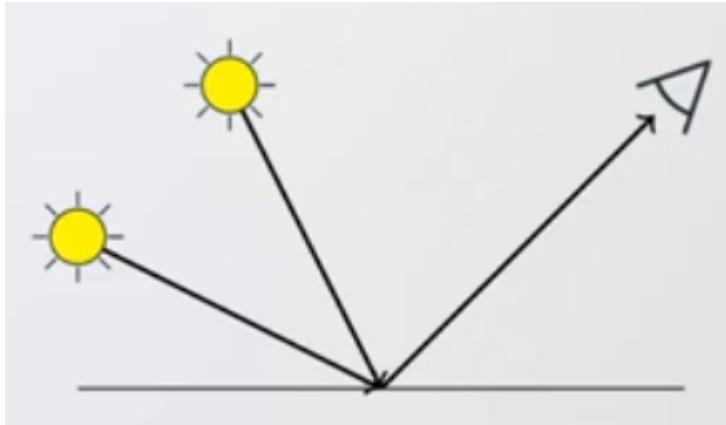
- Light sources
- Surfaces reflecting light

$$L_o(p, \omega_o) = L_e(p, \omega_o) + \int_{\Omega} f(p, \omega_o, \omega_i) L_i(p, \omega_i)(\omega_i \cdot n) d\omega_i$$



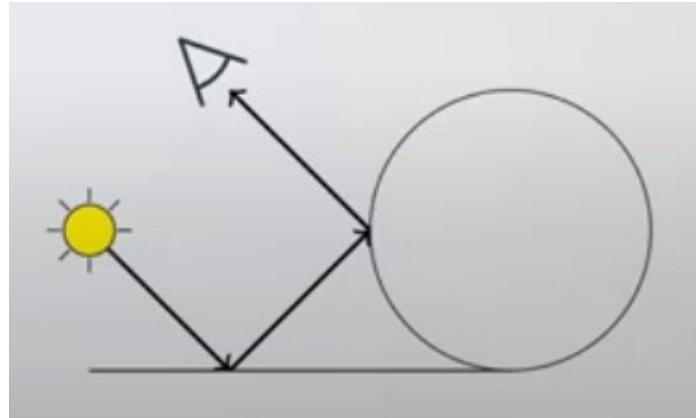
Rendering and light

- Rendering equation is very general equation: different rendering approaches solve different parts of this equation



Direct illumination: takes in account only light from light sources (e.g., non-physical light sources; point lights)

- Often used in rasterization-based rendering
- Problem: additional work needed for shadows, reflections, etc.



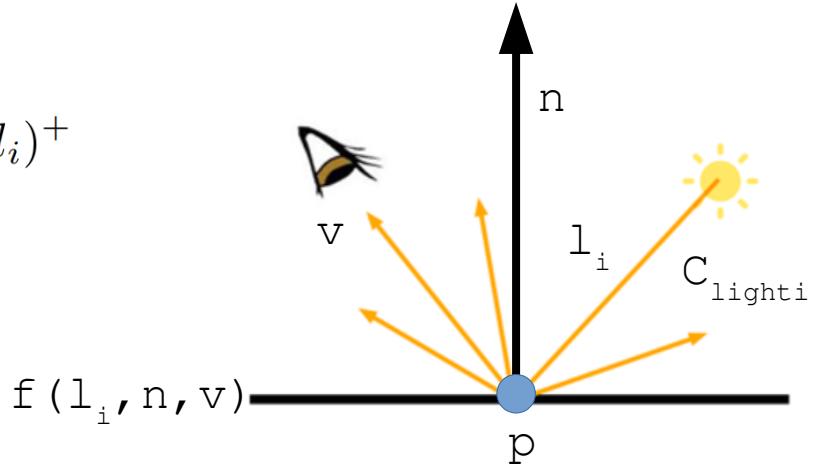
Global illumination: takes in account light from light sources and indirect light: reflections from other surfaces

- Often solved using ray-tracing-based rendering methods
- Multiple reflections, transmissions and shadows can be elegantly incorporated
- More advanced methods, e.g., path tracing, further build on raytracing to solve this problem more correctly

Local (direct) illumination

- Local illumination: shading is considering only light coming from light sources (discrete incoming light angles)
 - l_i is given as color and intensity.
 - Direction of light source is given or calculated from its position

$$c_{shaded}(p, v) = \sum_{i=1}^n f(l_i, n, v) c_{light_i} (n \cdot l_i)^+$$



Local (direct) illumination

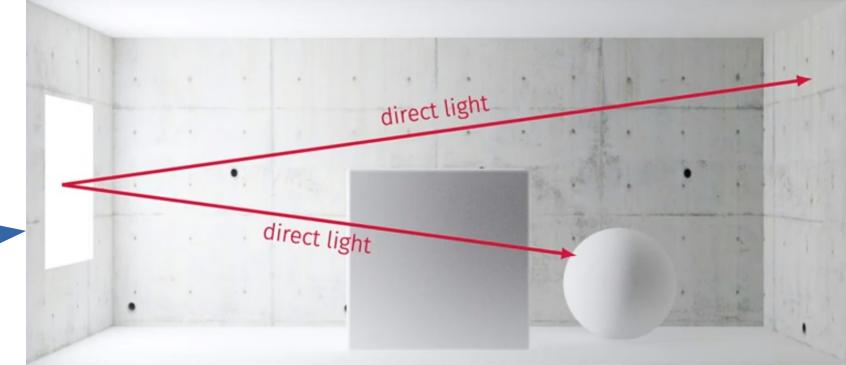
- If multiple light sources are present, they add up linearly
- Different types of light sources differ in light color and intensity and position/direction.



Local and global illumination

Direct illumination

- No shadows



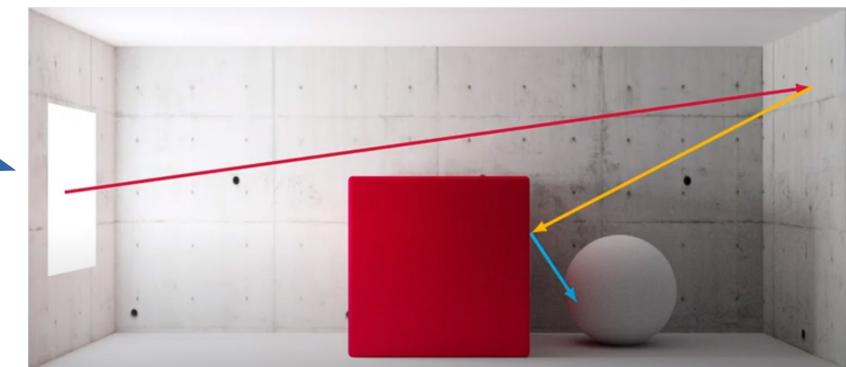
Direct illumination with shadows

- Areas in shadow are completely black



Global illumination

- Shadows
- Multiple reflections
- Color bleeding



Global illumination



Indirect light causes caustics.

Blender and appleseed:

<https://www.blenderdiplom.com/en/tutorials/all-tutorials/649-rendering-caustics-in-blender-with-appleseed.html>



Light refraction can not be solved with local illumination methods:

https://www.youtube.com/watch?v=rJghpp8RVmw&ab_channel=yuichiroyama



Participating media and volumetric objects can not be accurately represented with direct illumination:

<https://forums.jangafx.com/t/blender-cycles-test-rendering-of-the-free-vdb-cloud-pack/1273/3>

Light and rendering (shading)

- Light sources must define
 - Position or direction of light
 - Color and intensity of light
- Additionally, light sources can have:
 - Size
 - Shape



Light models



Representation and modeling of lights **physically**, with shape and size → **geometric area light or physical lights***. Physical lights are computationally expensive.

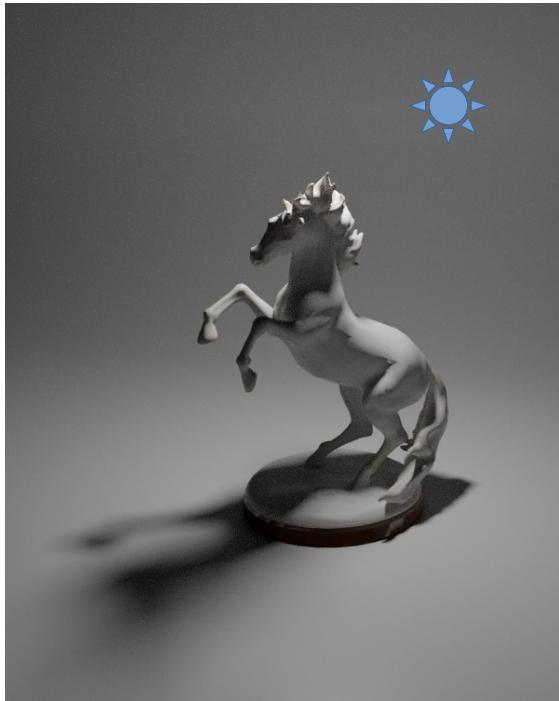
Simplification of lights with no physical size, are called **delta or non-physical lights**.

* When we discussed surface material we mentioned scattering and absorption. Also, materials can be emissive. Therefore, physical lights have 3D shape and material which is emissive. Emissive material is modeled as black body meaning that it absorbs all light falling on it.

Light models parameters

Non-physical lights:

- Color: RGB in $[0, 1]$
- Intensity: float $[0, 1]$
(multiplier)
- Position → Direction



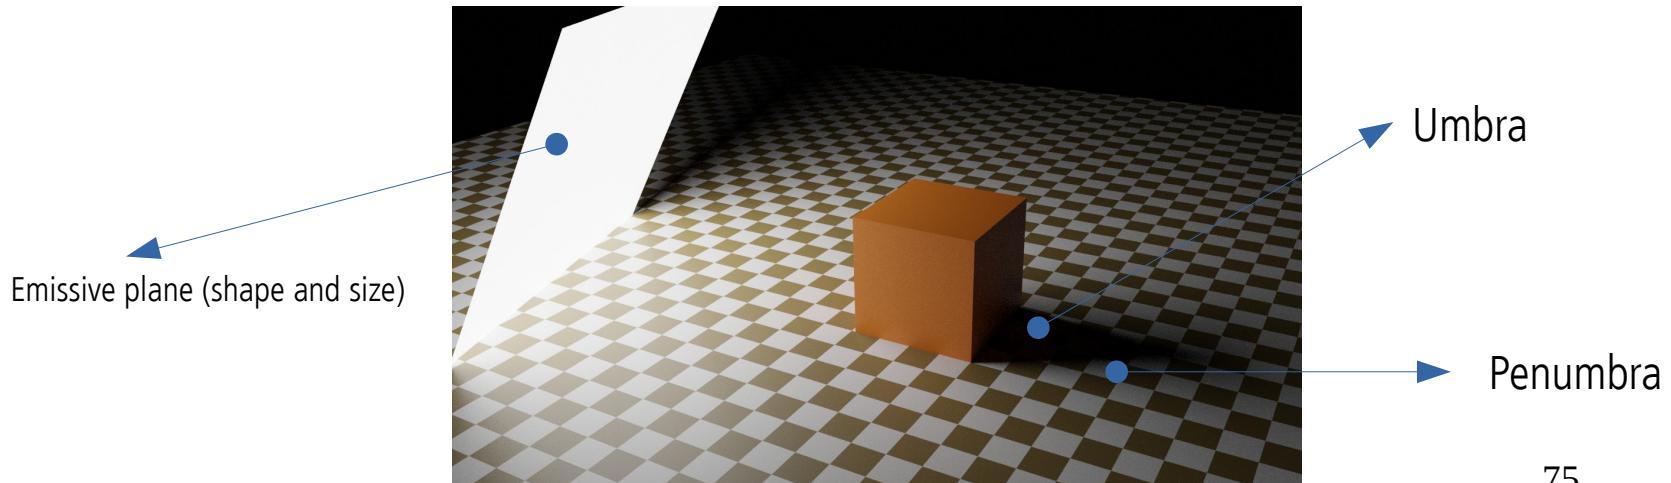
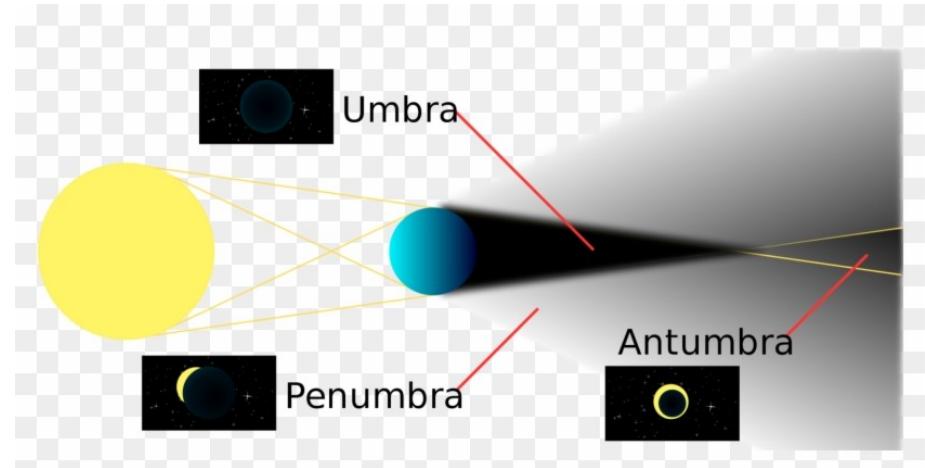
Physical Lights

- Color (RGB)
- Intensity (float)
- Position
- Shape
- Size

Physical lights

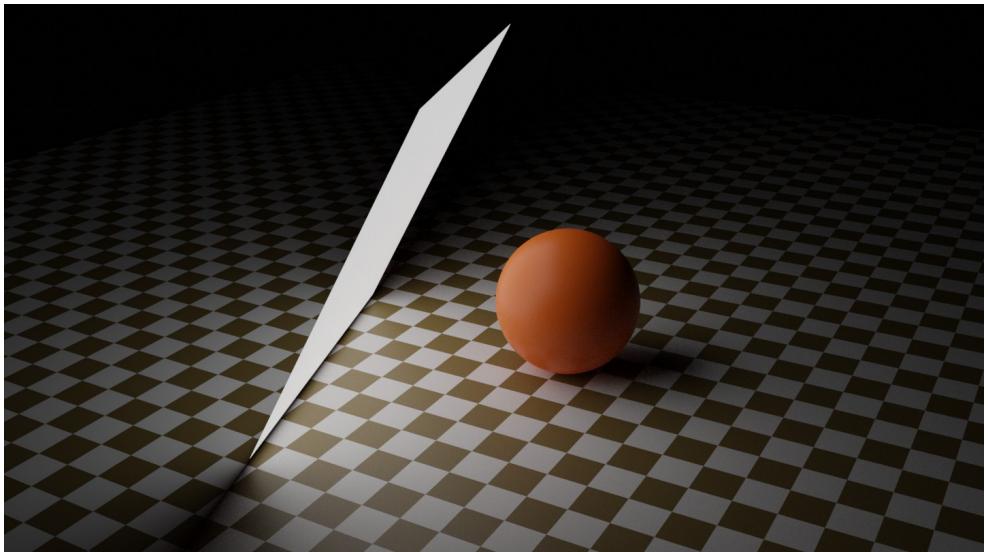
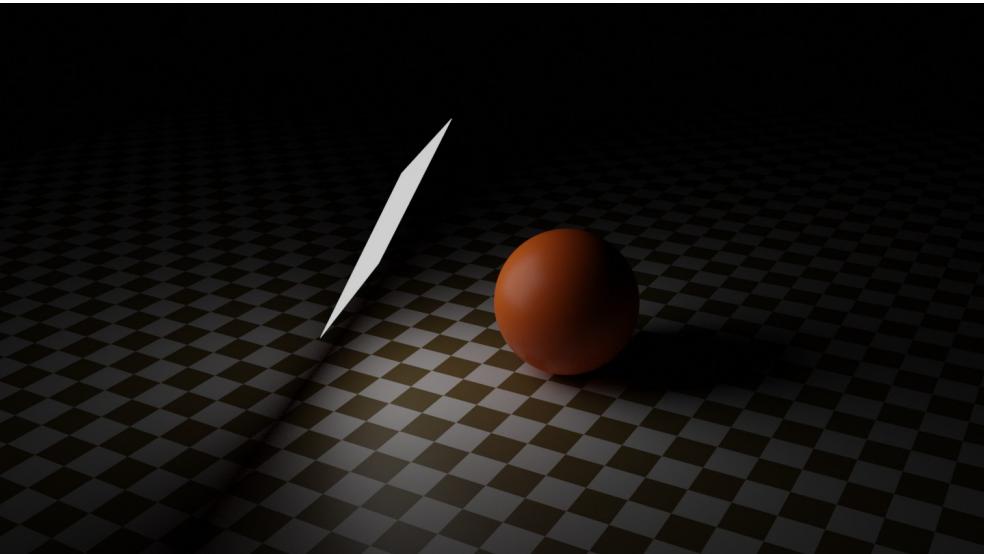
Physical lights

- Shapes that emit light from their surface (emissive material)
 - Color and intensity
 - Position
 - shape and size



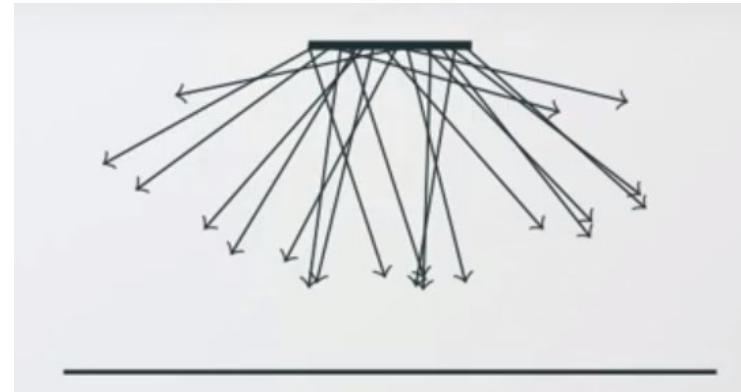
Physical lights

- Shape and size of light matters
 - Amount of brightness
 - Softness and size of shadows



Physical lights

- Intensity and color of emission is given as distribution over emissive surface
- **Diffuse physical light:** uniform spatial and directional intensity and color distribution
 - Light is emitted in direction defined with surface normal
- At each point in the scene, light from physical light source can be incident from many directions → integration over emissive surface is needed*



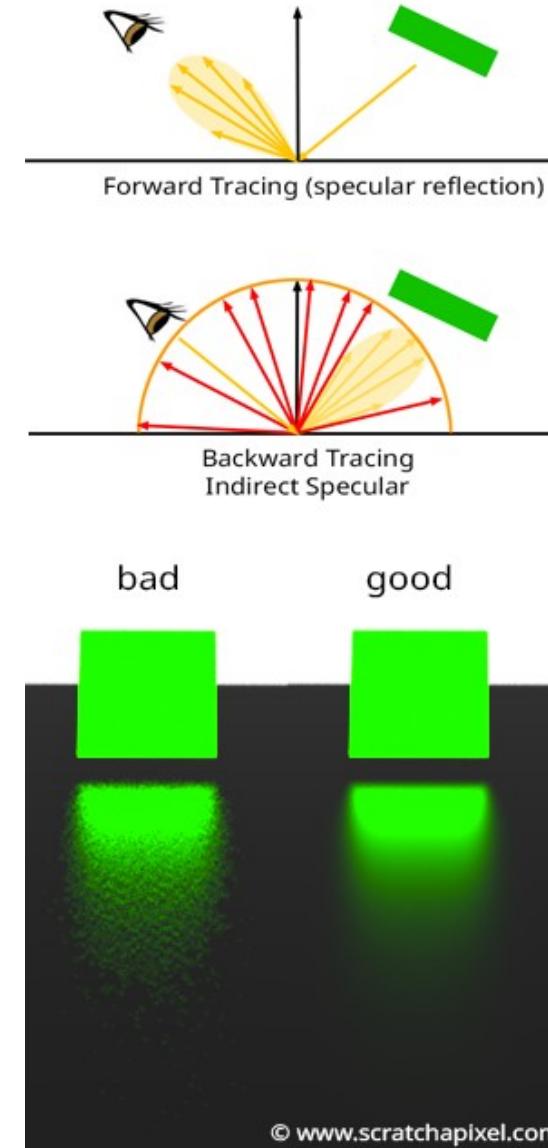
Physical lights

- Computing intensity of physical lights requires computing integrals over emissive surface and often can not be computed in closed form
- Ray-tracing based approaches try to approximate this problem by sampling only smaller number of directions
- Rasterization-based rendering simplifies and approximates effects of physical lights

Physical lights: ray-tracing-based rendering

- Ray-tracing-based rendering simulates physical flow of light, emission from physical lights and their effects on 3D scene (e.g., shadows) can be elegantly solved
- Physical lights have area which has to be sampled
 - Approximation-based methods such as Monte Carlo integration must be used
 - Global illumination methods: path-tracing

$$L_o(p, \omega_o) = L_e(p, \omega_o) + \int_{\Omega} f(p, \omega_o, \omega_i) L_i(p, \omega_i) (\omega_i \cdot n) d\omega_i$$



Light Sampling and Shadow/Direct Specular Noise

AA_Samples = 1. Light samples = 2 (4 samples).

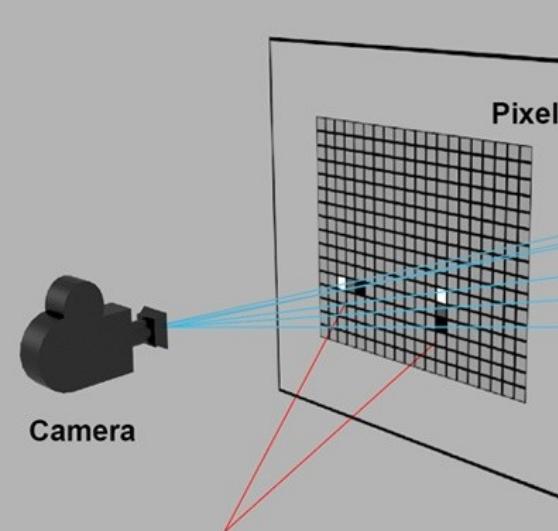
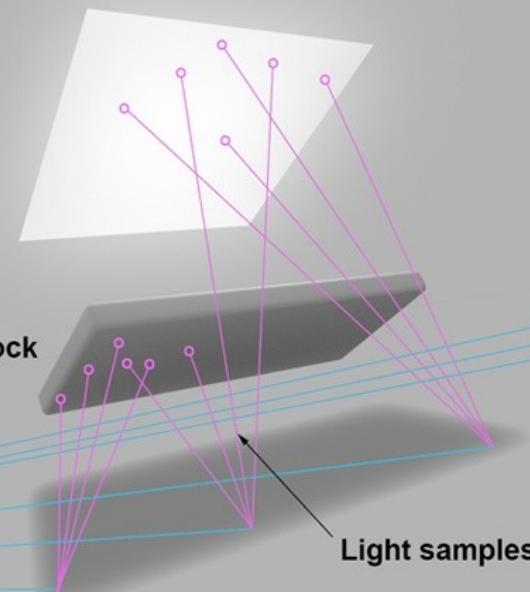
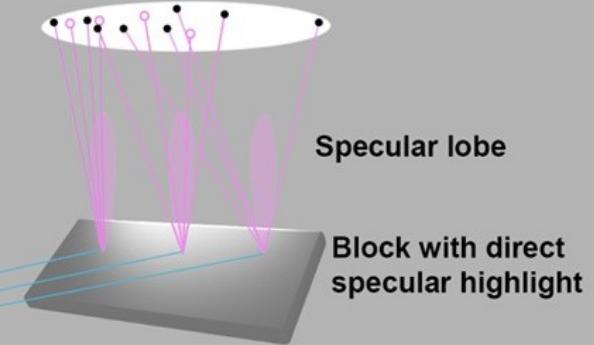


Image noise due to low sampling
(pixel value = average subsamples).

Area light



Area light



Direct Specular Noise

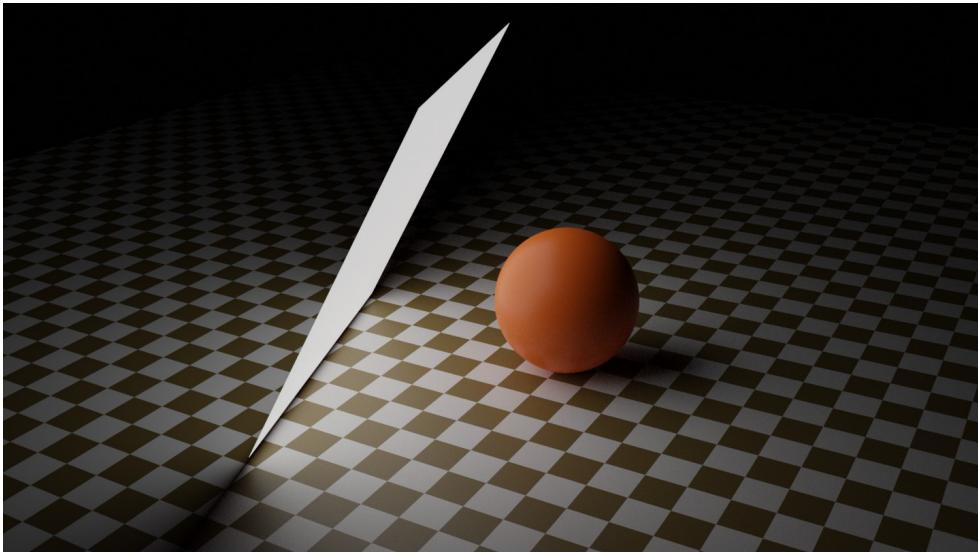
When a light source has a large radius, rays are fired from the surface to sample the surface of the light. However, when the specular roughness of the shader is large and there are not enough samples, noise will appear on the surface. Noise can also appear when the light source is too close to the surface creating high intensity samples.

Shadow Noise

Camera rays intersecting with geometry fire shadow rays that check occlusion. If the radius of a light > 0 , those rays are fired within the angle of the 'light sampling angle' (shaded area with dotted lines); shadows can be noisy as Arnold attempts to resolve the area light with insufficient rays. In this example, the 4 samples fired toward the light randomly hit either the block (resulting in a low value) or the light (resulting in a high value).

Physical lights: rasterization-based rendering

- Rasterization-based rendering is not simulating natural flow of light. Therefore, physical lights and their effects require simplifications and approximations.



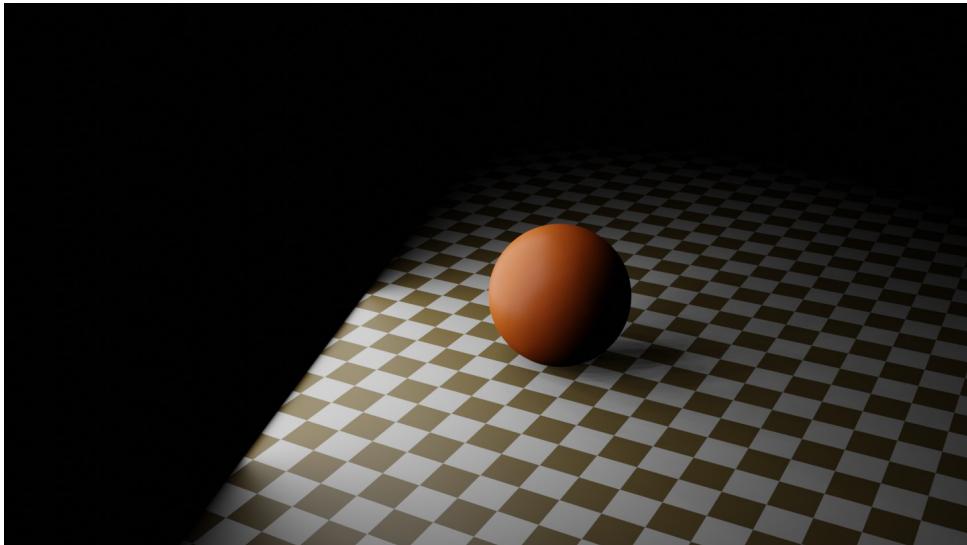
Physical lights in ray-tracing (path-tracing, Blender, Cycles)



Physical lights in rasterization (Blender, EEVEE)

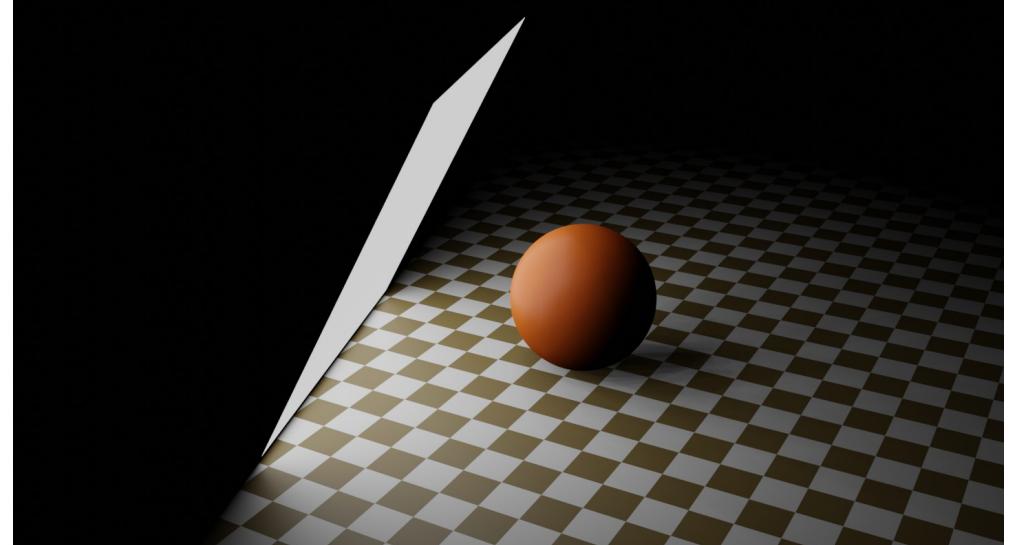
Physical lights: rasterization-based rendering

- Special cases of physical lights - rectangular planes, can be approximated in rasterization-based renderer*



Rectangular plane emission (Blender, EEVEE):

https://docs.blender.org/manual/en/latest/render/lights/light_object.html#area-light



Emissive plane and rectangular plane emission approximation (Blender, EEVEE)

Physical lights: rasterization-based rendering

- Alternative is to pre-compute effects of physical lights using path-tracing and store it in texture which is then used during shading*
 - <https://docs.blender.org/manual/en/latest/render/cycles/baking.html>
 - <https://catlikecoding.com/unity/tutorials/rendering/part-16/>
- Effects of physical lights can be approximated using non-physical lights

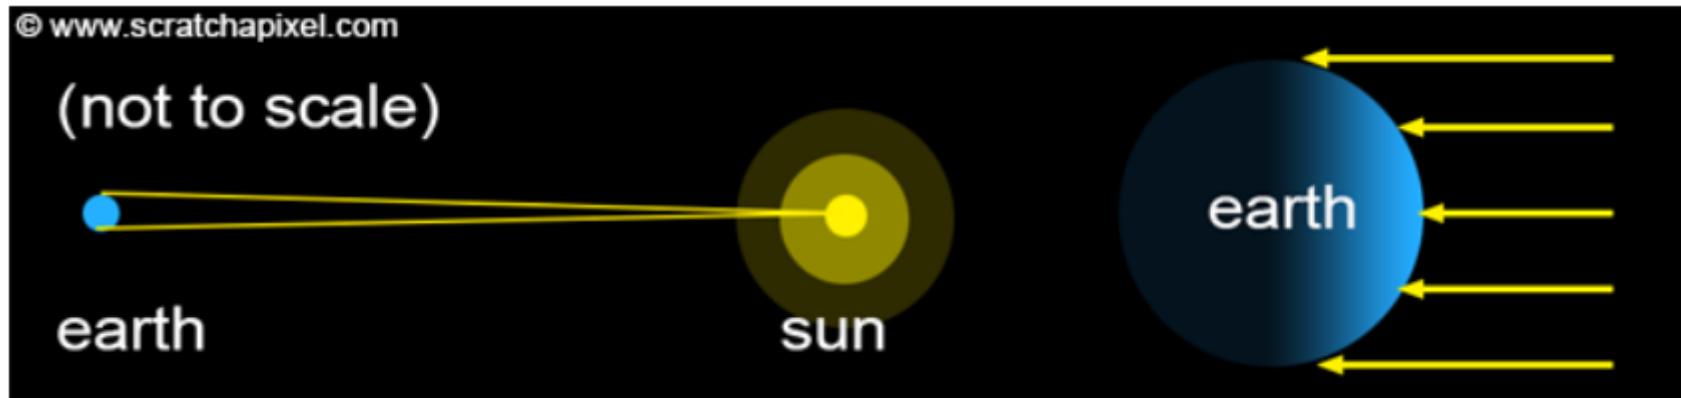
Non-physical lights

Non-Physical lights

- Non-physical lights simulate effects of physical lights: light fall-off with distance, which objects are illuminated, which objects cast shadows, etc.
- Types:
 - **Directional** (distant, parallel) lights
 - **Point** (spherical, omnidirectional) lights
 - Spot lights
 - Other variations

Directional (parallel) lights

- Directional lights are considered so far from the objects in a 3D scene that they can be represented by parallel rays.
 - We only care for the **direction** of those parallel rays

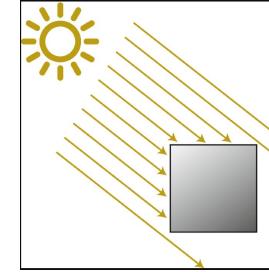


Example of directional light: is Sun light on Earth.

- Sun has spherical shape, but it is so far and Earth is so small compared to it that light rays reaching Earth can be considered parallel (small cone of directions – solid angle).
- For scenes that cover small area of Earth's surface, Sun light can be assumed parallel.

Directional (parallel) lights

- Directional light is simplest model of light source.
 - **Direction (\mathbf{l})** – unit vector (not affected by position!)
 - **Color (\mathbf{c})** – RGB vector in $[0, 1]$
 - **Intensity (i)** – float value in $[0, \infty]$
 - Direction (\mathbf{l}), color (\mathbf{c}) and intensity (i) are constant over 3D scene
- Intensity may be attenuated by shadowing and surface orientation

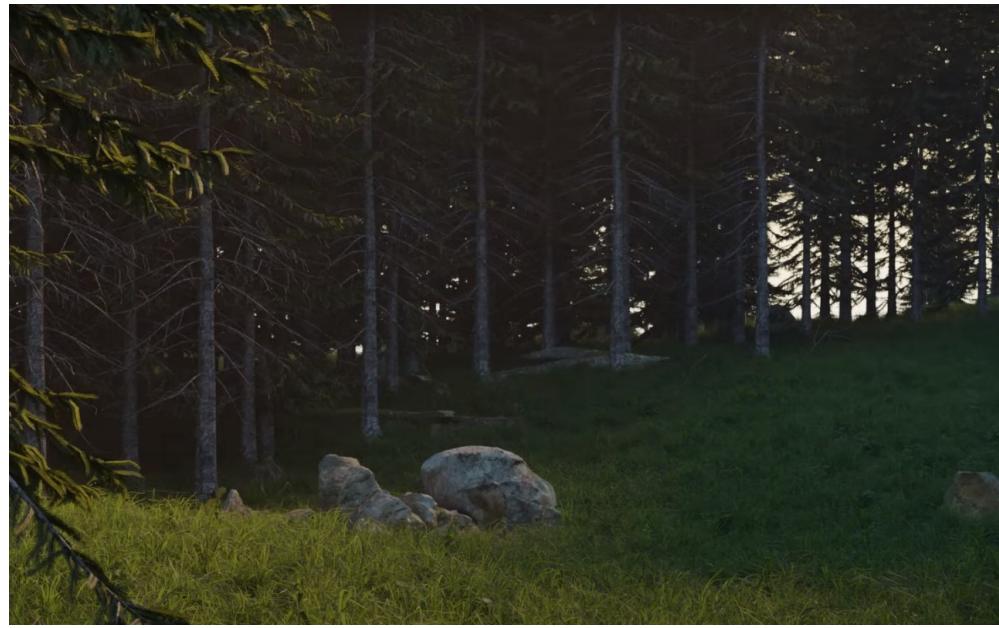


https://link.springer.com/chapter/10.1007/978-1-4842-4457-9_12



Practical tip: modeling with lights

- Orientation of light greatly determines scene appearance



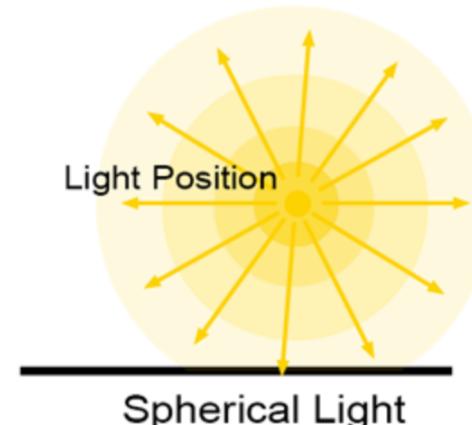
Point lights

Point light

- Infinitesimally small in size → no shape
 - Light is emitted uniformly in all directions (omnidirectional, isotropic)
-
- Parameters
 - **Position** – vector determining position in world space (lights can be also transformed using 4x4 matrices)
 - **Color** – RGB vector in [0,1]
 - **Intensity** – float value in [0,inf]
 - Not affected by scale or rotation.

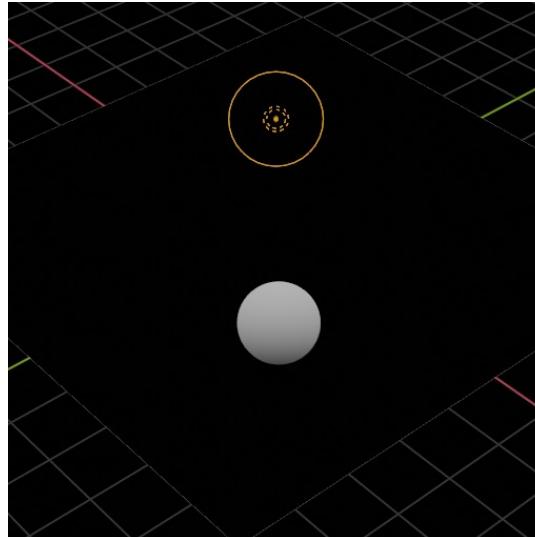
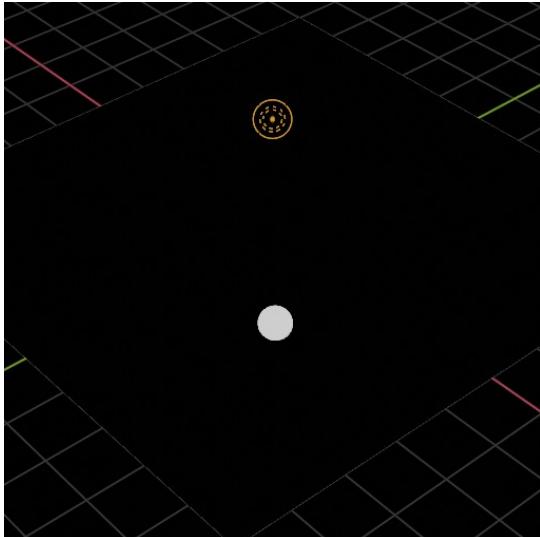


Most common light sources in nature and around us are spherical (e.g., light bulbs) or can be approximated as spherical light sources.



Non-Physical lights: note

- For correct physically-based rendering, non-physical light should be avoided
- Example: size of reflection of object by glossy or mirror-like surface depends on its size and distance to reflective surface. If light source has no shape nor size how reflection should look like? Hack: size parameter which is only used during shading.



Point lights have parameter "size" which determines light reflection size:
https://docs.blender.org/manual/en/latest/render/lights/light_object.html#point-light

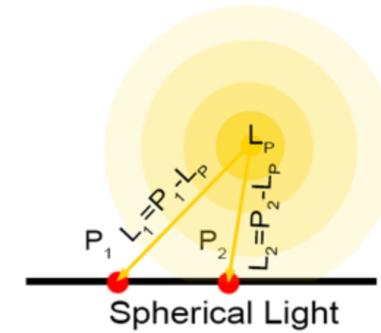
Point lights

- Point light position determines **direction and distance of incoming light** ray for each surface point of objects in the scene

$$\text{PointToLight} = \mathbf{P} - \mathbf{L}_p$$

$$\text{LightDistance} = \text{length}(\text{PointToLight})$$

$$\text{LightDirection} = \text{normalize}(\text{PointToLight})$$

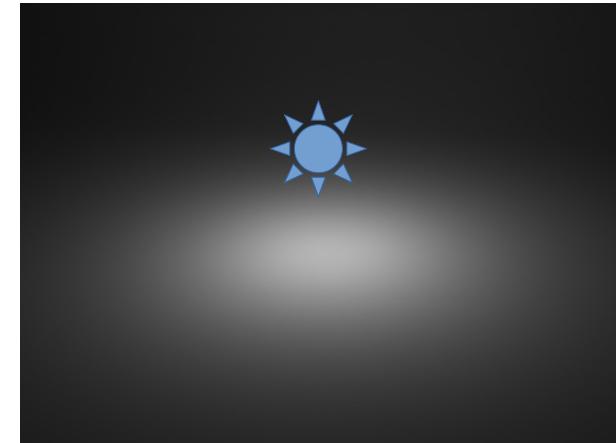
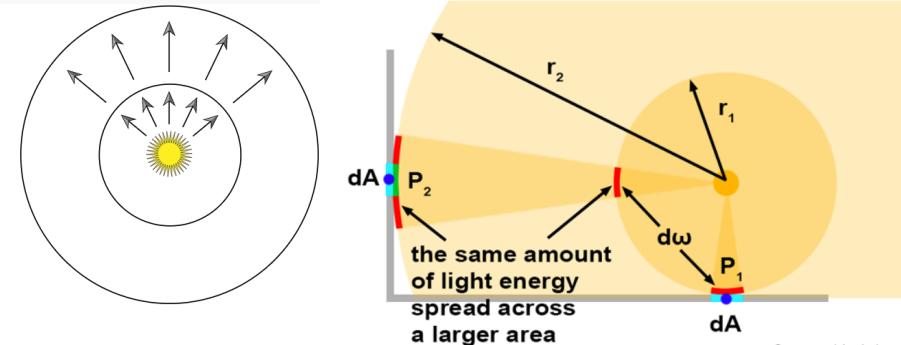


© www.scratchapixel.com



Point lights: direction and light falloff

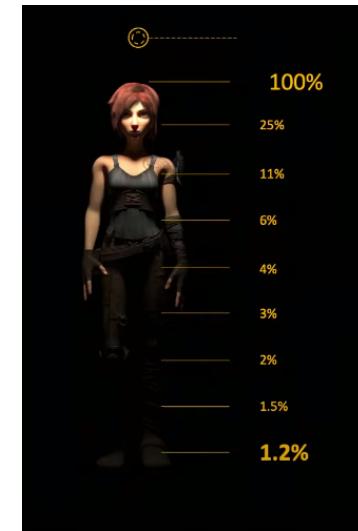
- **Distance** from point light to surface point determines **light falloff**
- Point light sources emit light radially
 - Energy of light source is redistributed over sphere
 - As the sphere keeps expanding in the space, energy becomes spread across much larger area → more distant objects in the scene receive less light → **light falloff**.
 - Distance of point light to surface point determines how much is that point illuminated → **intensity of point light varies as a function of distance.**



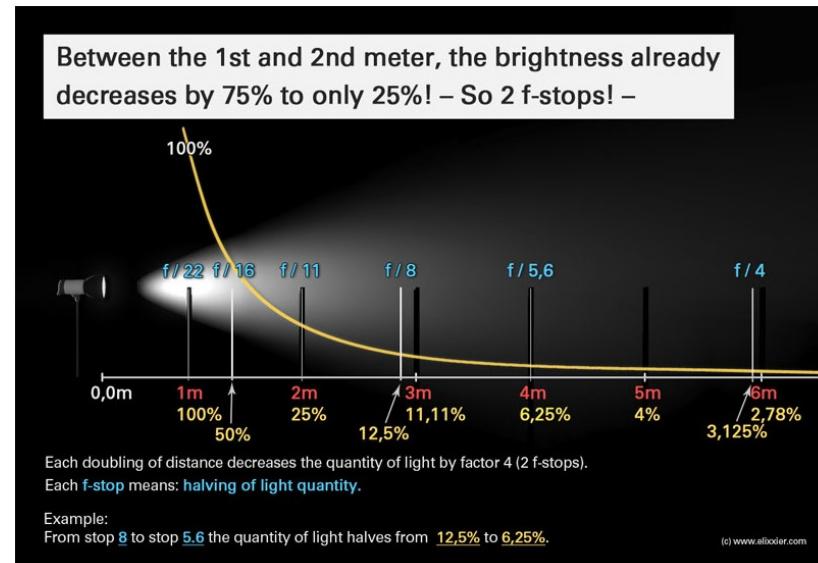
Point light - light falloff

- Point light distributes energy spherically.
- Larger distance → larger sphere → smaller energy
→ attenuated light - **inverse square law falloff**:

$$L_i = \frac{\text{light intensity} * \text{light color}}{4\pi r^2}$$



https://www.youtube.com/watch?v=Z8AAX-ENWvQ&t=2s&ab_channel=Blender



<https://petapixel.com/inverse-square-law-light/>

Point light - light falloff



True Grit (2010)



Toy Story 4 (2021)

https://www.youtube.com/watch?v=Z8AAX-ENWvQ&t=2s&ab_channel=Blender

Point light - light falloff

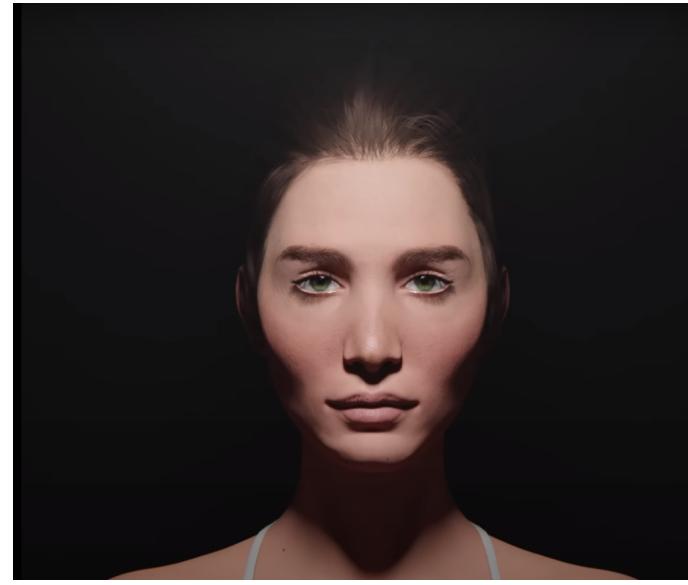
- Problems:
 - **Small distances**: as $r \rightarrow 0$, light intensity \rightarrow infinity. Fix:

$$L_i = \frac{\text{light intensity} * \text{light color}}{4\pi(r + \text{epsilon})^2}$$

- **Large distance**: for efficient rendering, it is desired to light reach 0 at some point. Otherwise, light calculation should be done no matter how far the camera is from the light
- Notes:
 - Light falloff function can be any other function of distance
 - For stylized appearance different functions may be more interesting

Practical tip: modeling with lights

- Orientation of light greatly determines scene appearance



Point lights: directional light falloff

- Next to distance light falloff, real-world illumination intensity also **varies by direction**, generally:

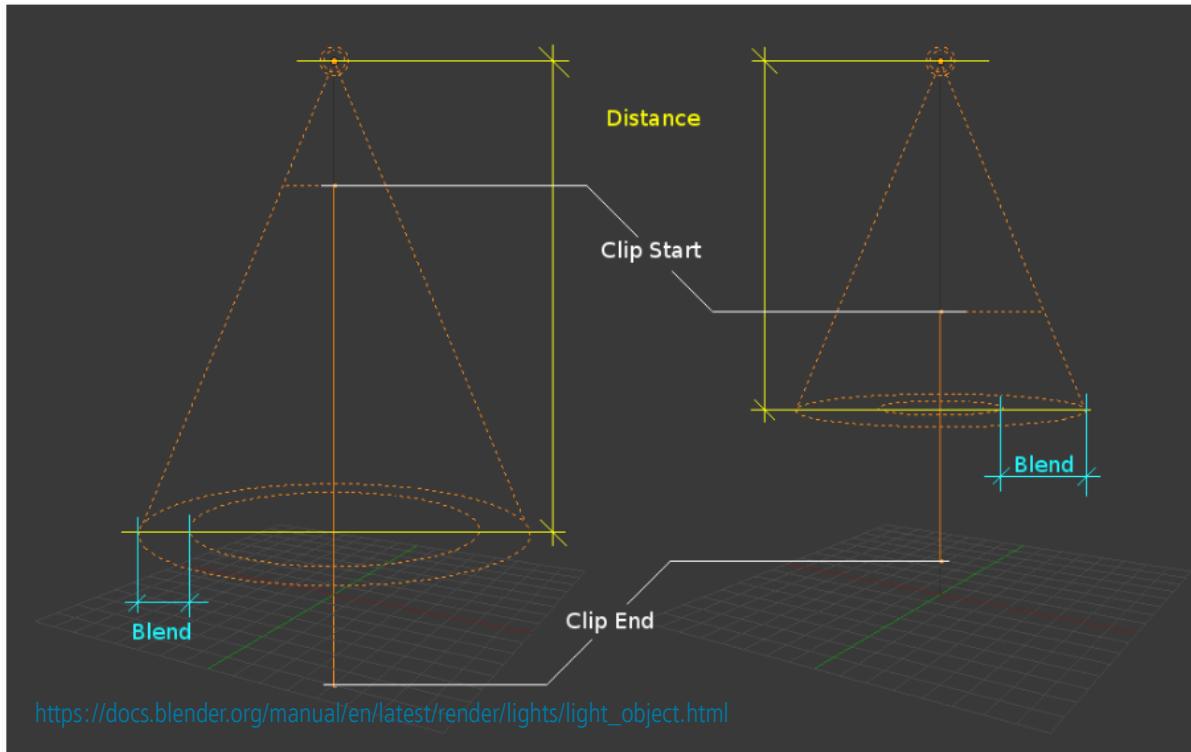
$$L_i = (\text{light intensity} * \text{light color}) f_{distance}(r) f_{direction}(l)$$

- Different choices of distance and directional light falloff functions produce different light sources:
 - **Spotlights** - emit light in a cone of directions from their position

$$L_i = (\text{light intensity} * \text{light color}) f_{distance}(r) \boxed{f_{direction}(l)}$$

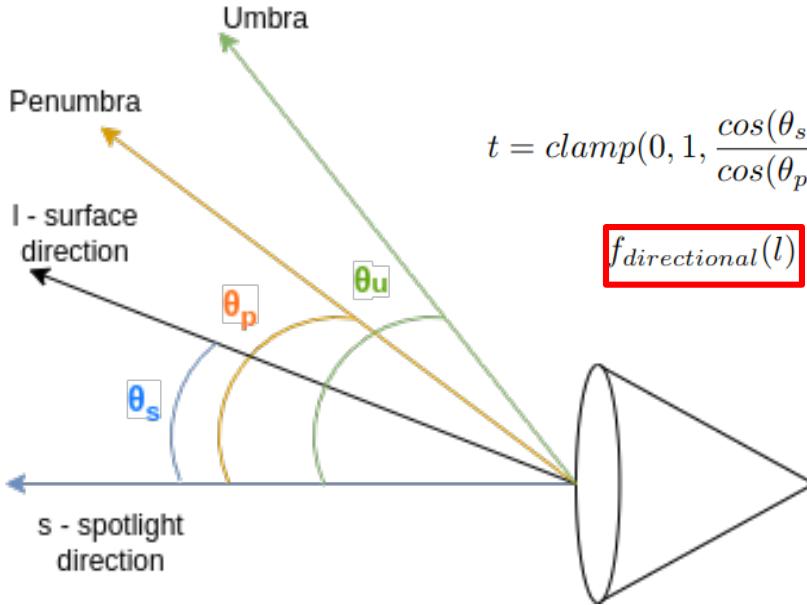
Spotlights

- Projects light in circular cone
- Directional falloff function is rotational symmetric around a spotlight direction

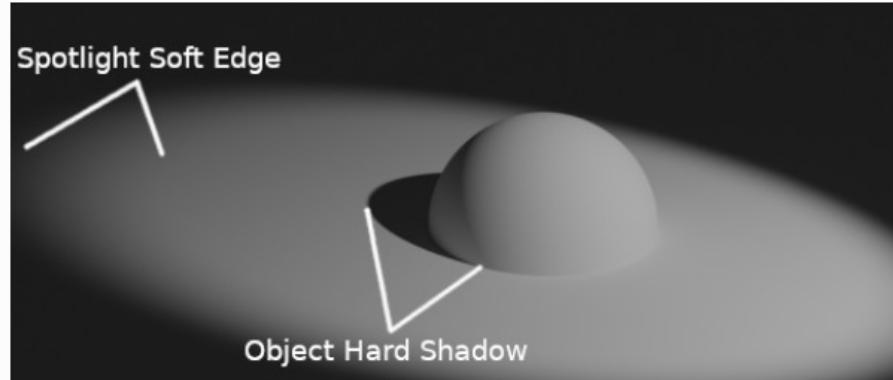


Spotlights

- Parameters:
 - Spotlight direction: s
 - Angle of umbra θ_u
 - Outer cone shell
 - Angle of penumbra θ_p
 - Inner cone where the light is at full intensity
 - Color: RGB vector in $[0, 1]$
 - Intensity: float in $[0, \inf]$



$$L_i = (\text{light intensity} * \text{light color}) f_{distance}(r) f_{directional}(l)$$



Other distance and directional light falloff functions

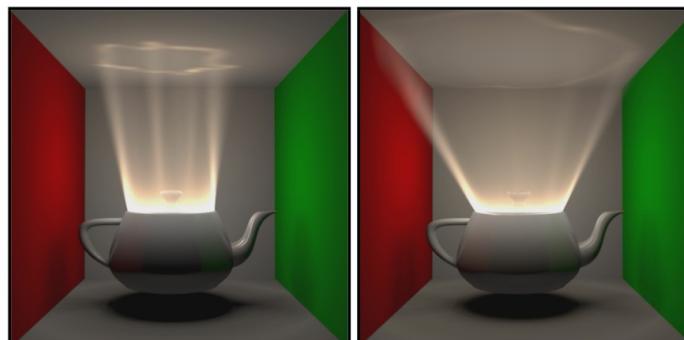
- **Textured lights**

- Use of **texture** for visual richness of light sources: intensity, color and directional/distance variations
- For **cone-type lights**: projective textures (slide projector effect, gobo/cookie lights)
- For **omni-directional-type lights**: texture for distance falloff – attenuation maps (e.g., light beams)

$$L_i = (\text{light intensity} * \text{light color}) f_{distance}(r) f_{direction}(l)$$



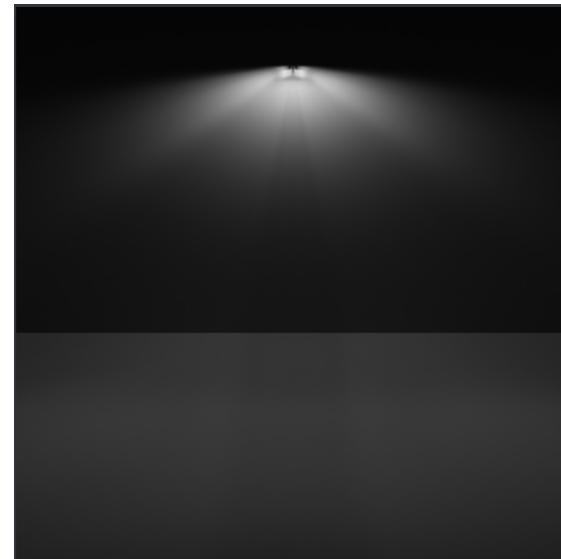
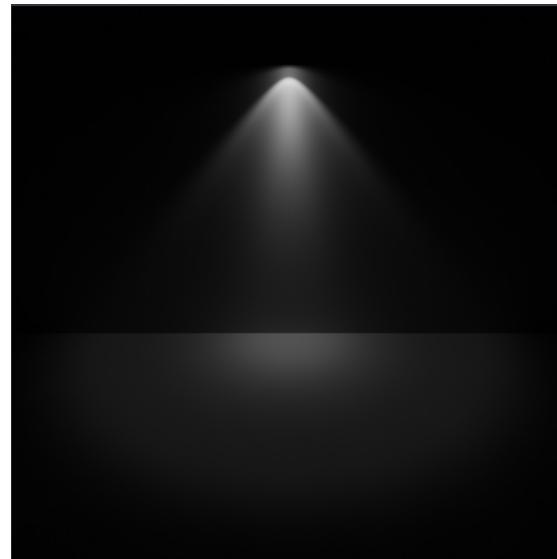
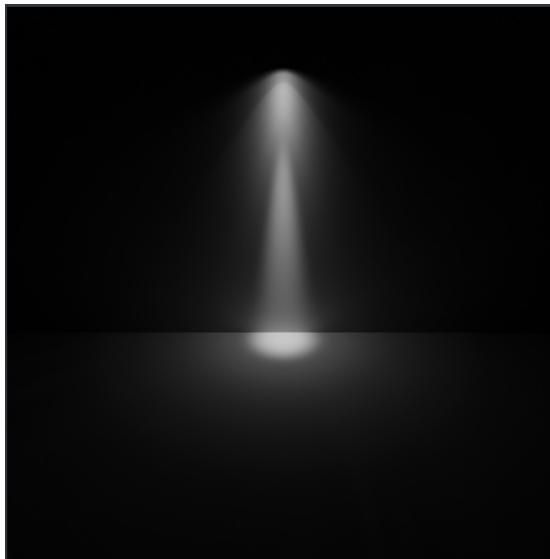
https://developer.valvesoftware.com/wiki/Env_projectedtext ure



https://www.researchgate.net/publication/220183756_A_Programmable_System_for_Artistic_Volumetric_Lighting

Other distance and directional light falloff functions

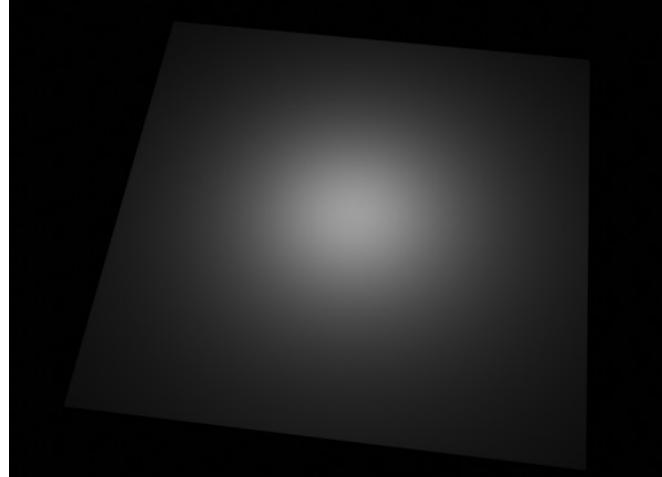
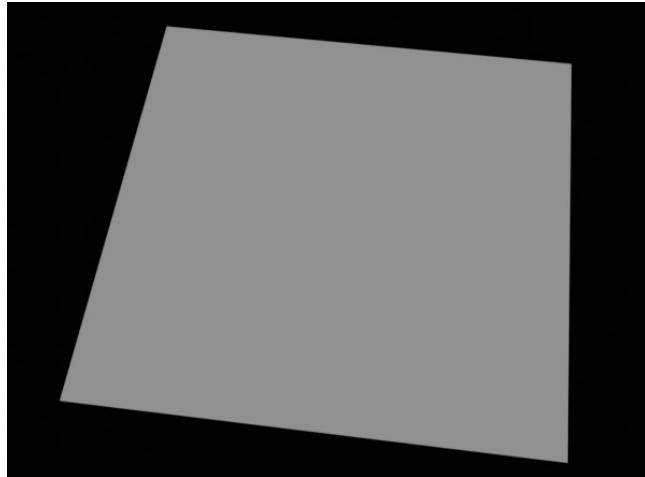
- Tabulated patterns
 - Standard: Illuminating Engineering Society (IES profiles)
 - Measured from the real world using: Goniophotometer



IES profiles: <https://ieslibrary.com/en/home>

Non-physical lights

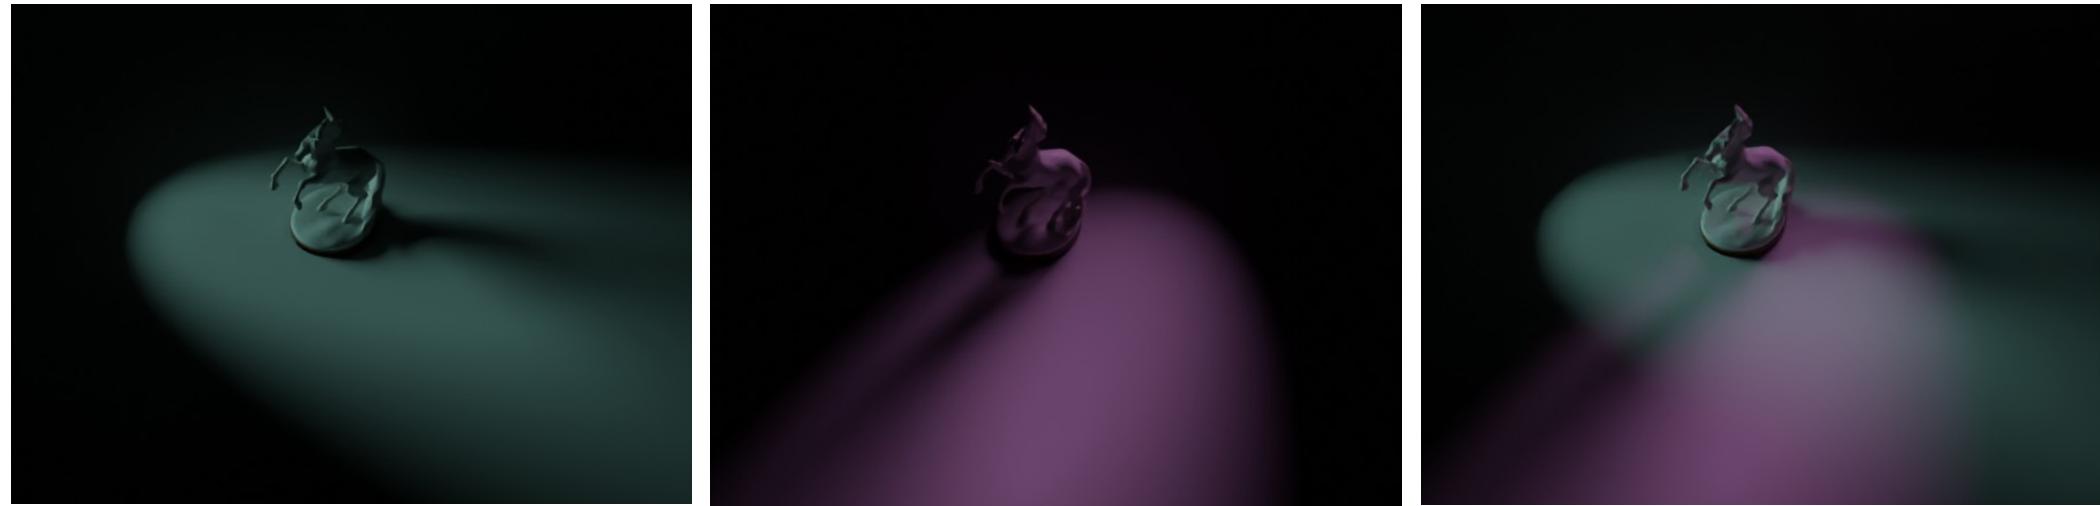
- Directional, point, spot light



Multiple lights

- Contribution of light adds up linearly
 - Basis of photo-realistic rendering where scene contains multiple lights
 - Required for compositing of rendered images: adding multiple rendered images where each has contribution from one light

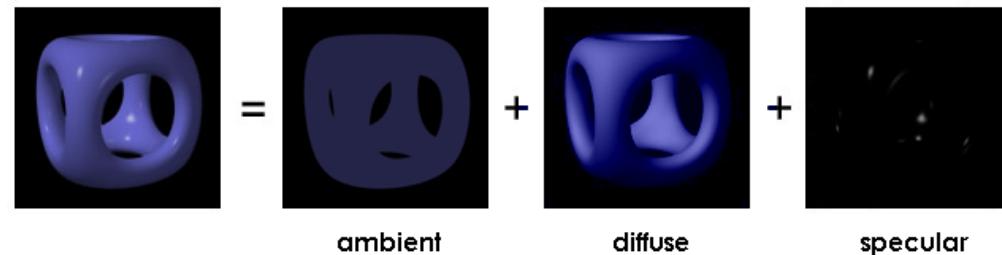
$$L_o(p, v) = \sum_{i=1}^n f(l_i, n, v) c_{light_i}(n \cdot l_i) \quad L_o(p, v) = \int_{l \in \omega} f(l, n, v) L_i(p, l)(n \cdot l) dl$$



Environment illumination

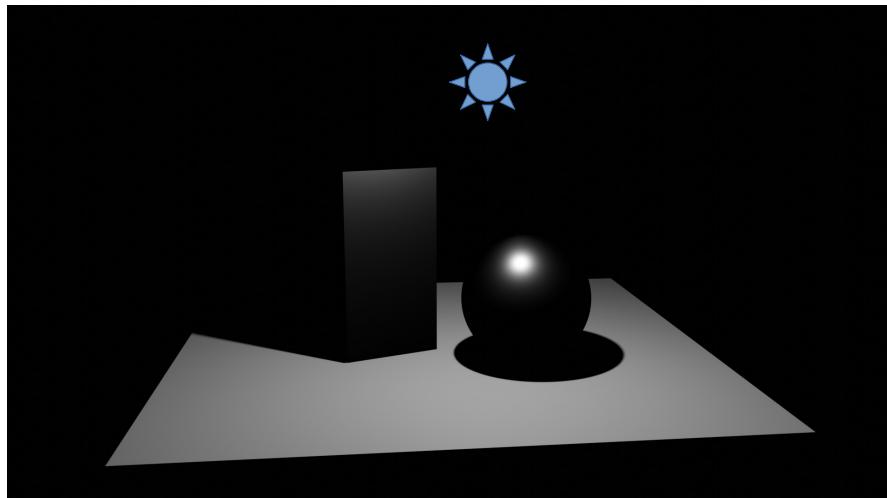
Environment illumination

- Until now, we have considered point and directional lights which limits surfaces to receive light from small number of discrete directions.
- In reality, scenes receive light from all directions otherwise surfaces would appear too dark or completely black if facing away from light sources
 - Example: ambient term in Phong model; constant color added to each point of surface
- To approximate light coming for different directions and large sources of light such as sky and clouds, environment lighting is used.
 - Such emitters solve the problem of objects being too dark

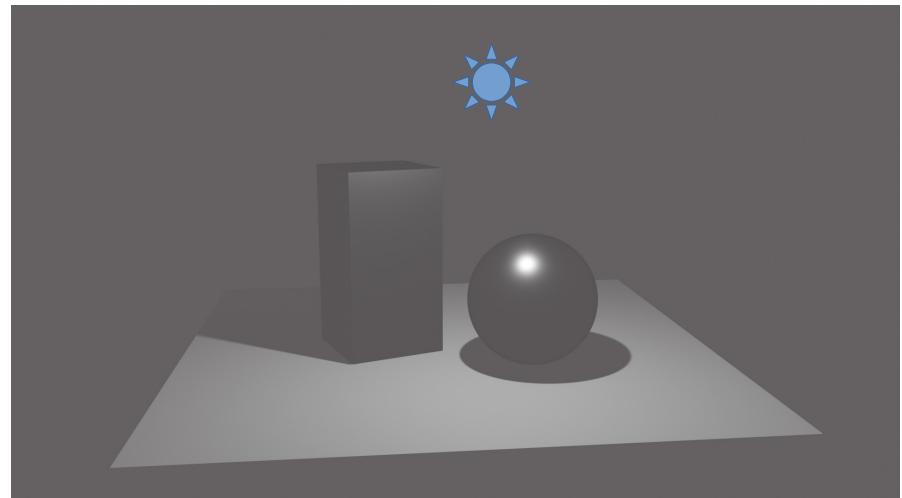


Environment illumination

- Light coming from point or directional light sources can be described with discrete incoming light directions to surface point
- Environment illumination is described with light coming from all directions to surface point



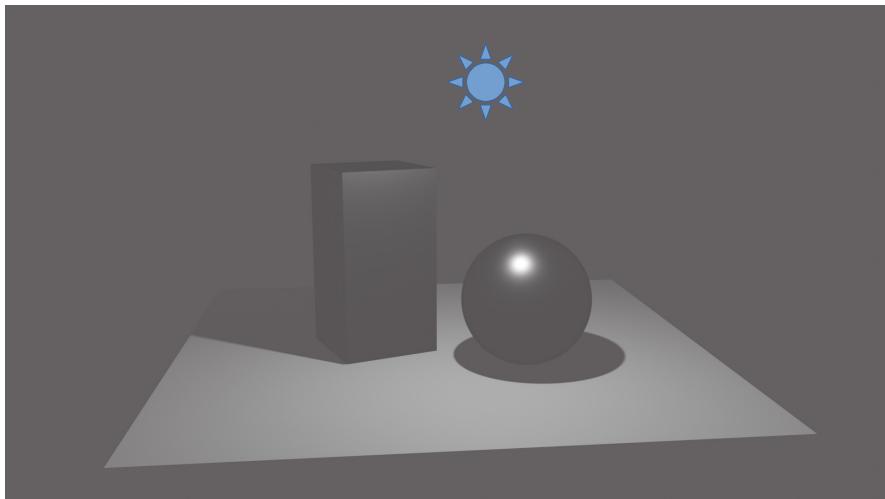
Only point light



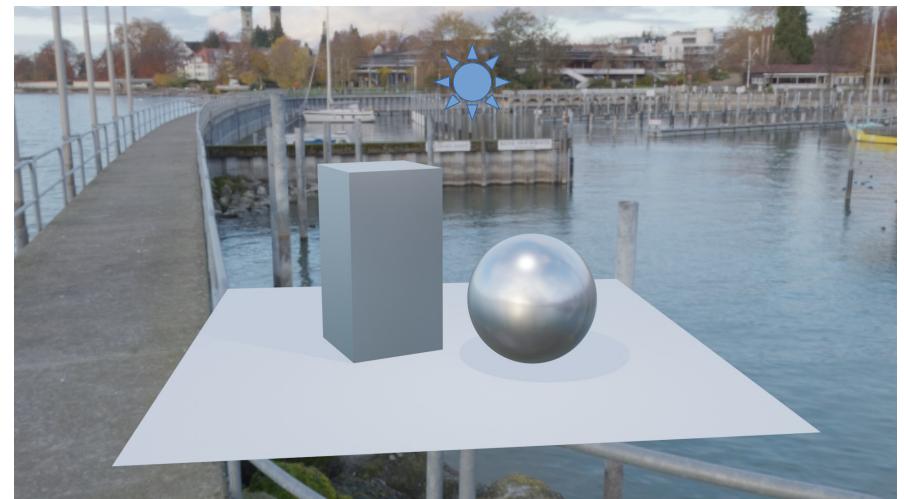
Point light and constant environment

Environment illumination

- Environment illumination is described with light coming from all directions to surface point
- Variation in intensity from different directions can be obtained using texture



Point light and constant environment



Point light and textured environment

Environment illumination

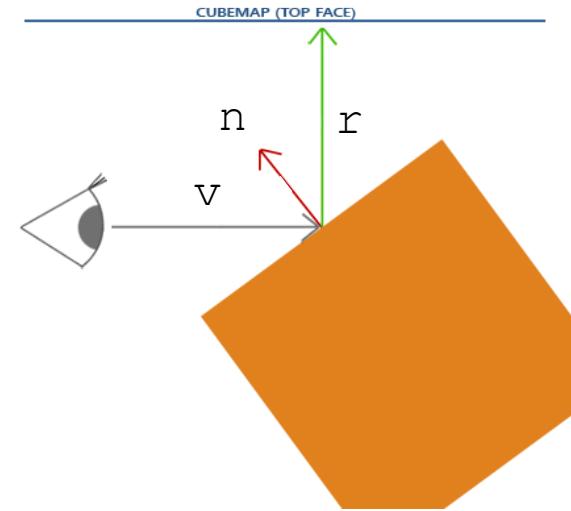
- Note that environment illumination is not global illumination
 - Knowledge of other surfaces which might reflect light is not used
 - Only more directions of light from light emitter are used
- However, this approximation can be used for simulating global illumination effects, that is, indirect light.

Environment illumination

- Variation in intensity depends only on incoming direction and not position
 - Assumption: environment light is infinitely far away.
- Such illumination can be represented by spherical functions
 - Defined over surface of unit sphere or space of directions
- Spherical function representations:
 - **Tabulated forms:** discrete points/directions over sphere containing intensity values which are interpolated
 - **Spherical basis:** spherical radial basis functions, spherical harmonics, etc.
 - **Environment mapping:** storing spherical function into texture, e.g., image texture
 - Most widely used, more memory costly but fast compared to other representations

Environment mapping

- Environment illumination is stored into image texture → **environment map**
 - Unlike surface textures, environment textures have color intensities larger than 1 → high dynamic range
- Global spherical function assumption: **light intensity only depends on a direction**
- Basic case of environment mapping: **reflection mapping**
 - For each reflective material, compute normal (n) at the location on the surface of the object
 - Compute reflection vector (r) from view vector (v) and normal vector (n)
$$r = 2(n \cdot v)n - v$$
 - Use reflected view vector (r) to compute environment map texture location
 - Read value from environment map texture and use it as light intensity

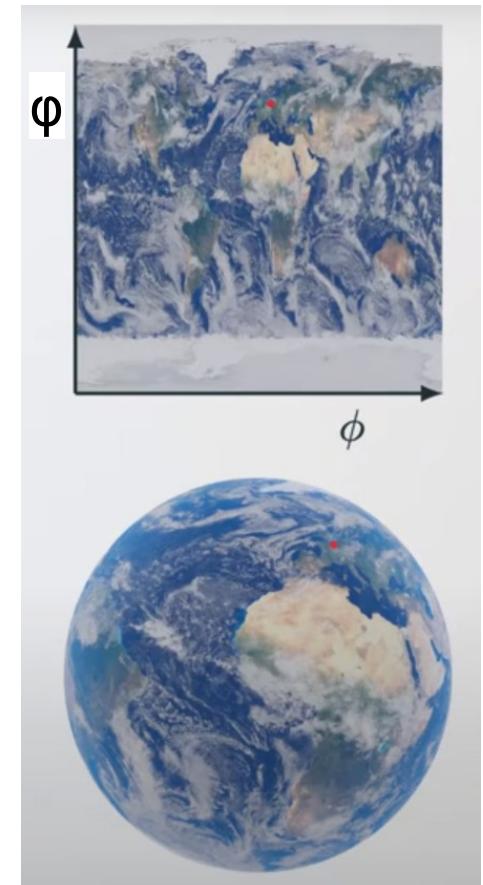


Reflection mapping

- Variety of projector functions mapping from reflected view vector (r) to environment texture location:
 - Latitude-longitude mapping
 - Sphere mapping
 - Cube mapping

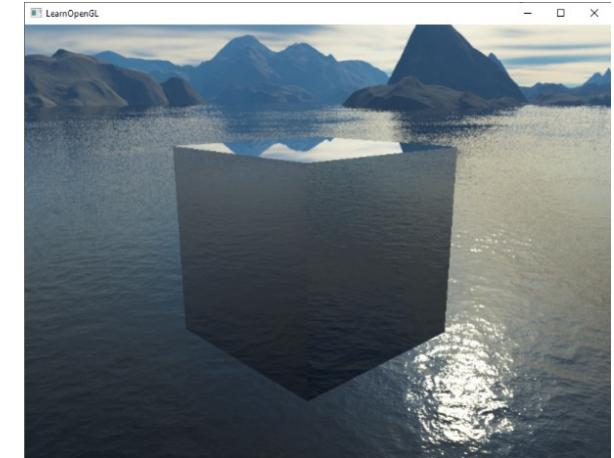
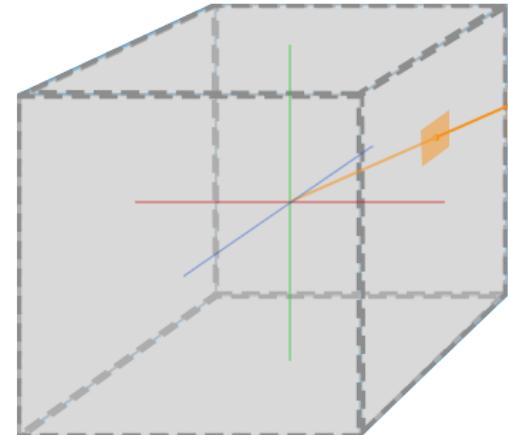
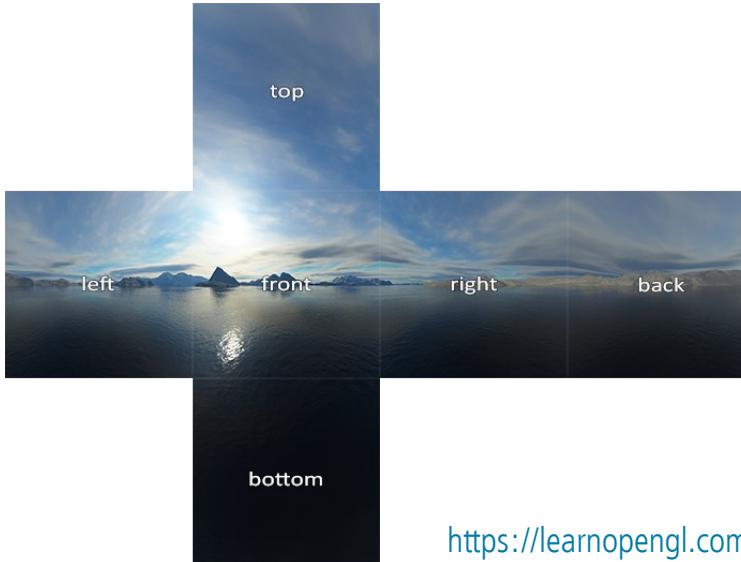
Latitude-longitude mapping

- Blinn and Newell method, first environment mapping method, 1976
 - Environment map idea: imagine that globe is viewed from inside
- Reflected view vector r (r_x , r_y , r_z) is converted to spherical coordinates:
 - Longitude Φ in $[0, 2\pi]$ $\rightarrow \Phi = \text{atan2}(r_y, r_x)$
 - Latitude ϕ in $[0, \pi]$ $\rightarrow \phi = \arccos(r_z)$
- Spherical coordinates are used to access the environment map

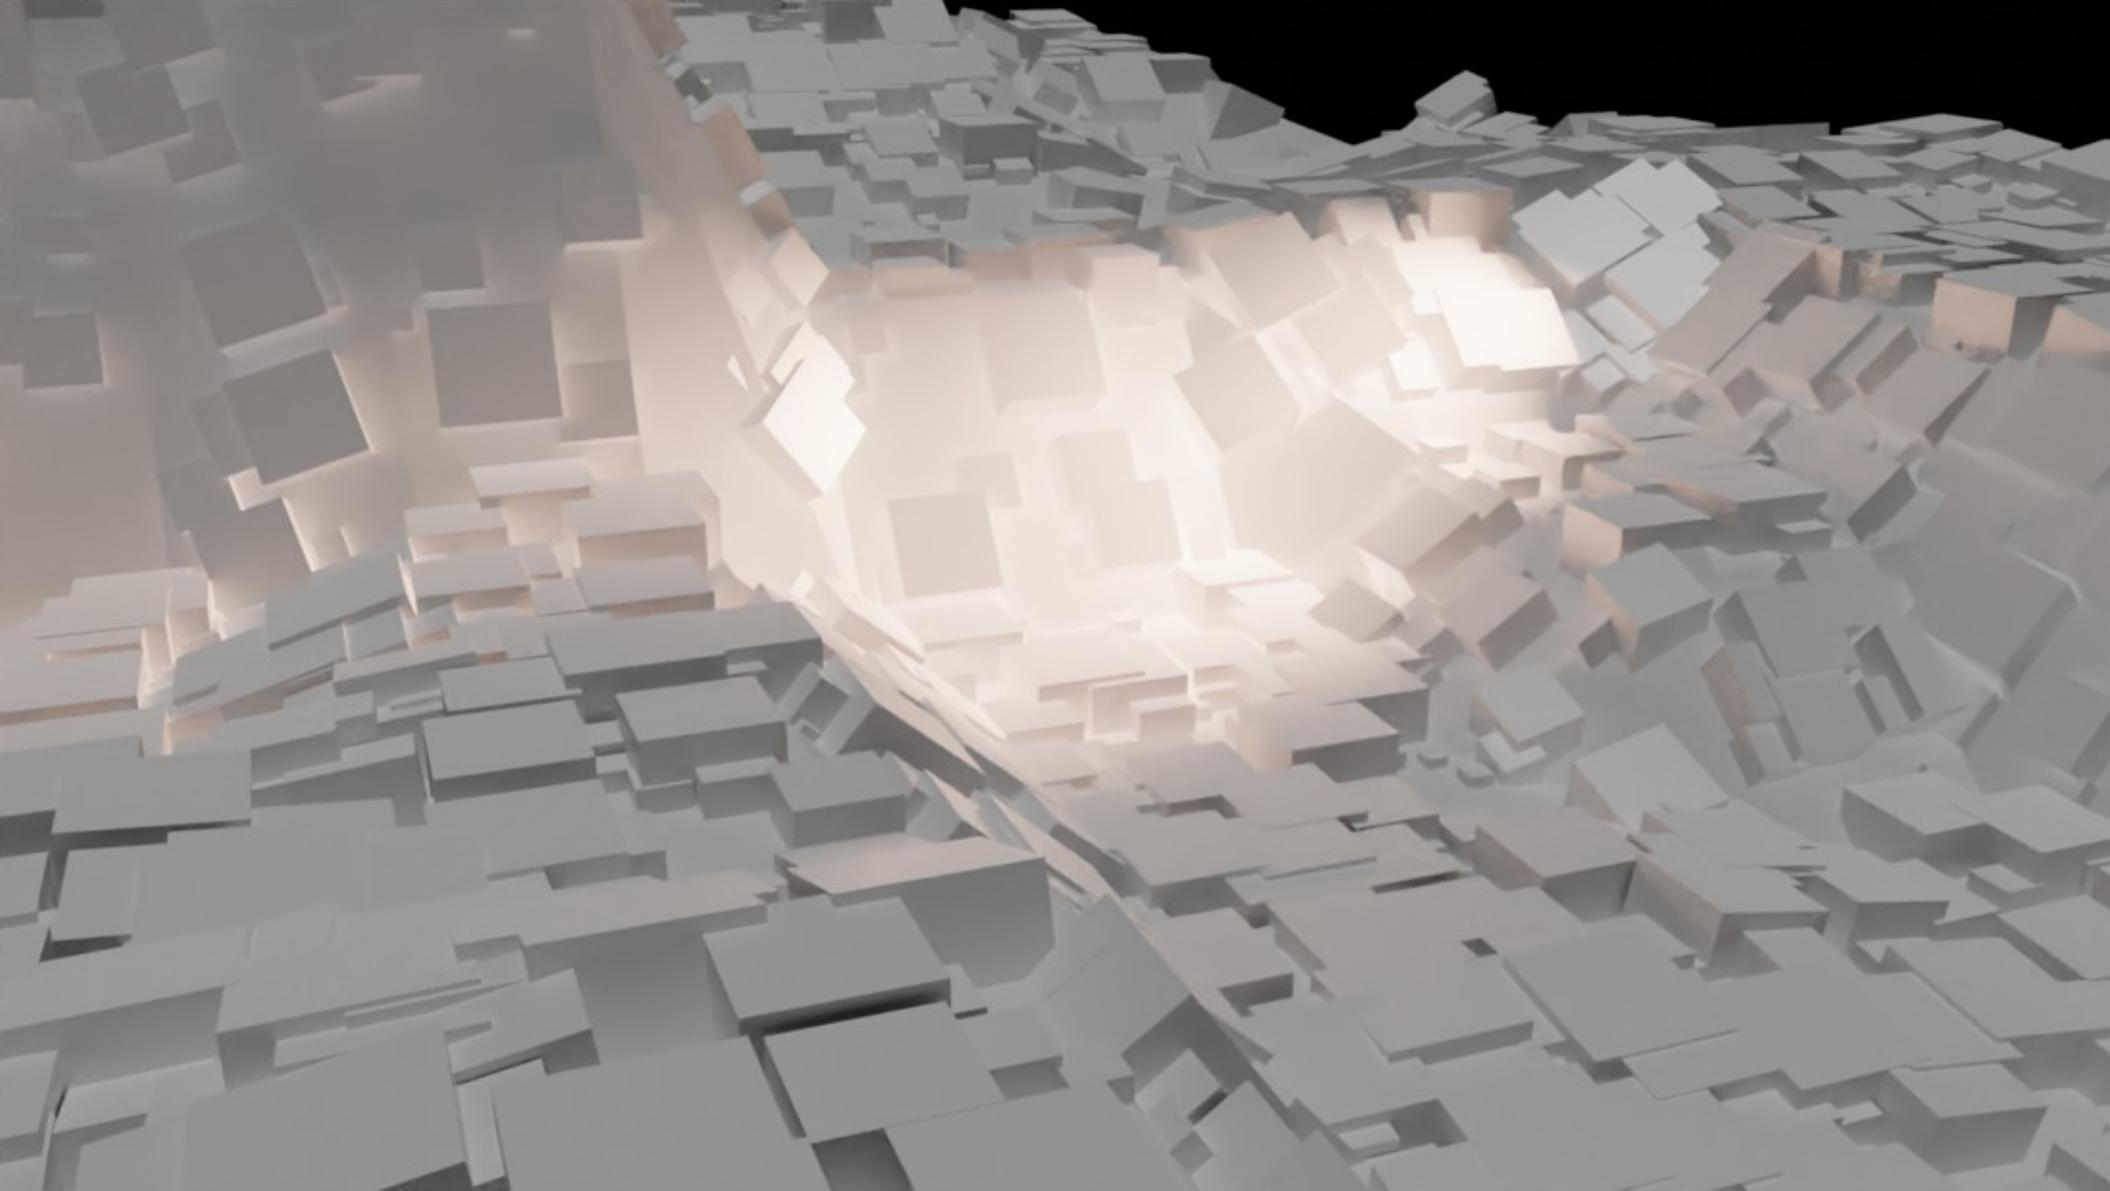


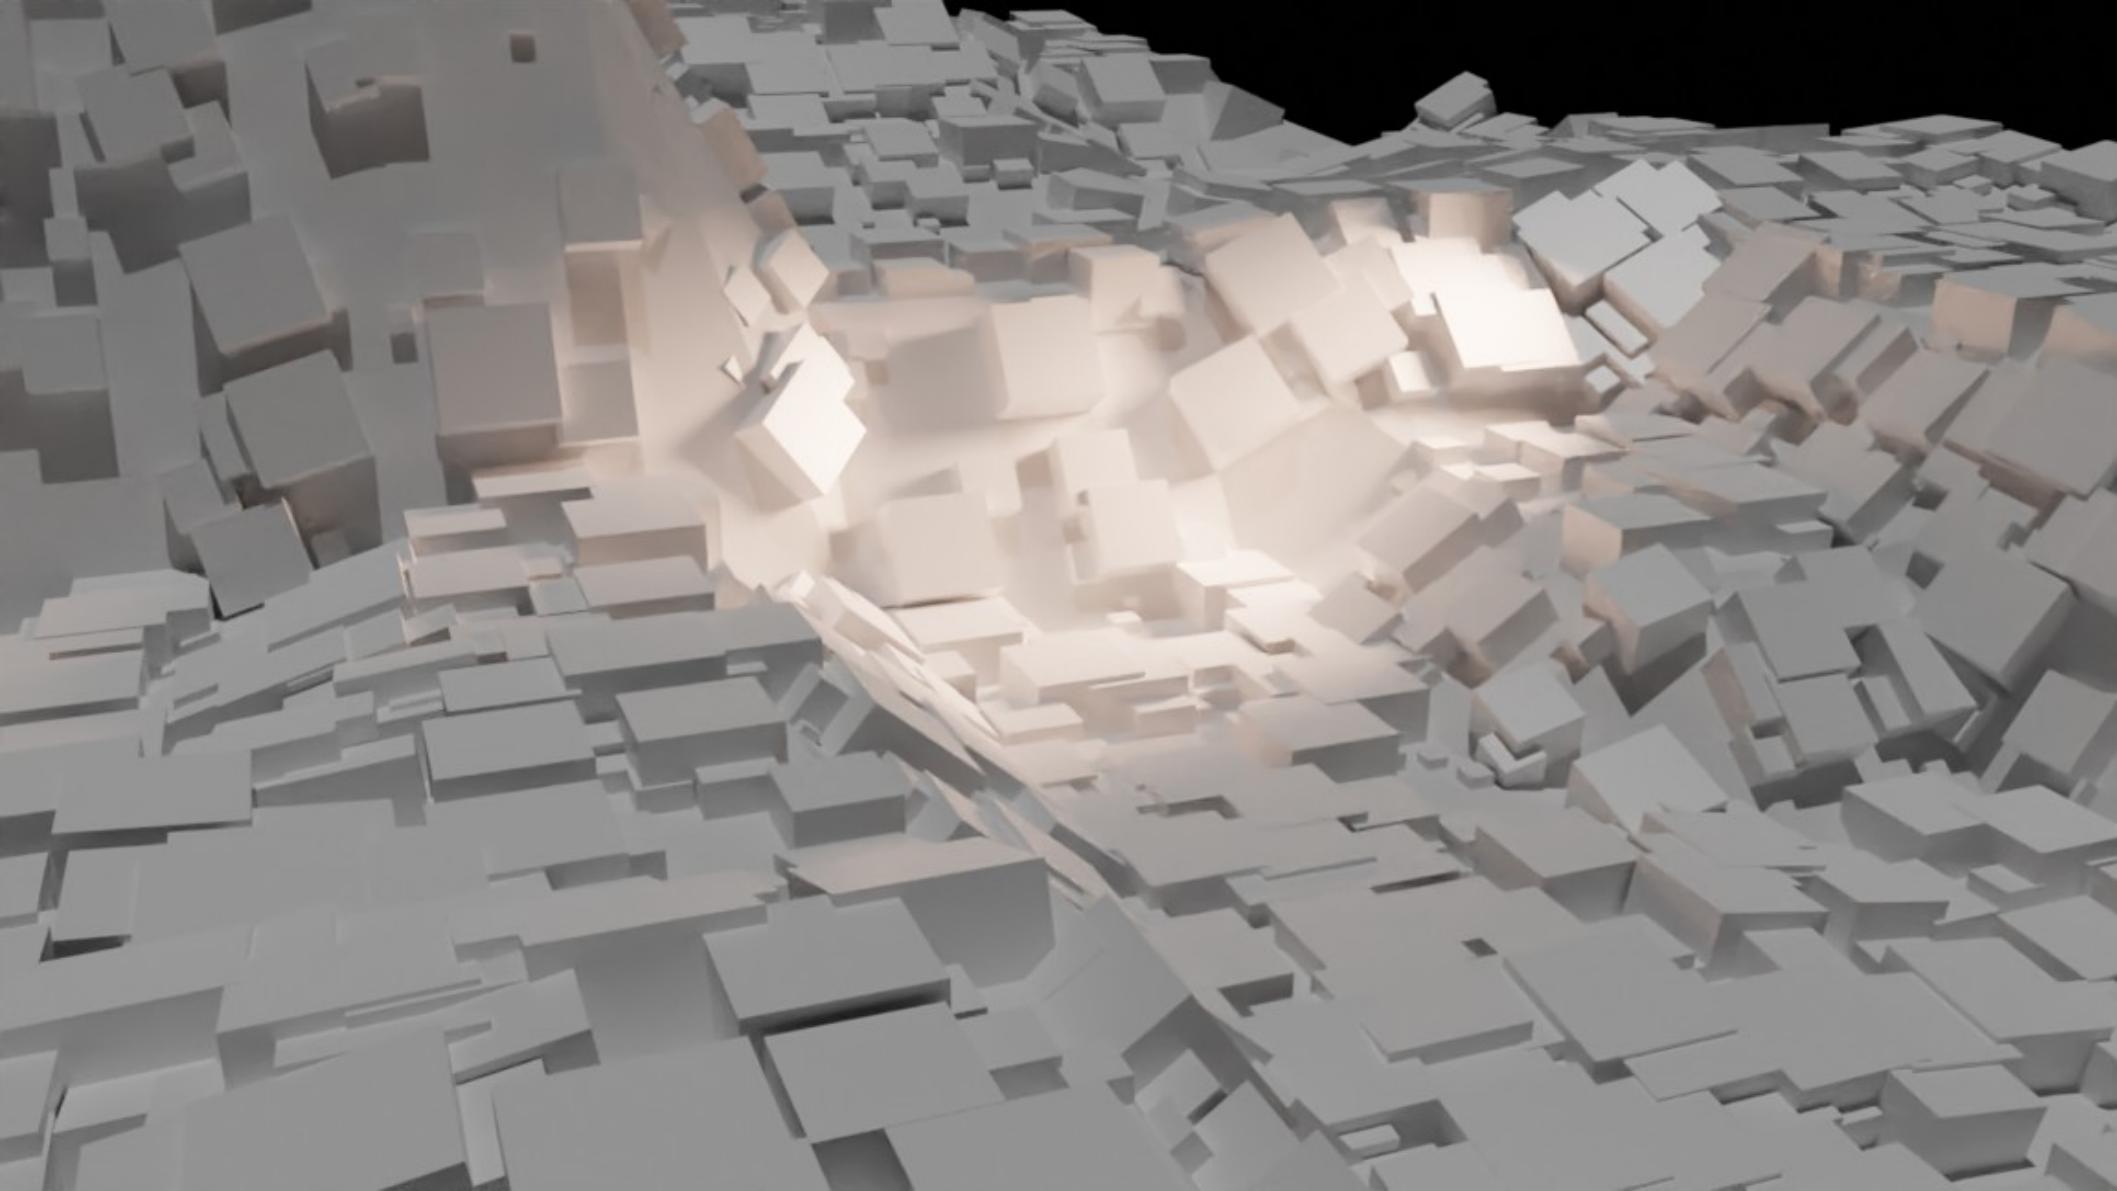
Cube mapping

- Most widely used method, Greene 1986
- Cube map is created by projecting environment on the sides of a cube positioned with its center at camera's location
- Visualized as "cross" diagram, stored as six square textures
- Reflected view vector (r) can be directly used as three-component texture coordinate to fetch intensity data stored in cube map



Shadows







Scene from:

<https://casual-effects.com/data/index.html>



Shadows

- Shadows are important for:
 - Providing user with visual cues about object placement and relation to other object
 - Realistic image synthesis: needed for understanding details and surface characteristics



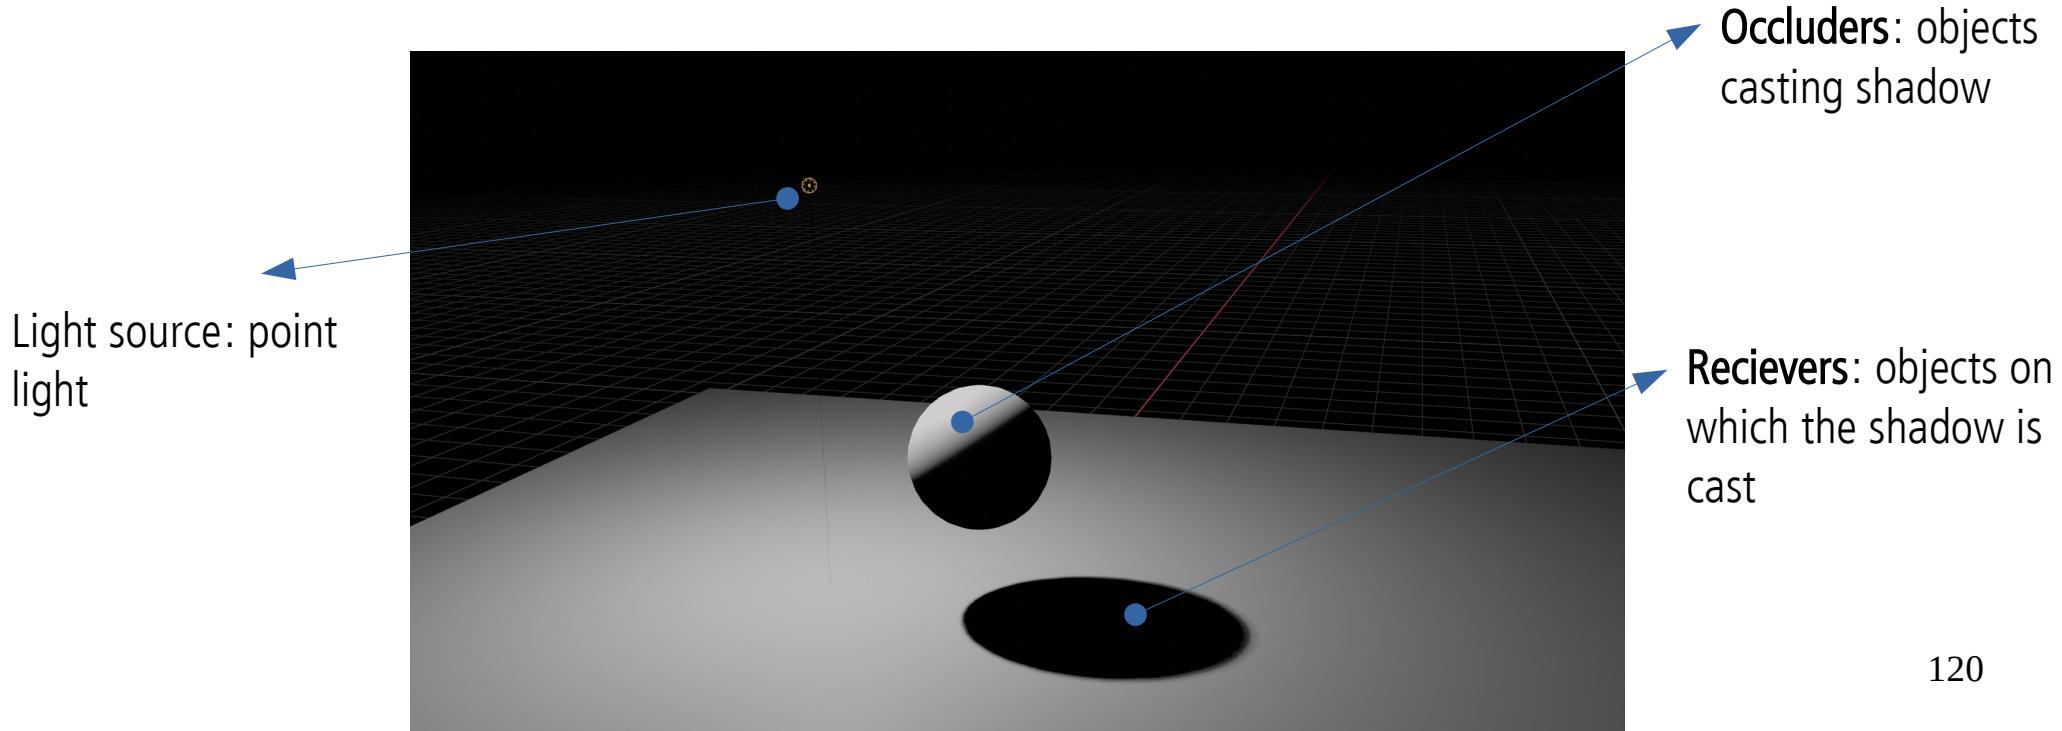
<https://creativecoding.soe.ucsc.edu/courses/cs488/reportsA/shadows.pdf>



<https://lumion.com/blog/feature-spotlight-sky-light-soft-shadows-and.html>

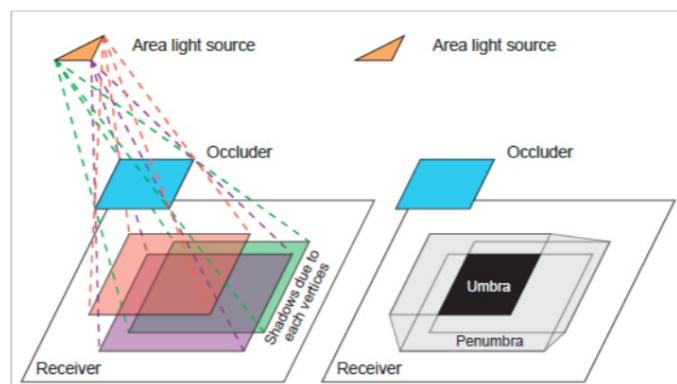
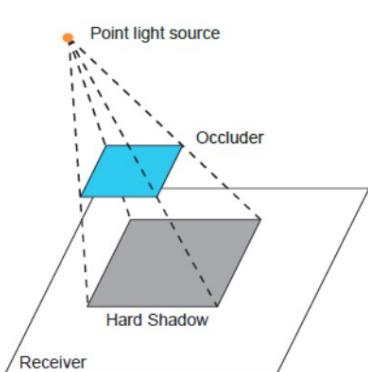
Light and shadow

- Light falling on object causes object to cast shadow – blocking light to fall on another surface



Sharp vs smooth shadows

- Size of light determines shadow smoothness
 - Hard-edge shadows: sharp shadow edges
 - Soft shadows: fuzzy shadow edges
 - Softer shadow edges are more realistic, but more expensive to render



Small light sources (e.g., point lights) → hard shadows; only umbra component

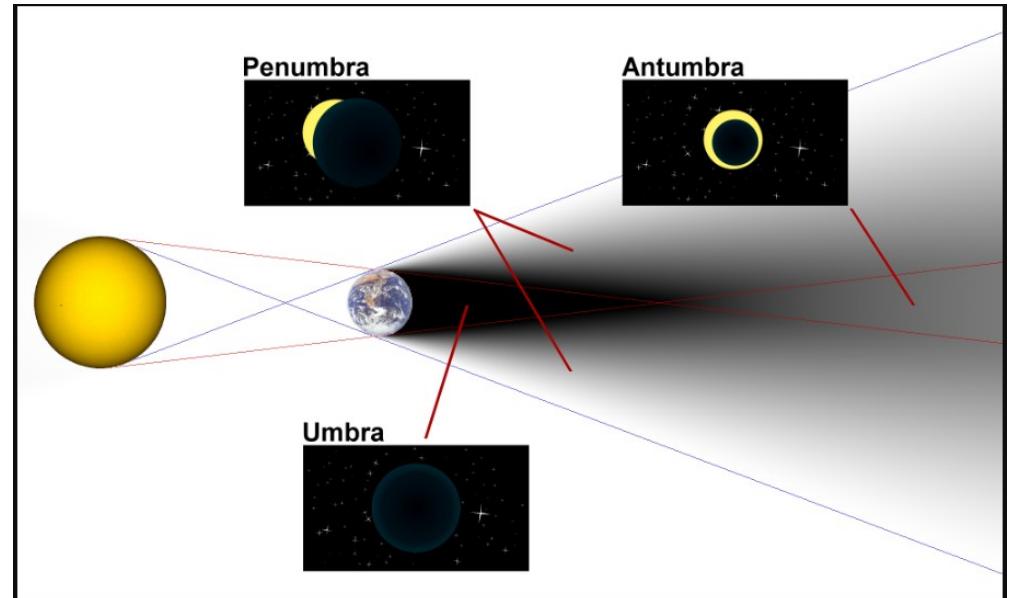


Large light sources (e.g, physical lights) → soft shadows; both umbra and penumbra component



Elements of a shadow

- **Penumbra**
 - Partially shadowed region
- **Umbra**
 - Fully shadowed region
 - Its region decreases in size as the light source gets larger
- NOTE: having any kind of shadow is more important than having soft shadows and penumbra.



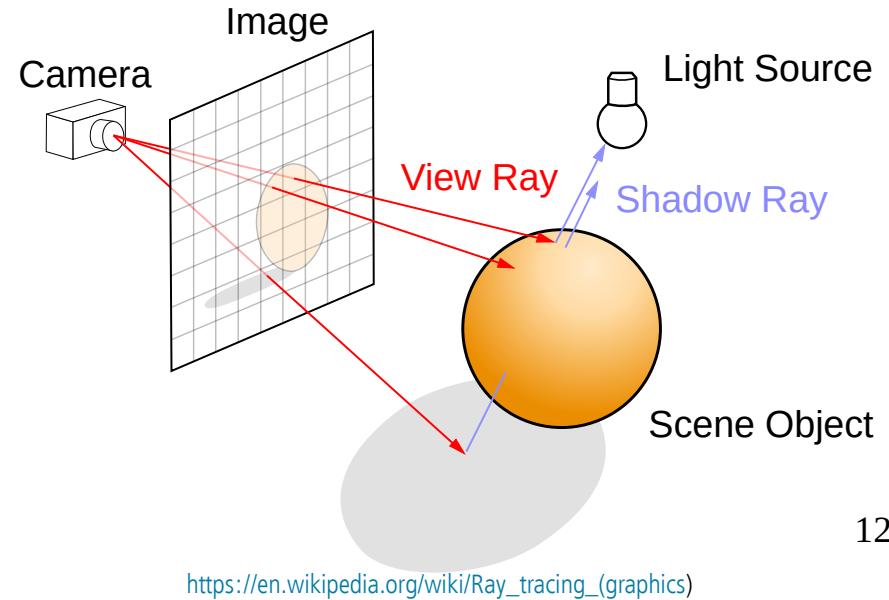
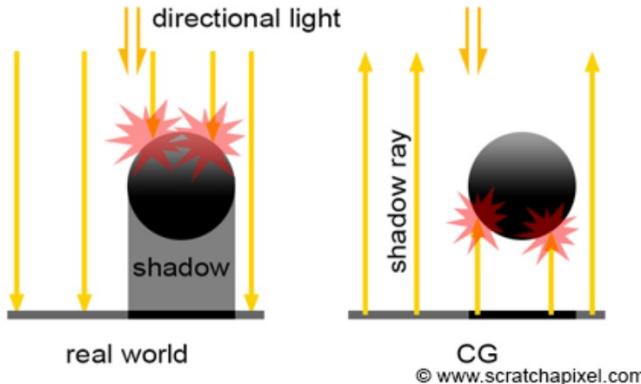
https://en.wikipedia.org/wiki/Umbra,_penumbra_and_antumbra

Rendering and shadow

- Simulating shadow depends on how rendering is solving visibility problem
- In ray-tracing based rendering shadows can be calculated in single pass without complex workarounds
- In rasterization-based rendering multi-pass rendering or additional work is required to simulate/approximate shadows

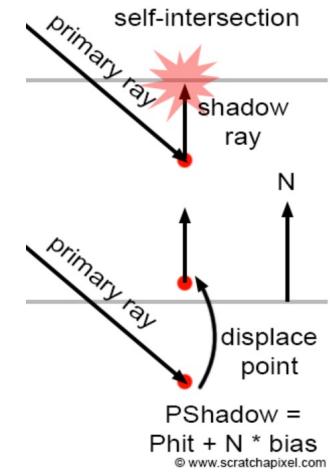
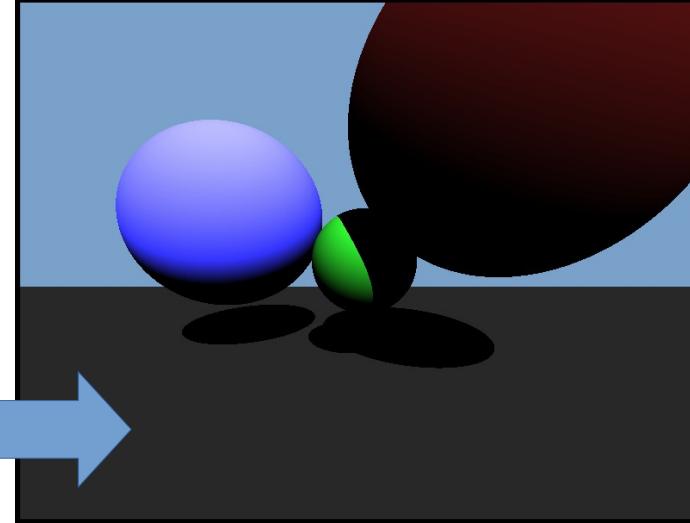
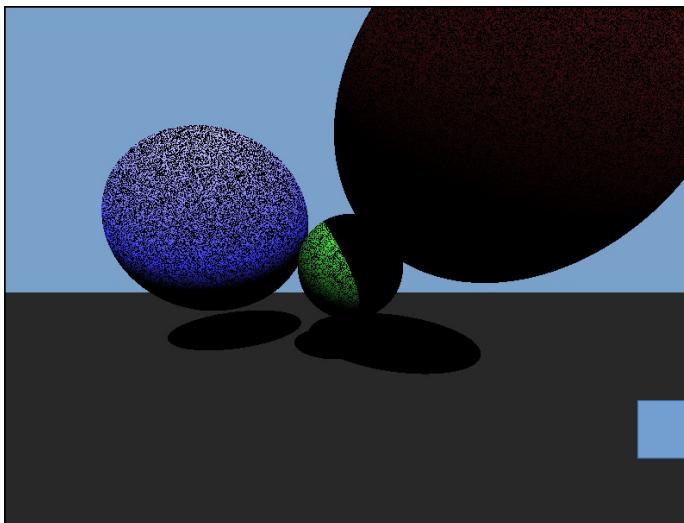
Shadow in ray-tracing

- For each pixel of virtual image plane, a ray is generated and tested for intersection with objects in 3D scene
- For found intersection, color is calculated using shading
- During shading, additional ray is casted towards light source – **shadow ray**
 - If shadow ray intersects object on its way, then shaded surface point is in shadow
 - If object is opaque, shadow is dark
 - If object is transparent, shadow is brighter



Shadow in ray-tracing: shadow acne

- Due to numerical limits in precision, some intersection point positions might end up under the surface
 - Casting shadow ray from this point causes intersection and thus light obstruction → shadow acne
 - Solutions:
 - Use double float precision
 - Add small displacement to intersection point in direction of normal: **shadow bias** parameter



<https://www.scratchapixel.com/lessons/3d-basic-rendering/introduction-to-shading/light-and-shadows>

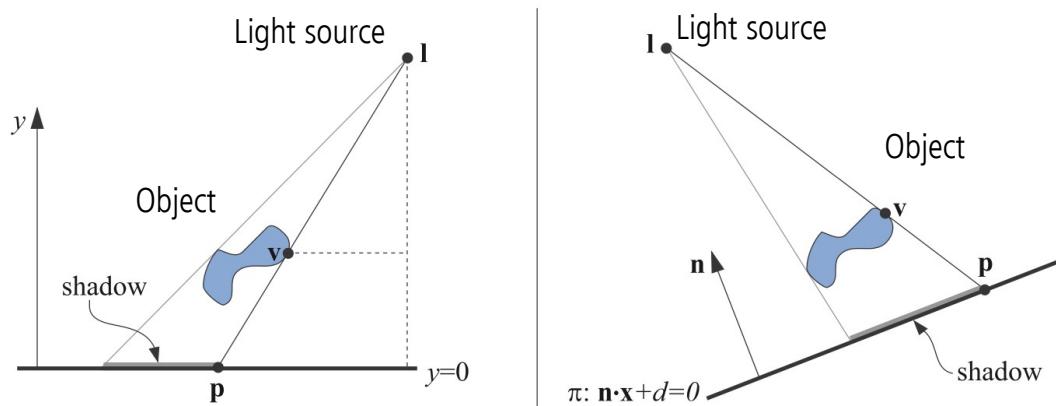
<https://bheisler.github.io/post/writing-raytracer-in-rust-part-2/>

Shadow in rasterization

- Shadow computation in ray-tracing is elegantly solved due to nature of visibility solving which simulates physical light flow
- In rasterization, all objects in 3D scene are projected onto virtual image plane and information of 3D scene is partially lost
- Therefore, various methods are developed to simulate and introduce shadows
 - Different methods solve different special and simplified cases of shadow occurrence
 - Planar shadows
 - Volume shadows
 - Shadow mapping

Planar shadows

- Simple case of shadowing: object cast shadow onto a planar surface
- Method: **projection shadows**



Projection matrix M is constructed to project object vertices v onto plane π with respect to the light source l giving shadow point p .

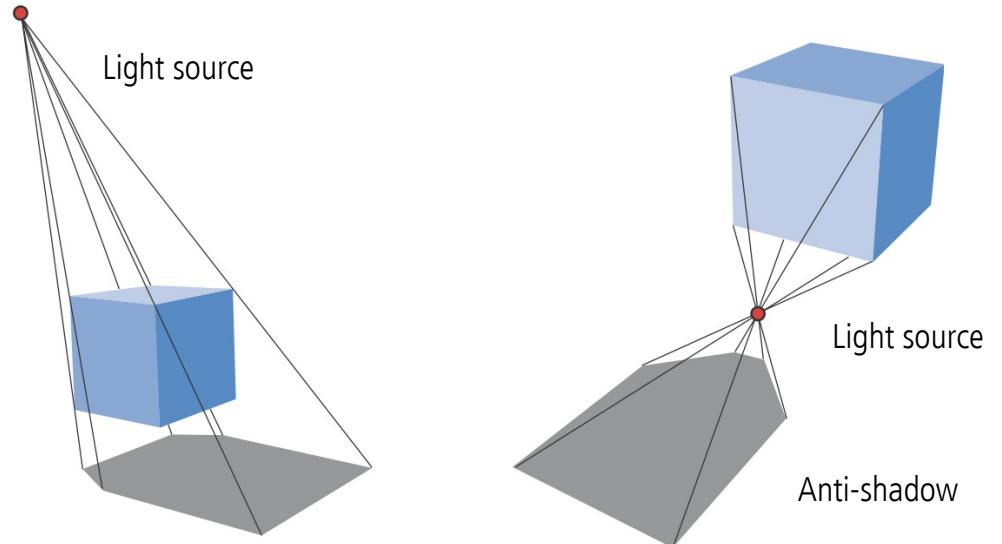
$$M = \begin{pmatrix} n \cdot l + d - l_x n_x & -l_x n_y & -l_x n_z & -l_x d \\ -l_y n_x & n \cdot l + d - l_y n_y & -l_y n_z & -l_y d \\ -l_z n_x & -l_z n_y & n \cdot l + d - l_z n_z & -l_z d \\ -n_x & -n_y & -n_z & n \cdot l \end{pmatrix}.$$

Projection shadows

- To render shadow, projection matrix M is applied to all objects (their vertices v) which should cast shadow on plane π
- Projected object is obtained which is rendered with darker color and no illumination.

Notes:

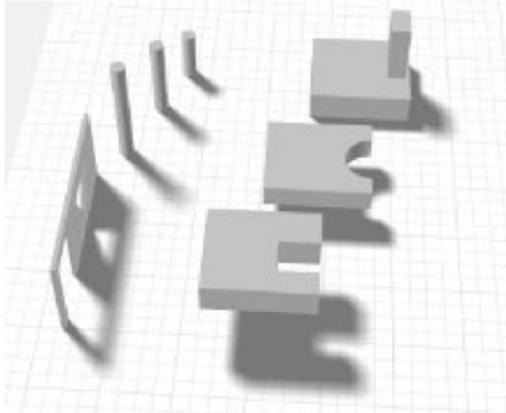
- Projected object must be above plane π (introduce bias parameter controlling distance)
- Shadows may fall outside of plane π – this breaks the illusion and clipping must be performed
- Anti-shadow can be generated if light is between object and plane π



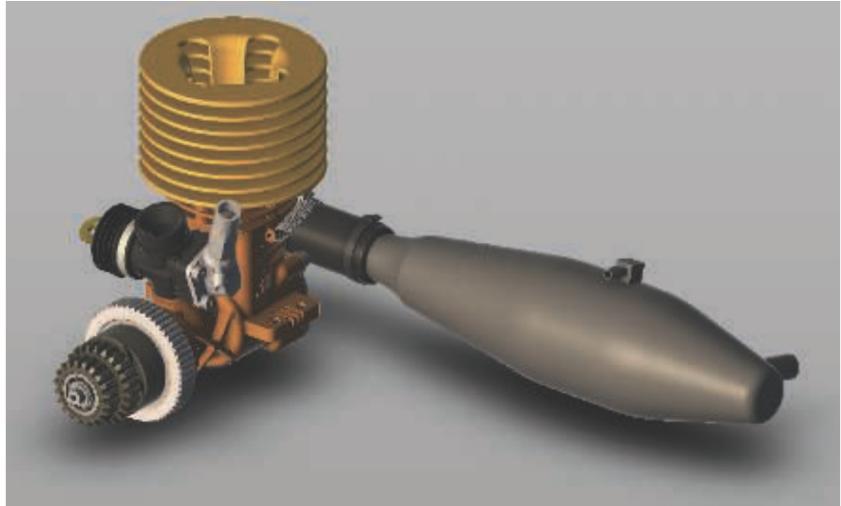
Projection shadows

- Soft shadows can be simulated.

Simulating area light: multiple point lights are placed where area light would be placed, resulting shadows are accumulated and averaged
(Heckbert and Herf's method)

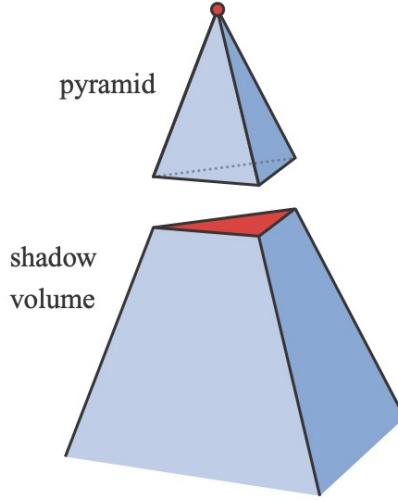
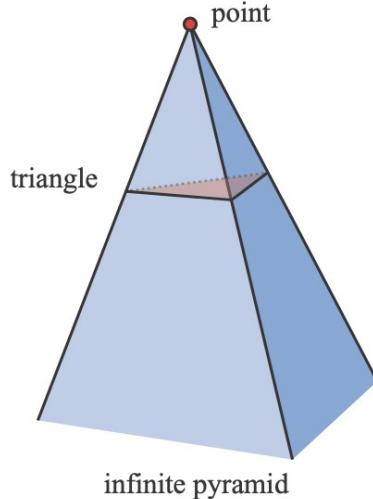


Blurring hard-shadow generated from a single point light.



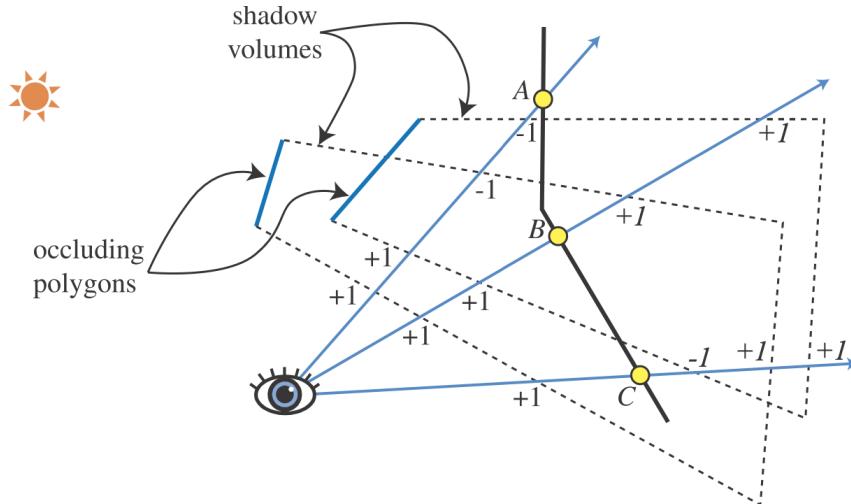
Shadow volumes

- More general approach for receivers of different shapes
- Volume of space is formed by occluder and light source
- Any object inside this volume is in shadow



Shadow volumes

- To determine if object is in shadow, counting method is used.
 - Ray is traced from view to object
 - Every time when ray passes through front-facing triangle of shadow volume the counter is increased
 - Every time when ray passes through back-facing triangle of shadow volume the counter is decreased
 - If counter is equal to zero than object is illuminated (not in shadow)

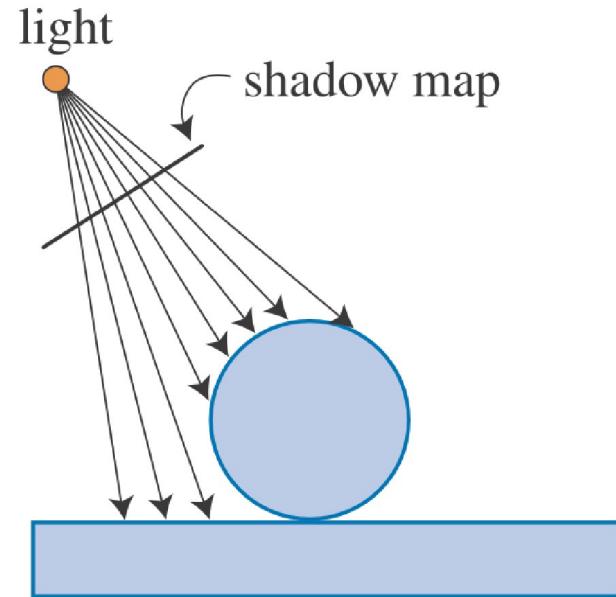


Shadow volumes

- Drawback: extreme variability
 - High difference in computation time depending on camera, object and light position
 - If camera is in shadow volume, the shadow volume fills the whole screen resulting in high computation → slow rendering
 - Consistent frame-rate is not guaranteed
- On GPU rasterization-based rendering counting is performed using depth and stencil buffers*

Shadow maps

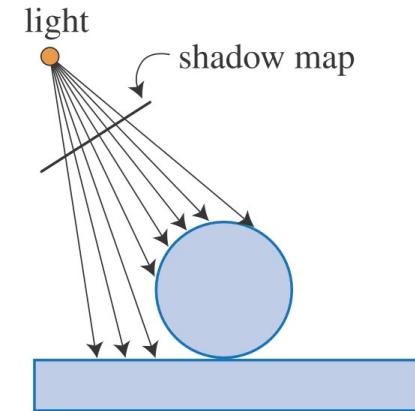
- Efficient rendering of shadows on various receivers
- Idea: solve shadow casting as **visibility problem**
 - Points visible to light source are illuminated
 - Points invisible to light source are in shadow
- GPU rasterization uses z-buffer to solve visibility and store distance to closest objects in depth buffer*



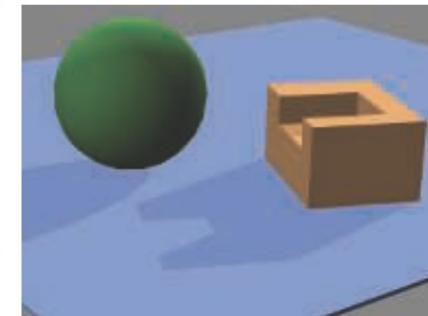
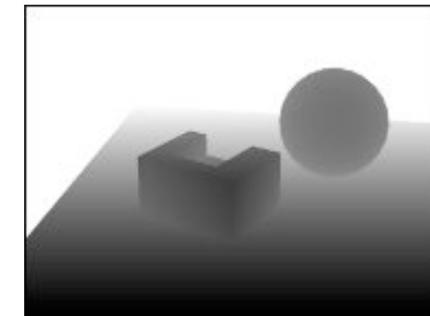
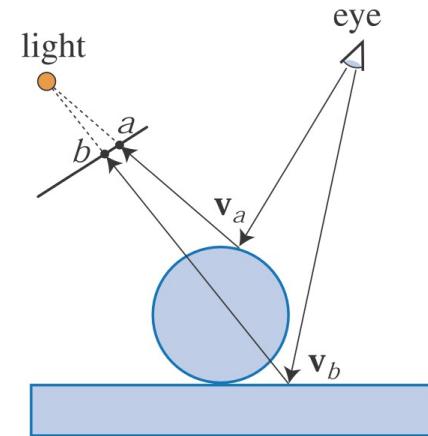
Shadow maps

- Render the scene from position of light source that is to cast shadows
 - Store distances of visible objects in a texture → **shadow map**
- Render the scene once more, now from camera point of view → **multi-pass rendering**
 - Distance of rendered points to light source are compared to distances stored in shadow map
 - If point is farther from light source than the corresponding value in shadow map, then point is in shadow, otherwise it is not

Rendering from light source



Rendering from camera



Shadow maps

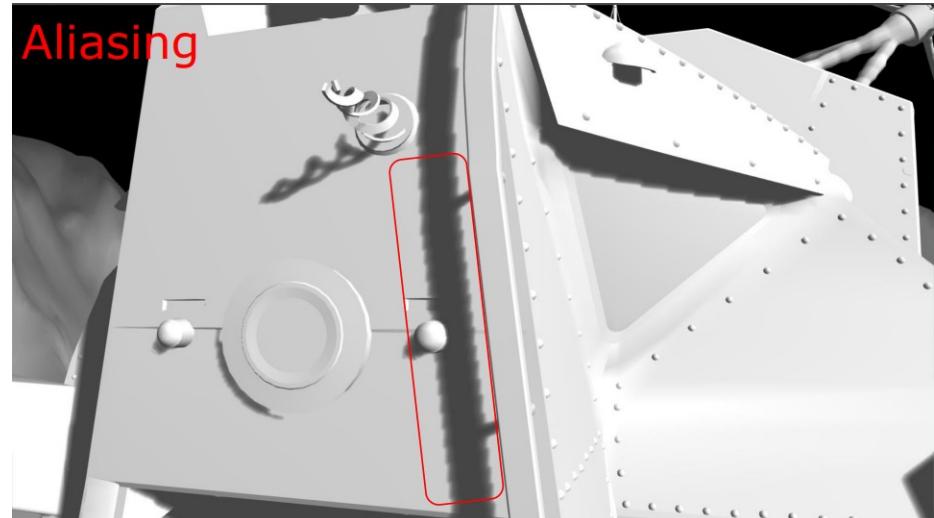
- **Shadow map creation depends on type of light:**
 - Distant lights: orthographic projection is used since light direction is parallel
 - Spotlight: frustum is defined with spotlight shape
 - Point light: omnidirectional shadow maps – six view cube (similar to environment mapping)
- **Shadow map approach is predictable**
 - Shadow map contains only objects that can cast shadows – objects inside light frustum
 - Cost of building shadow map is linear with number of triangles
 - Shadow map can be generated once and reused for all frames where light and objects are not moving

Shadow maps

- Quality of shadows depend on shadow map resolution
 - Quality of shadows depends on objects distance to light source → more samples for shadow map is needed for better approximation
 - Low shadow map resolution highlights aliasing

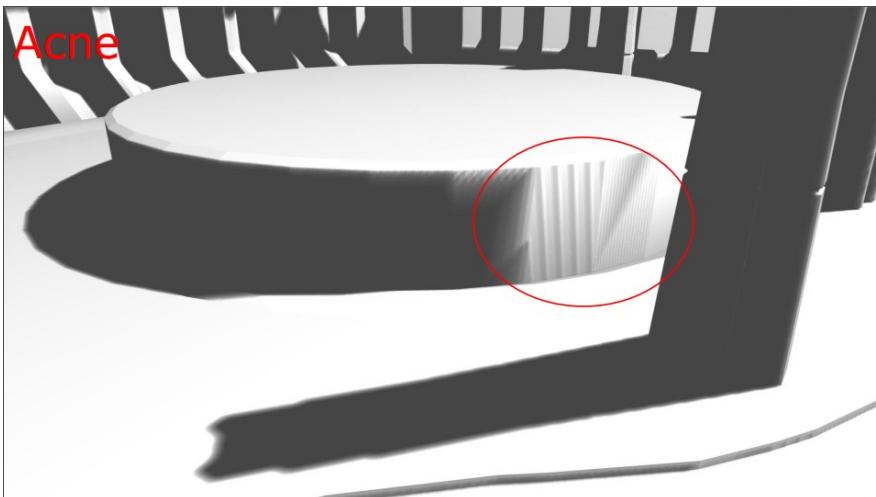


Quality of shadows depend on shadow map resolution

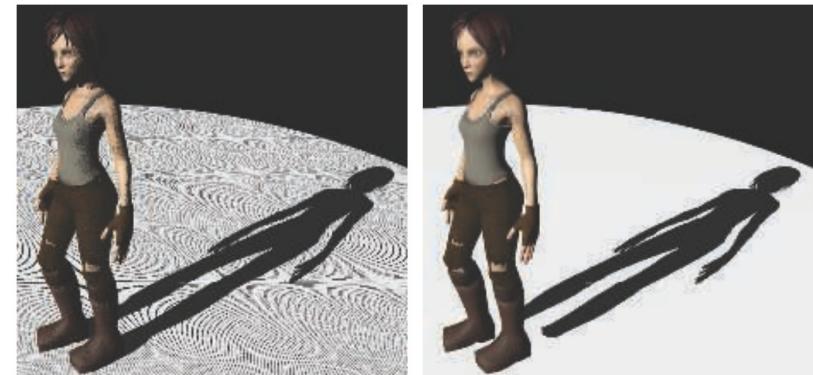


Shadow maps

- **Self-shadow aliasing and shadow acne** becomes apparent due to numerical limits and slight differences in light and camera view values
 - **Bias parameter** must be introduced which is correcting distances in camera and light view



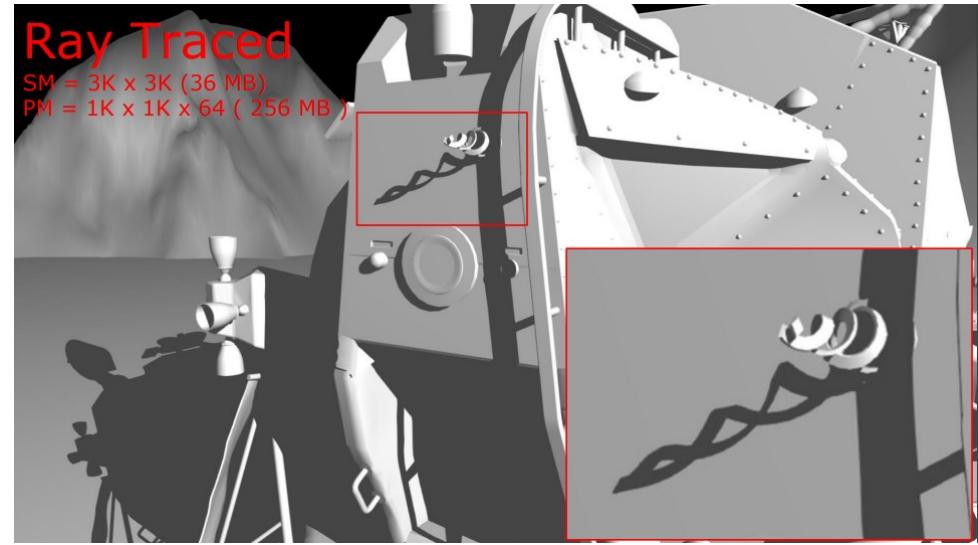
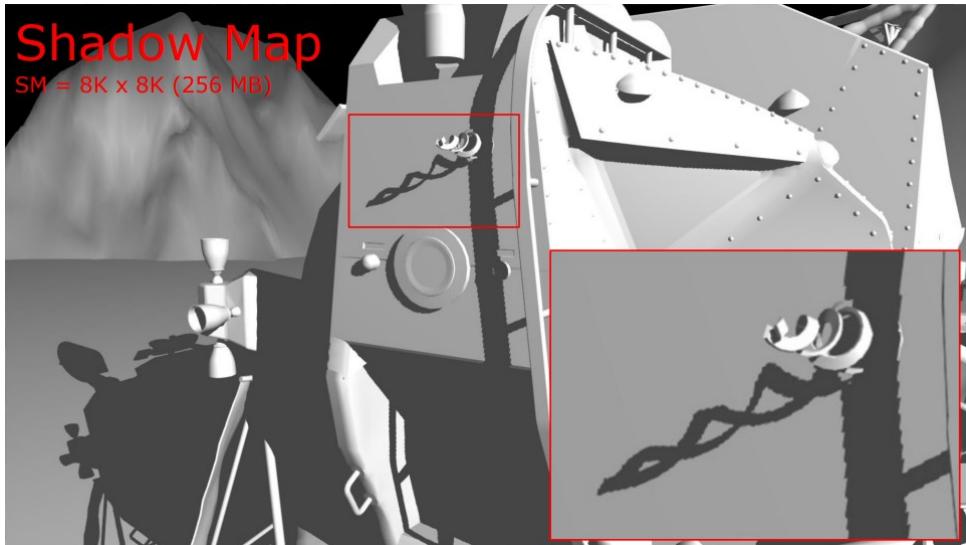
Self-shadowing (shadow acne) → moire pattern



Ground self-shadowing leads to aliasing (Moiré pattern)

Hybrid ray-traced shadows*

- Shadow maps are most often used method and still exhibit a lot of problems: aliasing, acne, etc.
- Solution is to combine rasterization-based rendering with ray-tracing for shadow computation



Light and Shadow: Tips and Tricks

TODO

- Non-physical light sources as well as shadows they cause can be simulated
- For real-time rendering, full physically based simulation of light transport is not feasible
- Non-physical lights can be placed to simulate global illumination
- When creating 3D scene, (environment) artists understand the how the scene should be illuminated given some main sources of light. Using this knowledge they place non-physical light to fake actual light flow through the scene
- <https://80.lv/articles/dishonored-interiors-lighting-props/>

More into topic

- Lighting Artstation: <https://www.artstation.com/channels/lighting>
- Lighting artist: <https://80.lv/articles/amazing-lighting-for-simple-scenes/>
- Lighting for artists: https://www.youtube.com/watch?v=Ys4793edotw&ab_channel=BlenderGuru
- Path-tracing lights
 - https://www.pbr-book.org/3ed-2018/Light_Sources/Area_Lights
 - http://graphics.stanford.edu/papers/veach_thesis/
 - <https://schuttejoe.github.io/post/arealightsampling/>
 - <https://www.scratchapixel.com/lessons/3d-basic-rendering/global-illumination-path-tracing/introduction-global-illumination-path-tracing.html>
- Area lights approximation in rasterization-based rendering (OpenGL): <https://learnopengl.com/Guest-Articles/2022/Area-Lights>
- Shadows in rasterization:
 - https://www.realtimeshadows.com/sites/default/files/Playing%20with%20Real-Time%20Shadows_0.pdf
 - <https://www.cs.montana.edu/courses/spring2005/525/students/Thiesen2.pdf>
 - https://web.cse.ohio-state.edu/~shen.94/781/Site/Slides_files/shadow.pdf
 - https://cg.informatik.uni-freiburg.de/course_notes/graphics_07_shadows.pdf
- Hybrid shadow computation methods (rasterization + ray-tracing):
 - <https://gpuopen.com/learn/hybrid-shadows/>
 - <https://developer.nvidia.com/content/hybrid-ray-traced-shadows>

Summary questions

TODO

Literature

- <https://github.com/lorentzo/IntroductionToComputerGraphics/wiki/Foundations-of-3D-scene-modeling>