

Lecture 14: Post-processing and Imaging pipeline

DHBW, Computer Graphics

Lovro Bosnar

29.3.2023.

Syllabus

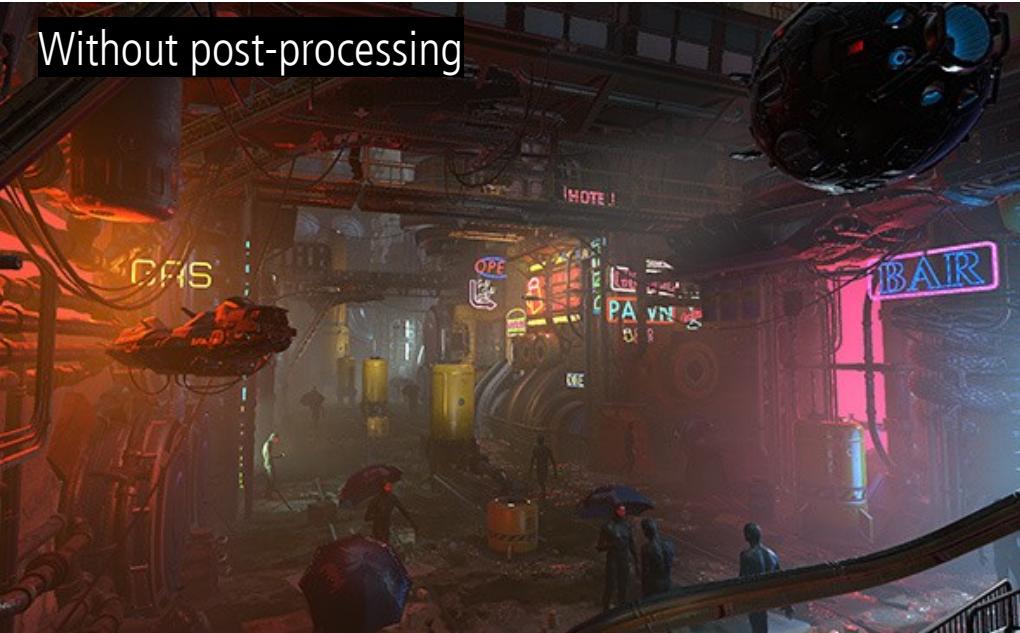
- 3D scene
 - Object
 - Light
 - Camera
 - Rendering
 - Image and display
-
- The diagram illustrates the syllabus structure. It starts with a list of topics: 3D scene, Rendering, and Image and display. A blue rounded rectangle encloses the 'Image and display' topic. A blue arrow points from the start of the 'Image and display' topic to the end of the '3D scene' topic. Another blue arrow points from the end of the 'Image and display' topic back to its beginning, forming a loop.
- Image and display
 - Image-space effects
 - Digital imagery
 - Aliasing and anti-aliasing
 - Scene to screen

Image-space effects

Image-space effects

- Rendered image can be directly sent to display or additionally modified before display → **post-processing**
- Post-processing is set of operations performed on a rendered image → **image-space effects**
 - This way, additional effects for achieving realism or artistic styles is possible

Without post-processing

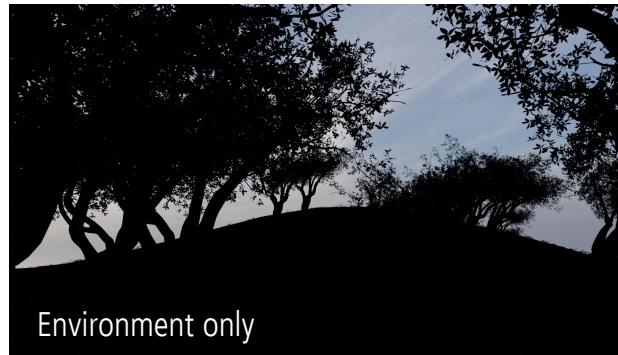
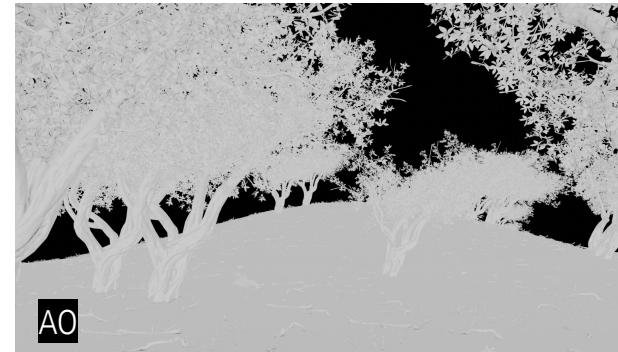
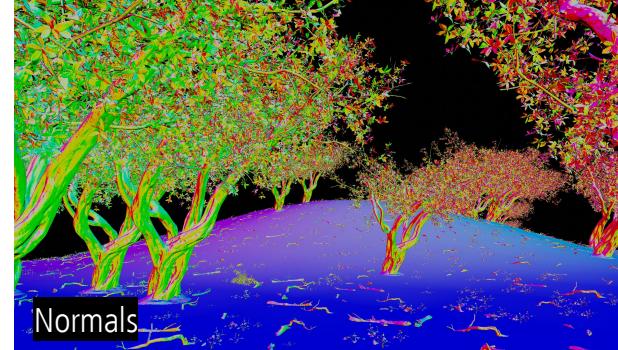


With post-processing



Post-processing data

- Post-processing of **one frame** can use multiple buffers (images) obtained during rendering containing 3D scene information:
 - Color, diffuse, normal, AO, depth, environment, etc.



Performing post-processing

- Depending on rendering environment, post-processing of frame can be performed:
 - **Real-time, multi-pass, render-time**
 - 3D scene is rendered to an **offscreen buffer**, such as color (image) buffer
 - Resulting image is treated as image texture which is **applied on screen filling quad**
 - Post-processing is done on this texture **using programmable GPU shaders** (e.g., fragment or compute shaders)
 - **Offline, compositing software**
 - 3D scene is rendered to image file
 - Post-processing is done on images which are stored as separate frames



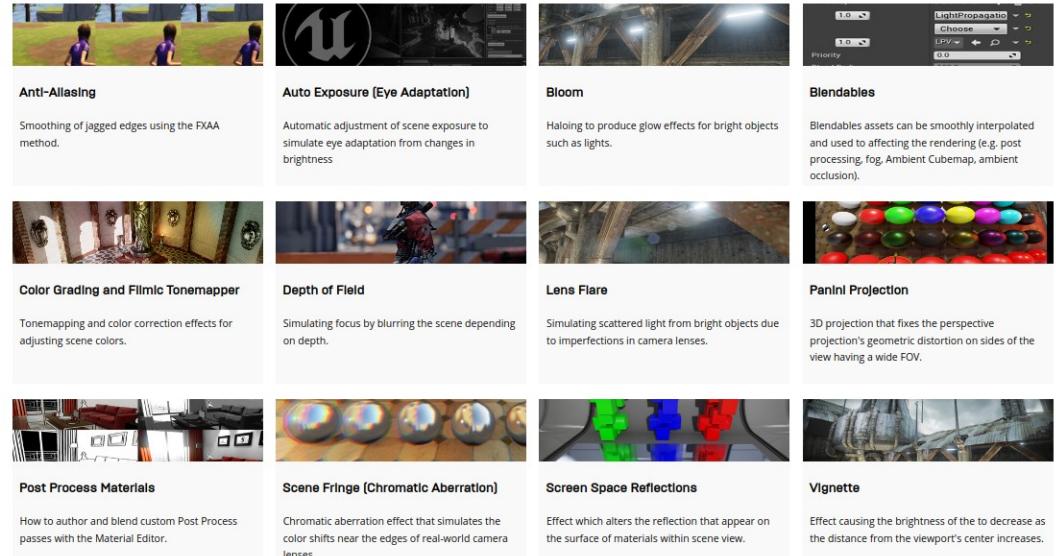
Post-processing in practice

- Real-time, multi-pass:

- Unity: <https://docs.unity3d.com/Manual/PostProcessingOverview.html>
- Godot: https://docs.godotengine.org/en/stable/tutorials/shaders/custom_postprocessing.html
- Unreal: <https://docs.unrealengine.com/4.27/en-US/RenderingAndGraphics/PostProcessEffects/>
- OpenGL: <https://learnopengl.com/In-Practice/2D-Game/Postprocessing>

- Offline:

- Blender: <https://docs.blender.org/manual/en/latest/compositing/introduction.html>
- Nuke: <https://www.foundry.com/products/nuke-family/nuke>
- After Effects: <https://www.adobe.com/products/aftereffects.html>
- Discussion: https://www.youtube.com/watch?v=7g4xCV0iv4w&ab_channel=InspirationTuts



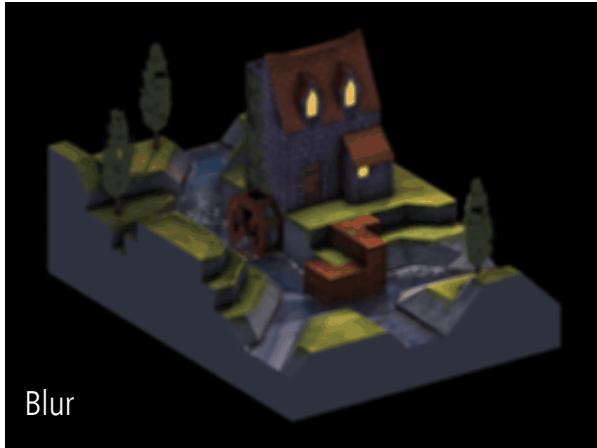
Unreal post-process



Nuke post-process

Post-processing using image processing

- Rendered image can be post-processed as any other image using image processing techniques (filtering kernels) in shader or CPU.



Blur



Posterization/quantization



Pixelization



Edge detect

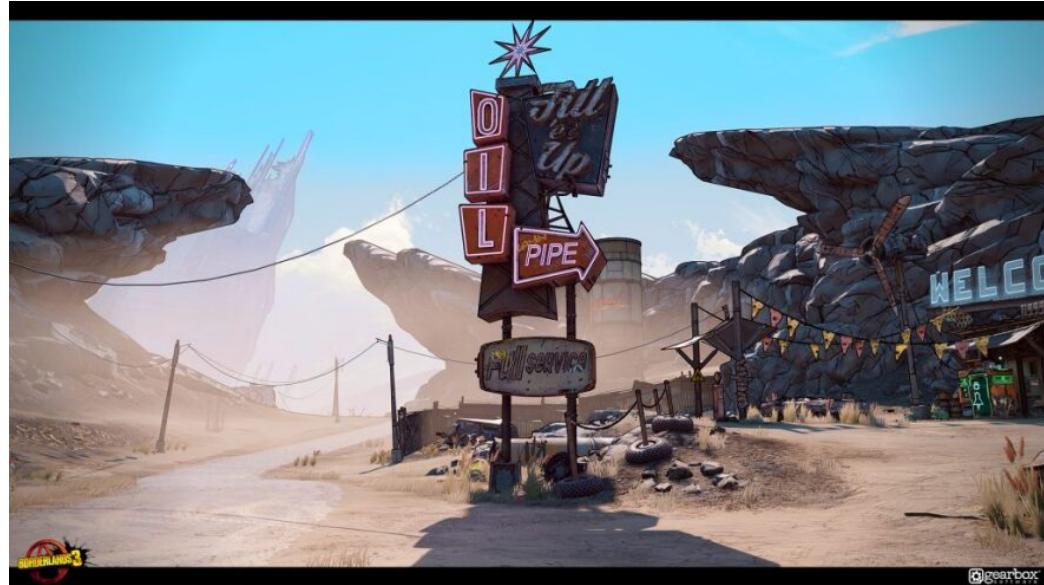


Dilatation

<https://lettier.github.io/3d-game-shaders-for-beginners/posterization.html>

Post-processing and NPR

- Often image processing techniques are used to achieve **non-photo realistic rendering effects**



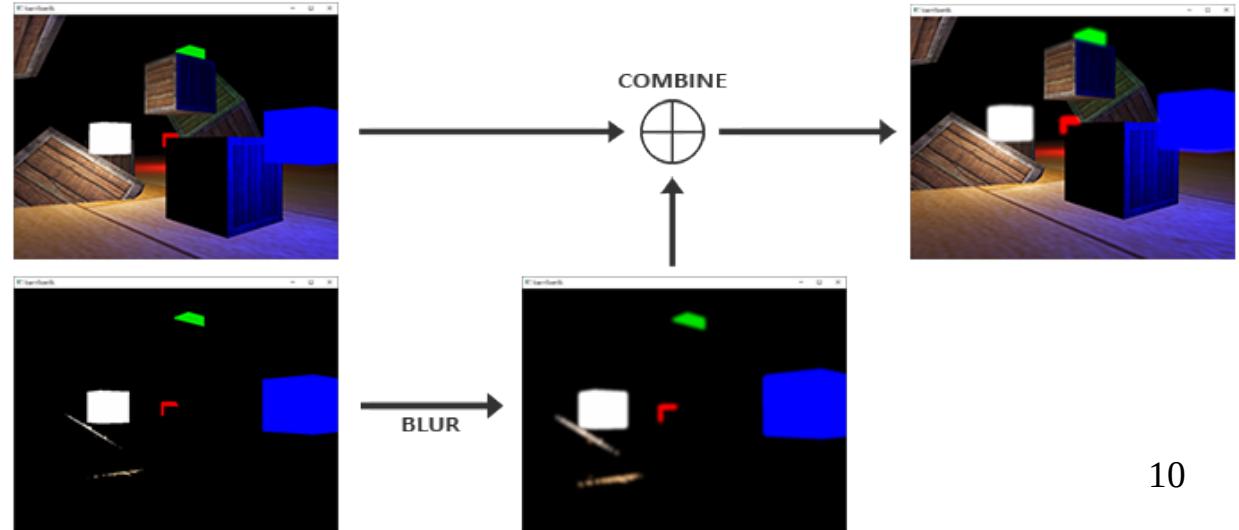
Edge detection (Borderlands)



Edge detection and quantization (Okami)

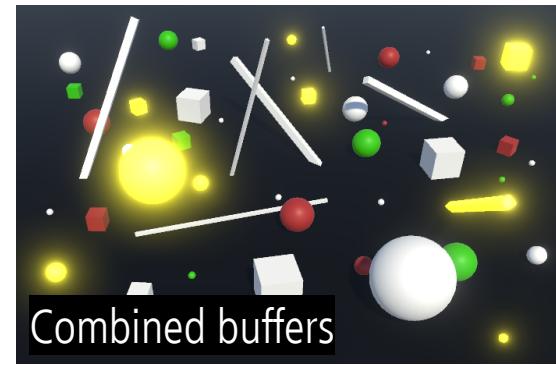
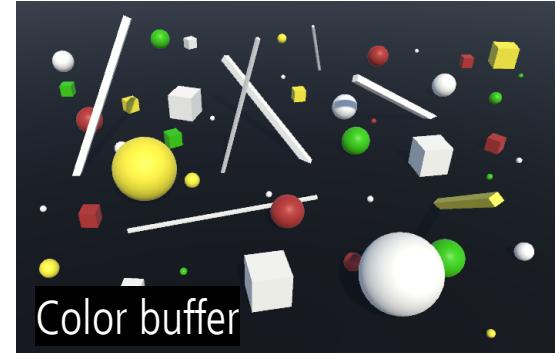
Post-processing: bloom

- Bloom occurs in camera when charged site in CCD gets saturated and overflows into neighboring sites
- Effect where **extremely bright area spills over adjoining pixels** → making image look overexposed where it is bright
- Concept:
 - Render whole scene in color buffer
 - Render only bright objects into separate buffer
 - Blur buffer with bright objects only (e.g., Gaussian Blur)
 - Combine buffers



Bloom: examples

- Bloom is one type of **glare effects**. Other glare effects are lens flare: halo and ciliary corona



Post-processing: lens flare

- **Flare phenomena:** caused by light traveling through a lens system
- Main types: halo (ring around light) and ciliary corona (rays radiating from a point)
 - Dependent on light source position
- Various approaches exist. Examples:
 - Set of squares with different textures, oriented on a line going from the light source potion screen through the screen's center
 - Light streaks from bright object can be simulated using steerable filter on a down sampled image



Lens flare (Battlefield)



Lens flare simulation <https://resources.mpi-inf.mpg.de/lensflareRendering/pdf/flare.pdf>

Depth of field

- Depth of field: range where objects are in focus. Outside objects are blurred
- Smaller camera aperture size → increased depth of field
 - Pinhole cameras: infinite depth of field
- Various methods exist. One solution:
 - Create separate image layers: e.g. near, far
 - Blur near or far layers
 - Final image is made by combining layers using back to front compositing



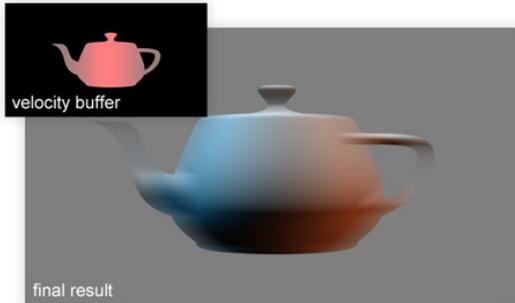
<https://www.versluis.com/2022/01/ue4-dof/>

<https://catlikecoding.com/unity/tutorials/advanced-rendering/depth-of-field/>

<https://developer.nvidia.com/gpugems/gpugems/part-iv-image-processing/chapter-23-depth-field-survey-techniques>

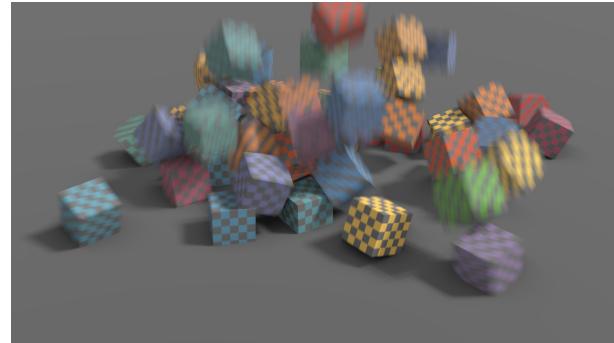
Motion blur

- Motion blur comes from **movement of an object** across the screen or from **camera movement**
- Different sources of motion blur require different methods
 - Camera orientation changes → e.g., directional blur based on orientation direction
 - Camera position changes → radial blur
 - Object position/orientation changes → motion of each object is computed using **depth or velocity buffer** for determining amount of blur



Object rotation and position change blur using velocity buffer

<http://john-chapman-graphics.blogspot.com/2013/01/per-object-motion-blur.html>



Directional blur due to camera orientation change (Tomb raider)



Radial blur due to camera position change (Need for speed)

Digital imagery

Digital imagery

- Most sources of digital imagery are:
 - Digital photographs or other types of 2D pictures scanned or loaded in computer
 - Images generated using computer graphics: 3D modeling and rendering software



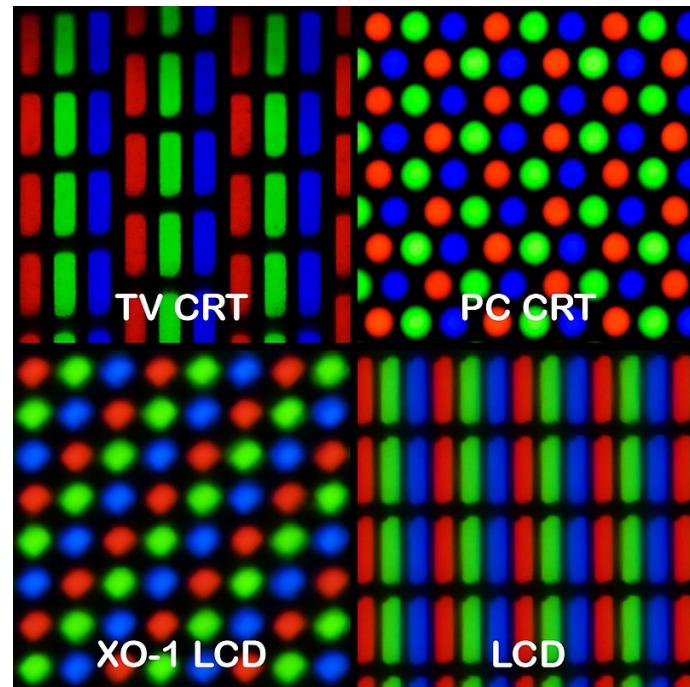
Digital photography <https://www.warnerbros.com/movies/blade-runner-2049>



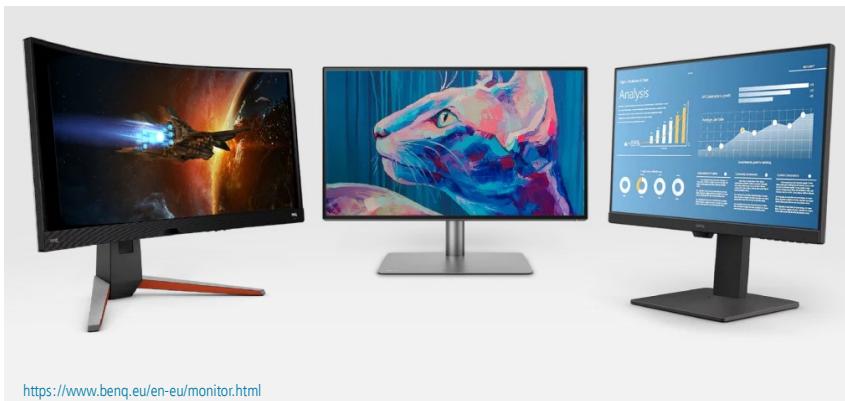
Rendered image: <https://www.thewitcher.com/en/witcher3>

Display device

- Digital images are displayed on a **raster display device**
→ **raster images**
 - e.g., monitor, television screen, etc.
- Raster display device is made of **arrays of pixels** → **discrete representation**
 - For example, each pixel of image created with digital camera (or rendering) stores red, green and blue value used to drive the red, green and blue values of screen pixel



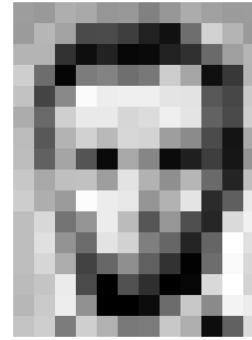
https://en.wikipedia.org/wiki/Pixel_geometry



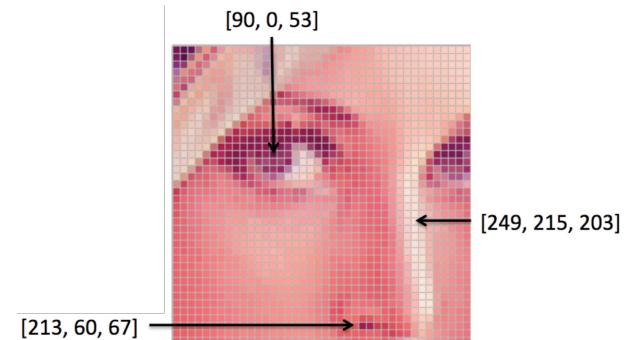
<https://www.benq.eu/en-eu/monitor.html>

Raster image graphics

- Pixels in raster image have the same type, examples:
 - Floating point numbers representing levels of **gray**
 - Floating point triplets representing mixture of red, green and blue (**RGB**)
 - Floating point numbers representing **depth** from camera
- Values in raster images can be interpreted in different ways
 - Arbitrary information can be stored/encoded in such array (e.g., normal maps)
 - Numbers in array do not have particular significance until stored in certain **standardized file format**



157	153	174	168	150	162	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	18	154
180	50	14	54	6	10	53	48	105	159	181	181
206	109	6	134	133	111	123	204	168	15	56	180
194	68	137	251	237	229	228	227	87	71	201	201
172	106	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	199	75	20	169
169	97	155	84	10	168	134	11	31	62	22	148
199	168	191	153	158	227	178	143	182	105	36	190
205	174	155	252	236	231	149	178	228	43	95	234
160	216	116	149	235	187	85	150	79	38	218	241
100	224	147	108	227	210	127	123	36	101	255	224
100	214	173	66	103	143	96	50	2	109	249	218
187	196	238	73	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	95	218



Naive image file format: PPM

- Portable Pixel Map (PPM): text-based (P3) image file format

- Header:

- Width (w) and height (h) of image is specified
 - Maximum color value

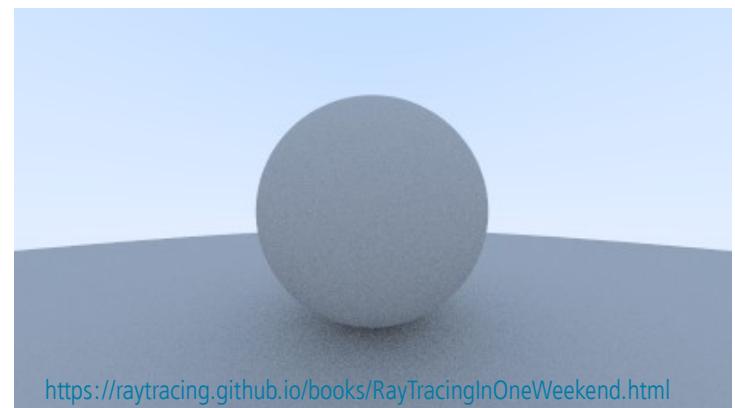
- Pixel values:

- $3 * w * h$ color values representing red, green and blue components of pixel
 - Pixel values ordering: left-to-right, top-to-bottom
 - Separated by white-space

- Naive image file format:

- Easy to write and analyze
 - Inefficient and highly redundant (compression is not present)
 - Little information about image

```
P3
# feep.ppm
4 4
15
 0  0  0   0  0  0   0  0  0   15  0  15
 0  0  0   0  15  7   0  0  0   0  0  0
 0  0  0   0  0  0   0  15  7   0  0  0
15  0  15   0  0  0   0  0  0   0  0  0
```



<https://raytracing.github.io/books/RayTracingInOneWeekend.html>

Raster image file formats

Format	Channel Depth	Alpha	Meta data	DPI	Extensions
BMP	8bit	✗	✗	✓	.bmp
Iris	8, 16bit	✓	✗	✗	.sgi .rgb .bw
PNG	8, 16bit	✓	✓	✓	.png
JPEG	8bit	✗	✓	✓	.jpg .jpeg
JPEG 2000	8, 12, 16bit	✓	✗	✗	.jp2 .j2c
Targa	8bit	✓	✗	✗	.tga
Cineon & DPX	8, 10, 12, 16bit	✓	✗	✗	.cin .dpx
OpenEXR	float 16, 32bit	✓	✓	✓	.exr
Radiance HDR	float	✓	✗	✗	.hdr
TIFF	8, 16bit	✓	✗	✓	.tif .tiff
WebP	8bit	✓	✓	✓	.webp

Raster image file formats

- Images are stored in many formats, closely related to the **display format**
- File formats use different **compression strategies**:
 - **Lossless** compression is one in which data occupies less space but from which original data can be reconstructed.
 - **Lossy** compression is resulting in even less occupied space, but original data can not be restored. Nevertheless, existing data is enough for intended use
- Content is described in terms of **channels**
 - e.g., red values for all pixels represents red **color channel**
 - e.g., depth values form **depth channel**
 - e.g., transparency values from **alpha channel**
- Some formats additionally store **metadata**
 - e.g., bit depth of color channel (e.g., 8 bit)
 - e.g. information on when and with which program the image was produced

File formats in production

- Digital painting and image manipulation
 - Krita: bmp, jp2, **jpeg**, ora, pdf, **png**, ppm, raw, **tiff**, xcf <https://krita.org/en/item/krita-features/>
 - Gimp: XCF, **JPG**, **PNG**, GIF, **TIFF**, Raw, etc. <https://www.gimp.org/tutorials/ImageFormats/>
- 3D modeling
 - Blender: BMP, Iris, **PNG**, **JPEG**, Targa, OpenEXR , **TIFF**, HDR, etc.
https://docs.blender.org/manual/en/latest/files/media/image_formats.html
 - Houdini: **png**, gif, **jpg**, **tiff**, etc. https://www.sidefx.com/docs/houdini/io/formats/image_formats.html
- Game engines:
 - Unity: BMP, **TIF**, TGA, **JPG**, **PNG**, PSD, etc. <https://docs.unity3d.com/2019.2/Documentation/Manual/AssetTypes.html>
 - Godot: BMP, OpenEXR, **JPEG**, **PNG**, SVG, etc.
https://docs.godotengine.org/en/stable/tutorials/assets_pipeline/importing_images.html

Raster image file formats

- Often used raster image file formats are:
 - Joint Photographic Experts Group - **JPEG**
 - Portable Network Graphics - **PNG**
 - Tag Image File Format - **TIFF/TIF**
- Advanced image file formats:
 - **OpenEXR** (<https://github.com/AcademySoftwareFoundation/openexr>)
 - **HDRI**: High Dynamic Range Image
- Library for handling different file formats for computer graphics applications:
<https://github.com/OpenImageIO/oiiio>

JPEG

- Recommended for display and storage of photography

Pros:

- Efficient file **compression**
- **Universally** supported for display
- Most **digital cameras** produce JPEG images

Cons:

- **Lossy**
 - Problematic for comparing the image due to differences in compression algorithm
 - Repeated editing degrades image quality
- **Artifacts** can be seen for computer generated graphics and text
- **Doesn't support transparency**
- Color channels are coded on 8 bits

PNG

- Developed as an improved replacement for Graphics Interchange Format (GIF)
- Recommended for web page widgets, computer graphics and screenshots

Pros:

- **Lossless** format
- Supports **transparency**
- **Small file size** for most computer graphics generated images
- Support by all **browsers**
- PNG file format is more compact than PPM and equally **easy to use**
- Support for **RGB(A)** and **grayscale images**

Cons:

- Complex images are heavier
- Color channels coded in **8 bits**

TIFF

- Used for storage and exchange of high quality images
- TIFF is ideal for representing intermediate or final results
 - In image editing and compositing tools, multiple images are often blended or laid atop one another

Pros:

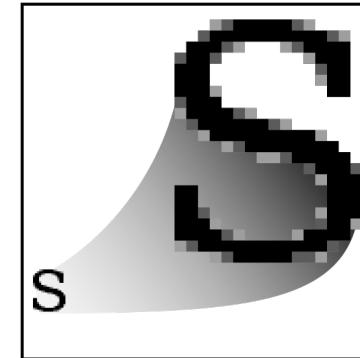
- Color coding in **16 bits**
- **Supported** by all image processing software
- TIFF images store **multiple channels** (images/layers)
 - Each channel store description of contents
- Various **compression** schemes
 - **uncompressed**
 - Compressed with **lossless** scheme
 - Compressed with **lossy** scheme

Cons:

- **Heavier** on complex images

Vector graphics images

- Alternative to raster graphics image is **vector graphics image**
 - Images are created directly from **geometric shapes defined on 2D Cartesian coordinate system**
 - Information in such image is **stored as points, lines, curves and polygons**
- Today, raster-based monitors and printers are typically used
 - Nevertheless, vector data and software is **used in applications where geometric precision is required: engineering, architecture, typography, etc.**
 - To display an image, **rendering** is performed to evaluate analytically defines shapes on raster display



Raster
GIF, JPEG, PNG



Vector
SVG

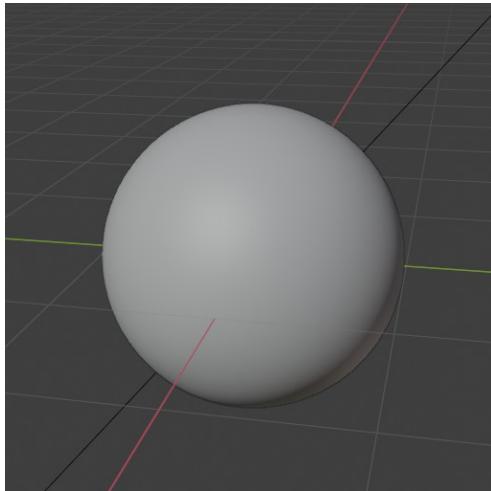
https://en.wikipedia.org/wiki/Vector_graphics

Aliasing and anti-aliasing

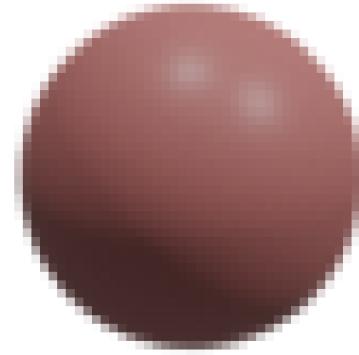
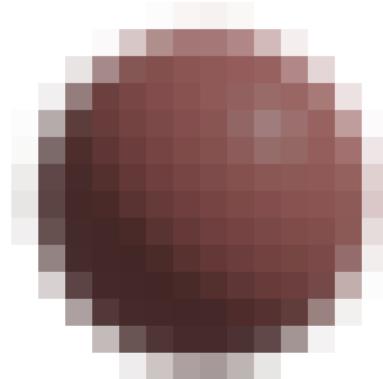
When continuous becomes discrete

Continuous and discrete

- Images are **discrete** array of pixels
- 3D scene objects, lights and thus color values are **continuously** changing
- Rendering process, in order to create a discrete image of 3D scene, **samples the pixels on image plane**, finds corresponding sample in 3D scene and calculates color of the sample.

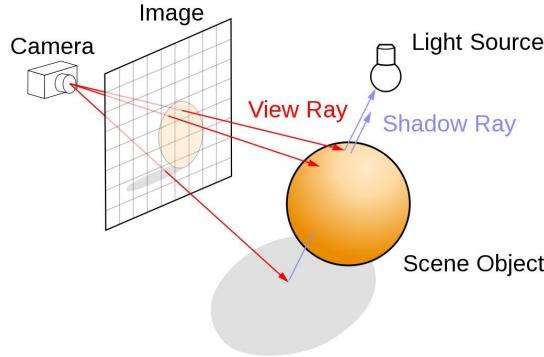


Continuous object in 3D scene



Scene sampling and Image buffer

- Process of **rendering** images is inherently a **sampling task**
- Rendering is the process of **sampling a 3D scene** in order to obtain **colors for each pixel** in the image
 - Using camera information, **samples per image pixel** are generated
 - For each sample, **corresponding 3D scene point** is found and color is calculated (ray-tracing, rasterization)
 - Color of each pixel for each sample is combined into **color buffer**

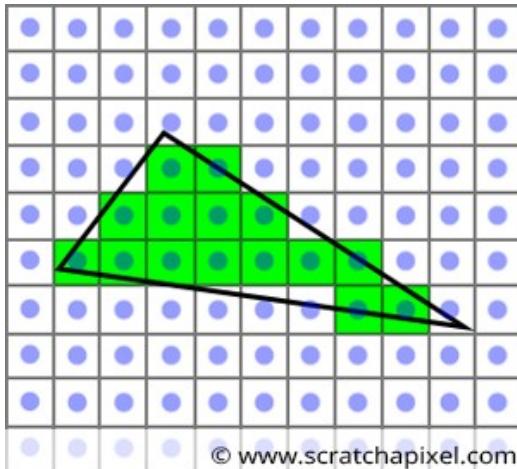


Note that pixel can have only one color!

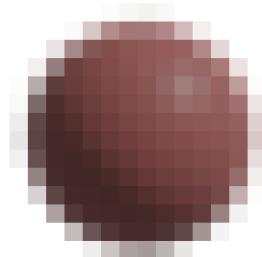
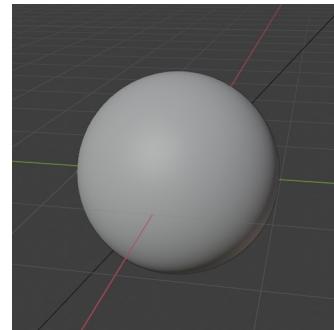
Pixels overlapping multiple objects/textures in 3D scene must be sampled with multiple rays to obtain correct representation.

Aliasing

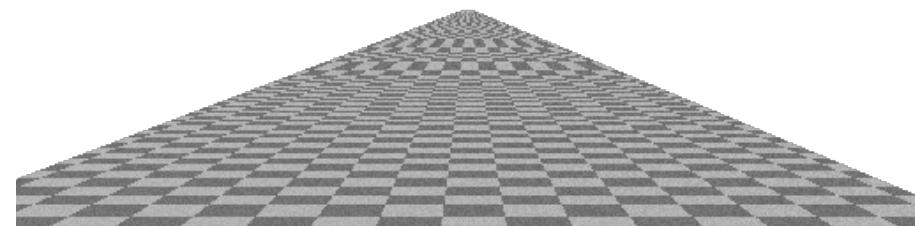
- Discretizing 3D scene (continuous information) may cause **Aliasing**
 - Jagged edges, flickering highlights (firelfys), Moiré pattern



Discretization



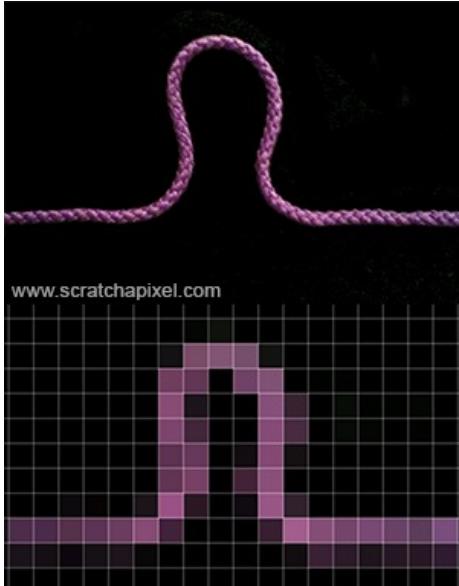
Problem: Jagged edges



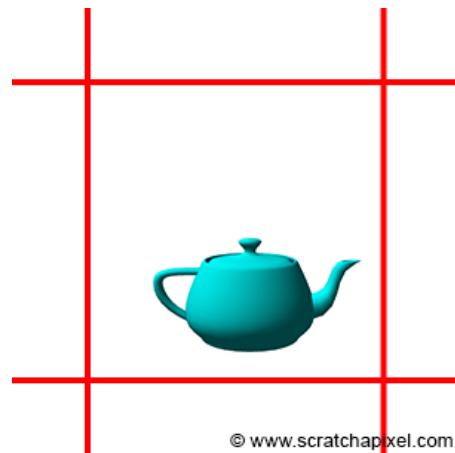
Problem: Moiré pattern

Aliasing: intuition

- Projecting pixel into 3D scene can cover larger portion of 3D scene which contain multiple objects and thus color variation.
- Any kind of rapid (high-frequency) change in geometry edges, textures, shadows, highlights can cause aliasing.



Thread contains large amount of details covered by only one pixel



Extreme example: one whole object in pixel footprint.

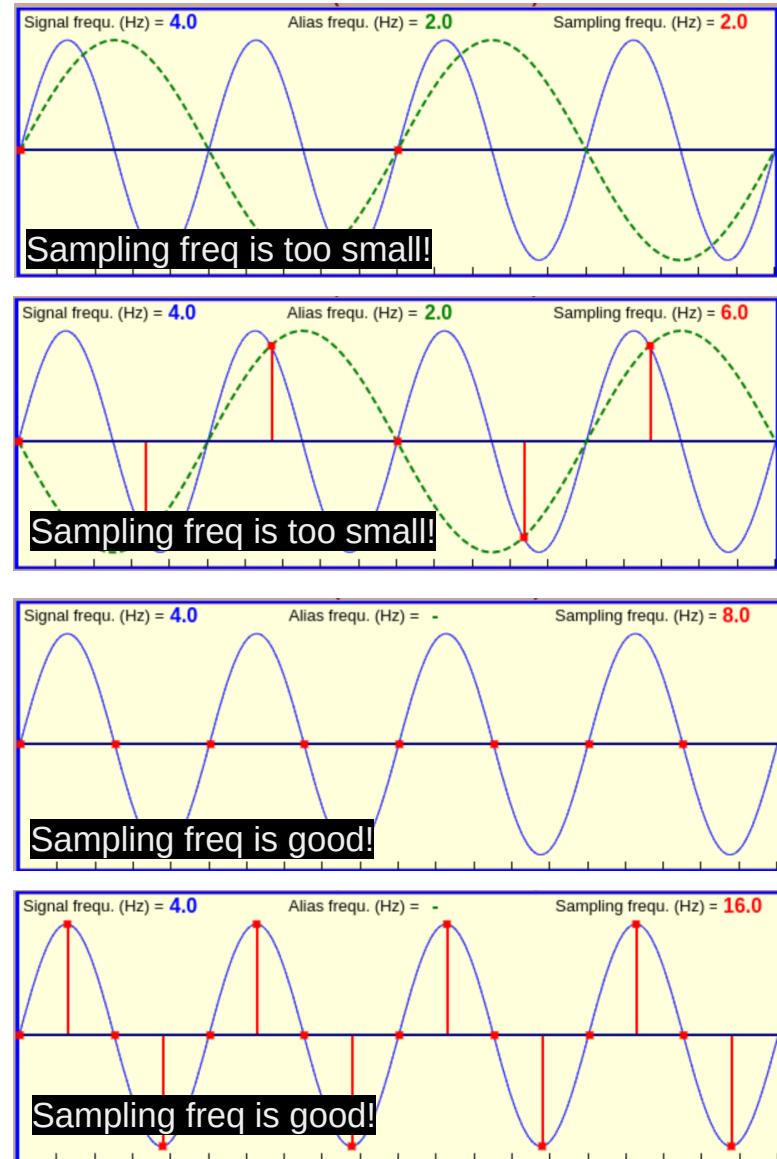


Changes in geometry, textures, highlights → sources of aliasing.

Aliasing

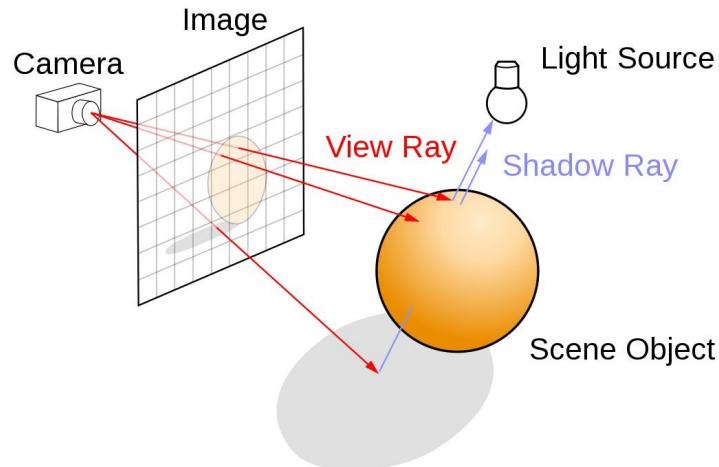
Blue: original signal
Red: samples
Green: reconstructed signal

- **Aliasing** – Occurs if signal is sampled at too low frequency
 - Reconstruction of signal from samples can not be performed correctly, e.g., reconstructing final image from pixel samples (**filtering**)
- **Sampling theorem: Sampling frequency must be twice the maximum frequency of sampled signal; Nyquist rate/limit**
 - As maximum frequency must be known for determining the sampling rate, **signal must be band-limited**
- **3D scenes are normally never band-limited** when rendering with point samples
 - **Sharp triangle or shadow edges** → discontinuous changes → infinite frequencies
 - **Rapid change of color**, e.g., specular highlights
- Point sampling is almost always used and entirely avoiding aliasing is not possible
 - Sampling pixel footprint which may contain lots of information requires a lot of samples



Anti-aliasing

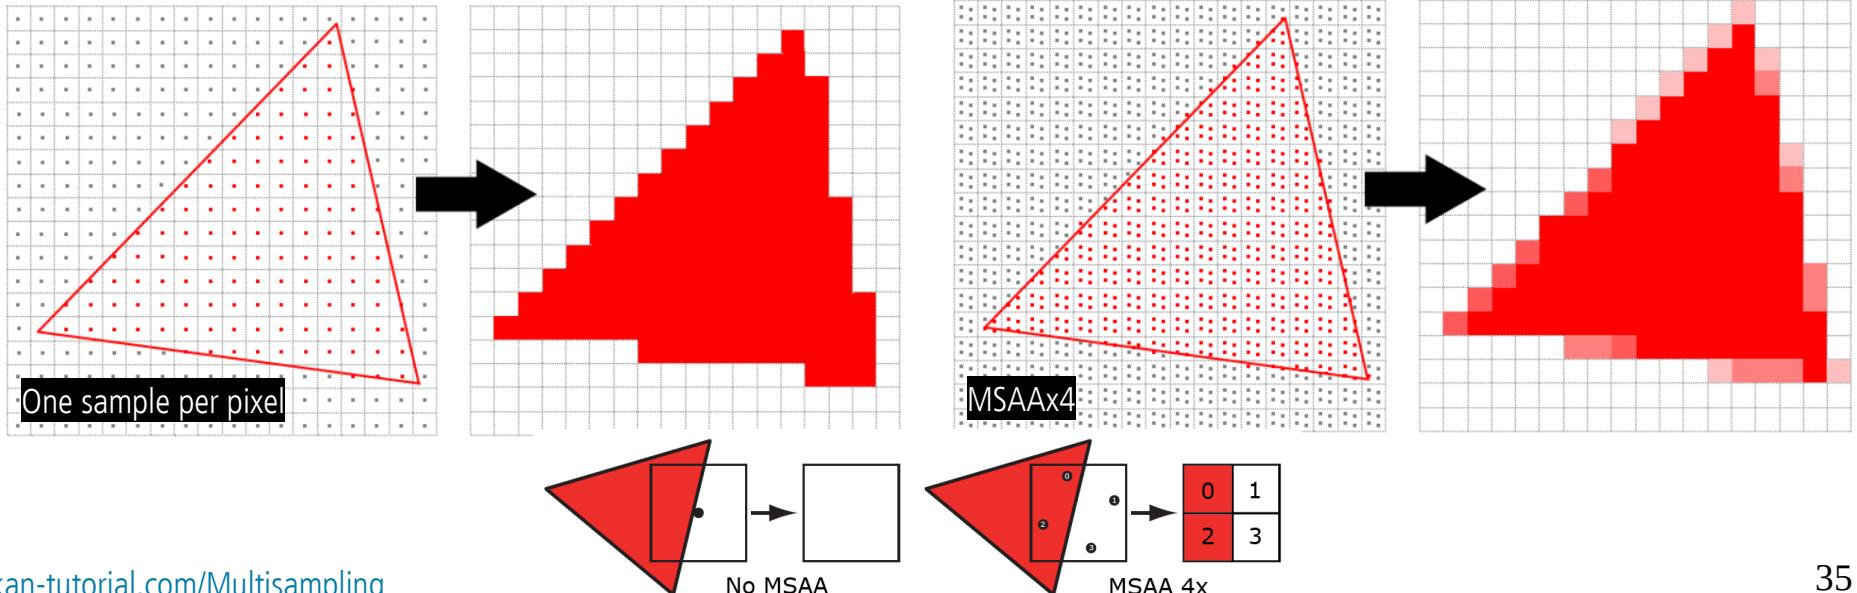
- Most common anti-aliasing method: **screen-based anti-aliasing**
 - Using multiple samples per pixel during rendering
- Anti-aliasing can be applied on: geometry, textures, shadows, highlights, etc.
 - At times, when it is possible to know if the signal is band limited and then anti-aliasing is performed
 - Trade-off between: quality, ability to capture sharp details or other phenomena, appearance, memory and speed



Instead of one ray per pixel, generate multiple rays - samples

Multisample anti-aliasing

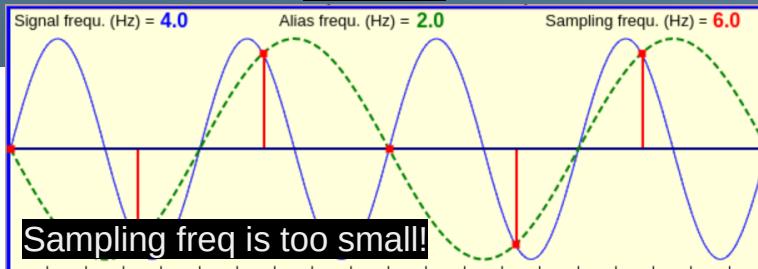
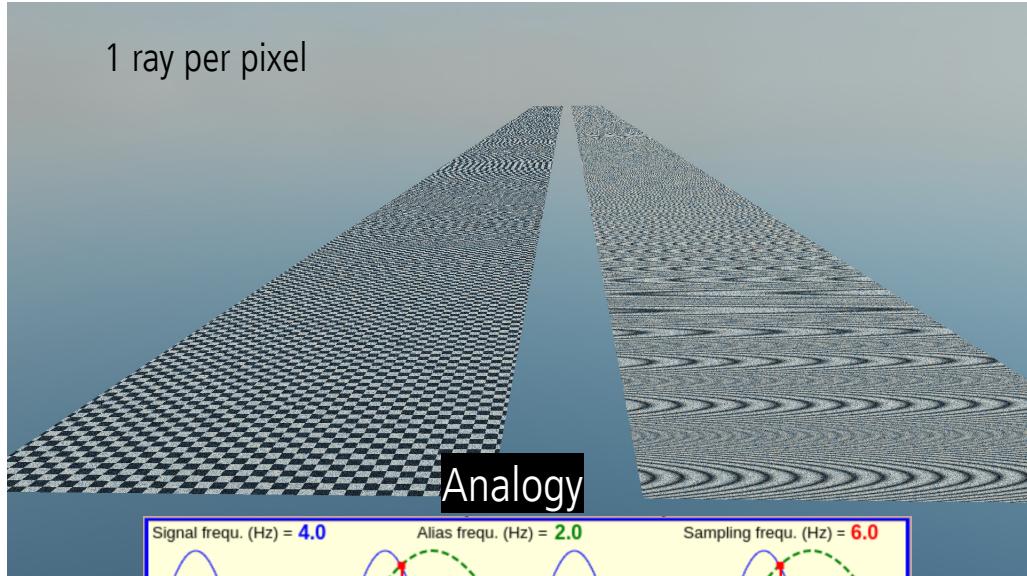
- Aliasing appears always when there are a lot of details under a pixel footprint in the scene
- To solve this, multiple sample per pixels can be used to “catch” high frequency details.
 - This method is called Multisample anti-aliasing (MSAA)



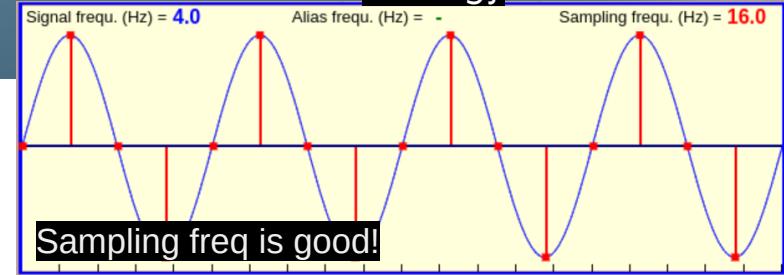
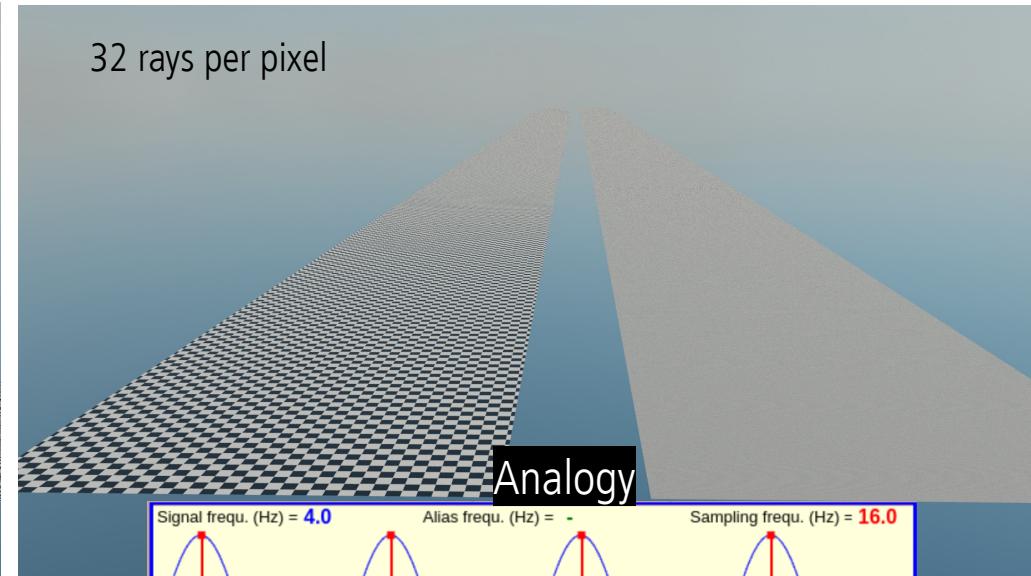
Multisample anti-aliasing

- Using more samples per pixel reduces aliasing.
 - One sample per pixel is often not enough to reconstruct the image

1 ray per pixel

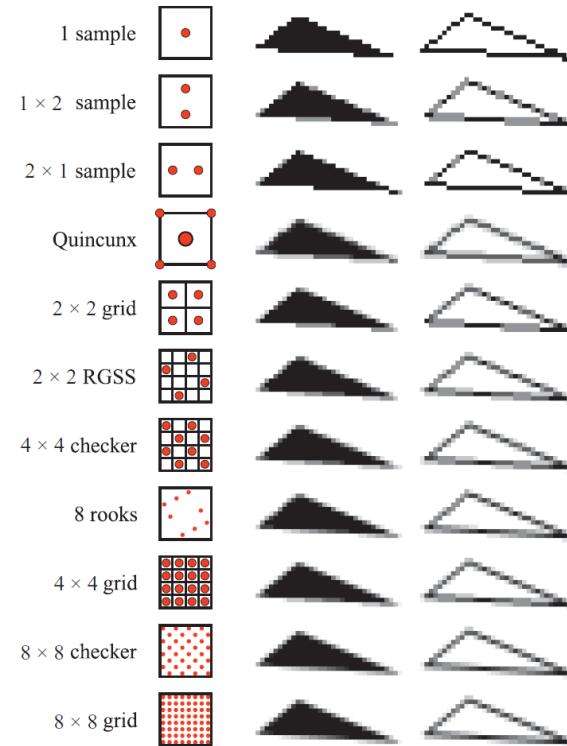


32 rays per pixel



Antialiasing: sampling schemes

- Using more samples per pixel and combining those results in more representative pixel color
- Different pixel sampling schemes exists: trade-off between quality and speed



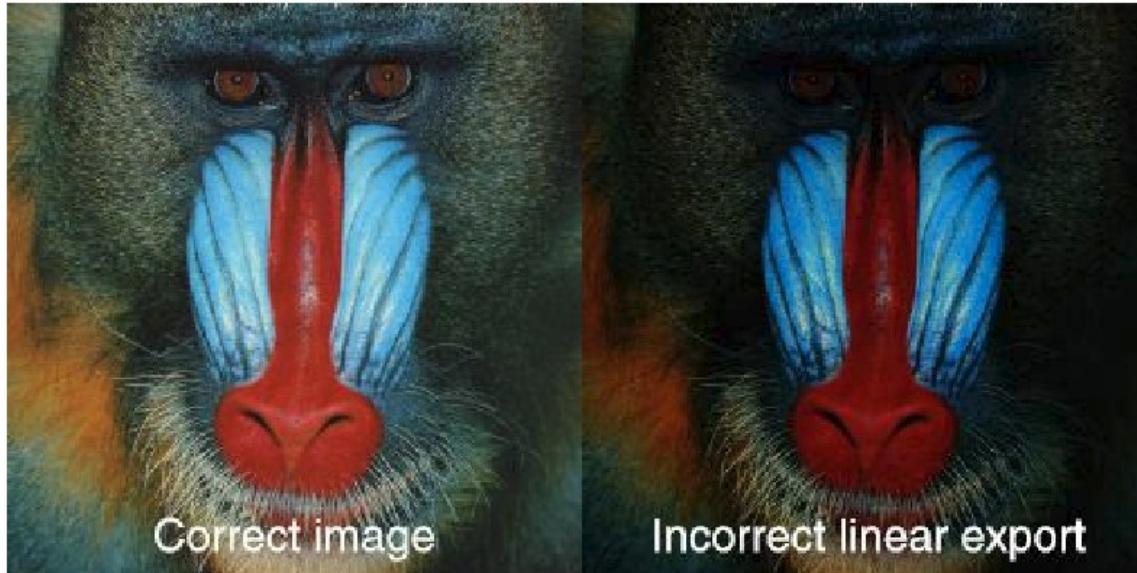
Scene to screen

Limitation of display device

- **3D scene contains:**
 - Continuous description of objects (e.g., triangles, parametric surfaces, etc.)
 - Strong light sources and reflections (e.g., point lights, specular reflections)
- **Display devices are limited (discrete):**
 - Resolution (number of pixels)
 - Brightness (intensity)
 - Contrast
 - Color (gamut)
- **Rendered images** may contain values which can not be directly shown on display devices
 - Idea: manipulate image values so that image looks as intended when displayed on particular display device

Problem: displaying rendered images

- All input values (e.g., texture image) for rendering and all values during rendering computation must be in **linear** colorspace
 - Linear values are needed for correct addition and multiplication operations
- Rendered image containing **linear color space** will not display correctly if shown directly on a display device
 - Linear values will appear too dim on the screen

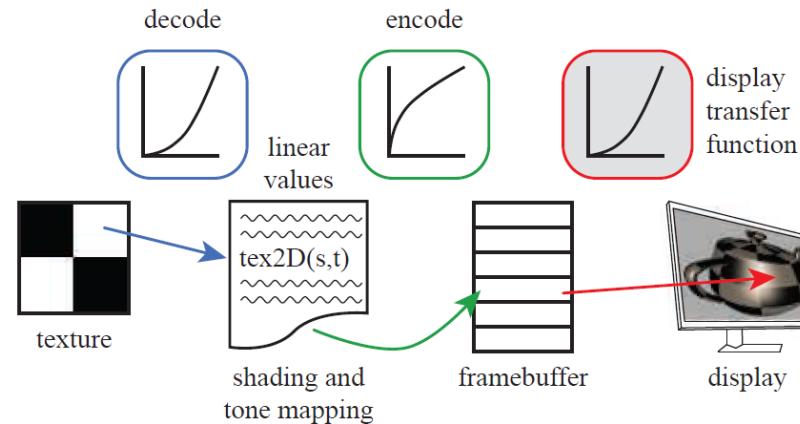


Display limitations

- Relationship between digital color values in image buffer and amount of light emitted from the display is described with electrical-optical transfer function (EOTF)
 - Standard dynamic range (SDR) displays (e.g., personal computer display) use sRGB display standard
 - High dynamic range (HDR) displays uses standard Rec. 2020 and Rec. 2100
- Display devices have nonlinear relationship between input voltage and display amount of light
 - As energy level applied to pixel is increased, the amount of light emitted doesn't grow linearly but according to power law → nonlinear relationship
 - e.g., pixel set to 50% will emit 0.5^2 amount of light

Solution: display encoding

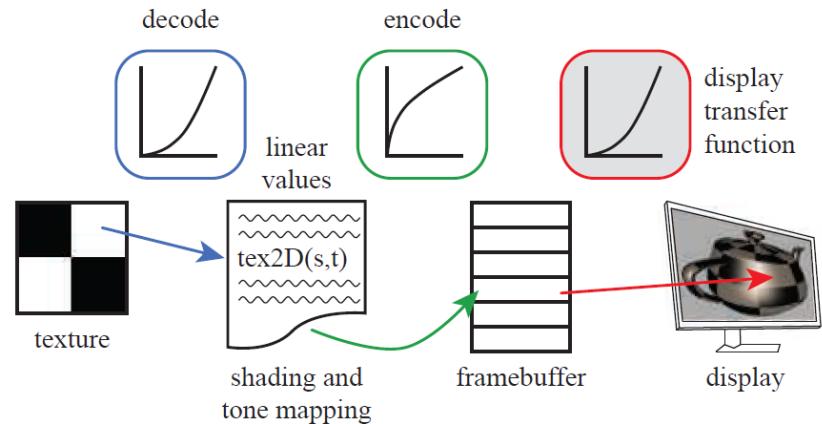
- The goal is to **cancel out the effect of display transfer function (EOTF)**
- Linear color values are **encoded** for display
 - Inverse of display transfer function (EOTF) is applied to color values in image buffer → gamma correction



Gamma correction

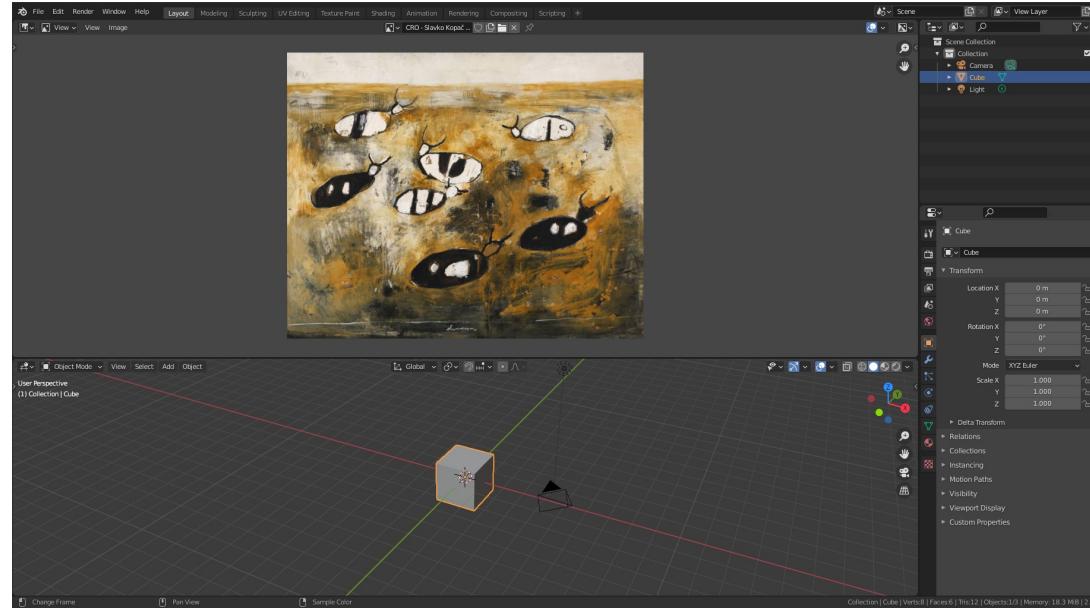
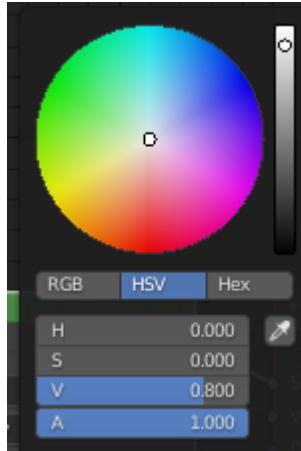
- Gamma correction is **last step in rendering and post-processing** – when everything is computed and image is ready for display.
 - To encode rendered image which is in linear colorspace (x) for sRGB display standard (y) the following formula is applied to all color values:

$$y = x^{1/\text{gamma}}, \text{ gamma} = 2.2$$



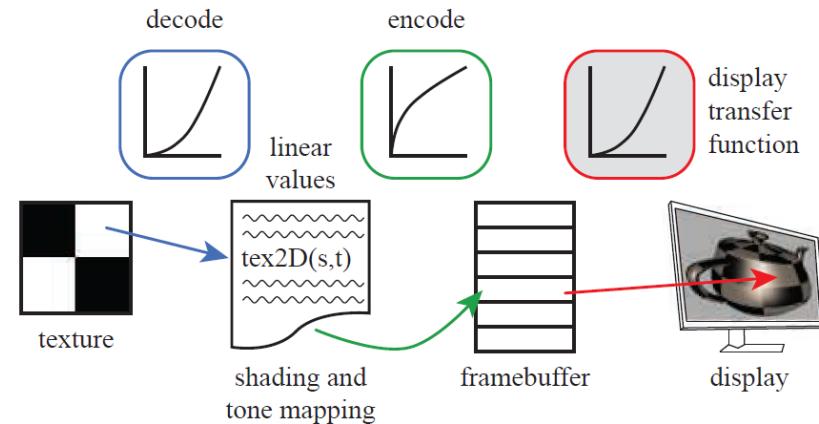
Problem: preparing data for rendering

- When working in modeling tools users pick texture images or colors on screen - those colors are **encoded for display device** so we can see them properly
 - Those colors **can not be used in rendering computation directly** because they are in non-linear space



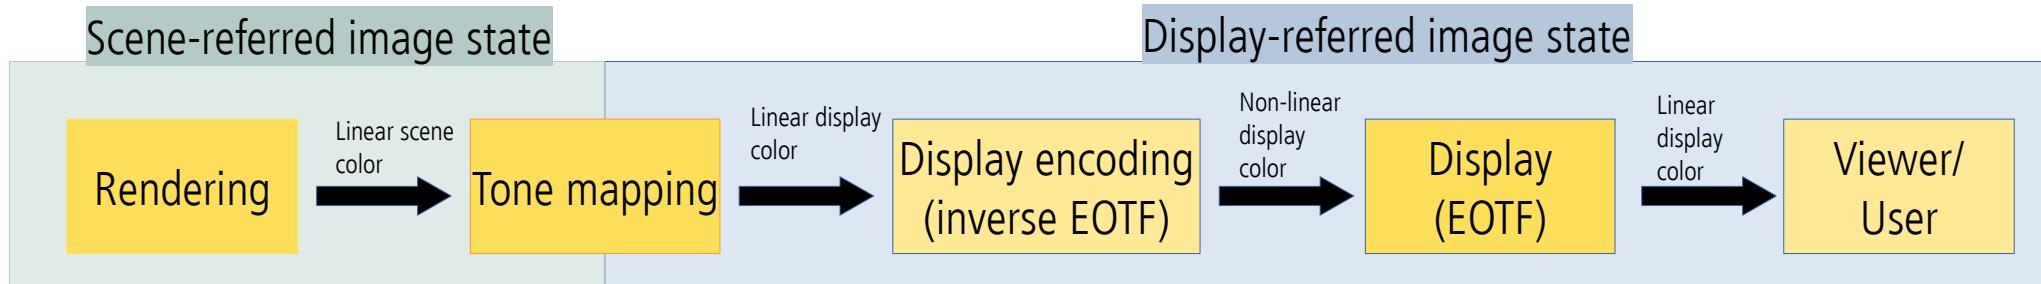
Solution: inverse gamma correction

- Everything we see on the screen (e.g., image textures or color-pickers in modeling tools) is display-encoded data and those values **must be decoded to linear values for rendering**
 - To decode sRGB display encoded values (y) to linear colorspace (x) the following formula is applied to all color values:
 - $x = y^{\text{gamma}}$, where $\text{gamma} = 2.2$



Imaging pipeline

- **Imaging pipeline:** describes color handling from initial rendering to final display.
- **Rendering:** creating 2D image based on 3D scene description.
 - The results are pixel values in **color buffer** – display results still need to be determined.
- **Display encoding:** converting linear color (radiance) values to nonlinear values for display
- **Tone mapping:** converting linear scene-referred values to linear display-referred values



- Defined in reference to scene color (radiance) values
 - Physically-based rendering computations
- Defined in reference to display color (radiance) values.
 - Appropriate for viewing

Tone mapping

- Tone mapping (reproduction) is process of converting scene color (radiance) values to display color (radiance) values
 - Transform applied on this step is called end-to-end or scene-to-screen transform
- Two main types:
 - Image reproduction - create display-referred image that **reproduces**, as closely as possible, given display and viewing properties, perceptual impression that viewer would have if they were observing original scene
 - Preferred image reproduction – create display-referred image that **looks better** (to some criteria) than original scene.

Tone mapping: (1) image reproduction

- During image reproduction it is important to keep in mind that scene luminance and saturation exceeds display capabilities, that is, **dynamic range of scene is much larger than display device dynamic range**
 - Therefore, (high) dynamic range of scene must be mapped to (low) dynamic range of the display
 - This mapping is done using **sigmoid (s-shaped) tone-reproduction curve**, researched for photochemical film, thus name “filmic”
- Further, **exposure** is critical to image reproduction
 - Exposure in **photography** refers to controlling the **amount of light falling on film/sensor**.
 - In **rendering**, exposure is a **linear scaling operation** performed on **scene-referred image** before tone reproduction transform is applied.
- **Tone reproduction transform** and **exposure** are closely tied together
 - Tone transforms are designed to be applied to **scene-referred values** which have been scaled by certain exposure

Tone-reproduction
transform

Scaling by exposure

Image reproduction

- Global tone mapping consists of two steps:
 1. Scaling scene-referred rendered image by exposure
 2. Applying tone reproduction transform applied on image
- In contrast, local tone mapping:
 - Uses different mapping pixel-to-pixel based on surrounding pixels and other factors

Global tone mapping: scaling by exposure

- Exposure scaling factor is determined by analyzing scene-referred luminance values
- Analyzing pixel luminance of the **scene-referred pixel values** in the rendered image
 - Metric 1: log-average of scene luminance
 - Problem: sensitive to outliers (e.g., small number of bright pixels that affect the exposure for the whole frame)
 - Metric 2: median of scene luminance by using histogram of luminance values
 - https://cdn.cloudflare.steamstatic.com/apps/valve/2008/GDC2008_PostProcessingInTheOrangeBox.pdf
- Analyzing light luminance alone, that is, exposure scaling based on light intensity
 - <https://www.gdcvault.com/play/1018086/Photorealism-Through-the-Eyes-of>



Different exposures of the same render:
https://docs.blender.org/manual/en/latest/render/color_management.html

Global tone mapping: tone reproduction transform

- **Tone reproduction transform:** 1D curve mapping scene-referred input values to display-referred output values

- Transform can be directly applied on R, G, B values: the result will be in display gamut, but hue and saturation are shifted
- Transform can be applied to luminance: hue and saturation will not be shifted, but resulting color might be outside of display gamut

- **Reinhard tone reproduction transform**

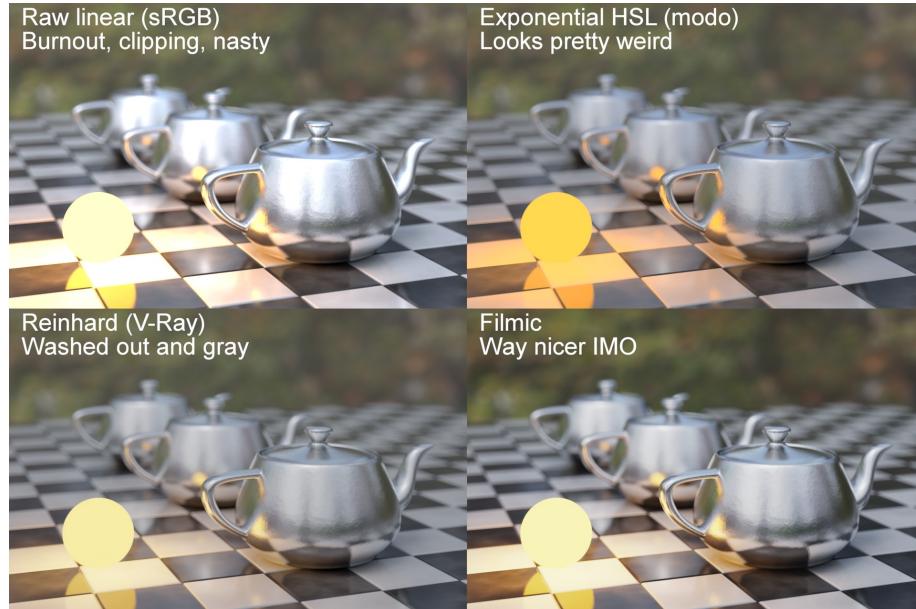
- Leaves dark values unchanged while bright values asymptotically go to white

- **ACES tone mapping**

- Standard for tone managing for motion picture and television
- Emerging in real-time applications: Unity and Unreal

- Many other tone mapping transform exist, examples:

- Adaptive Logarithmic Mapping: similar to Reinhard with adjustable output display luminance: <https://resources.mpi-inf.mpg.de/tmo/logmap/logmap.pdf>
- Filmic tone mapping:
<http://duikerresearch.com/2015/09/filmic-tonemapping-ea-2006/>



<https://blenderartists.org/t/cycles-tonemapping/566158/89?page=5>

Tone mapping: (2) preferred image reproduction

- **Color grading:** creative manipulation of image colors to obtain image desired artistic look
 - <https://www.gdcvault.com/play/1012248/Perfecting-the-Pixel-Refining-the>
- **Color grading is performed using look-up tables (LUTs)** contain desired tabular color transformations which are applied on image
 - https://renderwonderland.com/publications/s2010-color-course/hoffman/s2010_color_enhancement_and_rendering_hoffman_b.pdf
- Color grading can be performed on:
 - Scene-referred image data: produce high fidelity results
 - Display-referred image data: easier to set-up, often was used in real-time applications



More into topic

- Image space effects:
 - Motion blur: https://www.nvidia.com/docs/io/8230/gdc2003_openglshaderticks.pdf
 - Motion blur: <https://developer.nvidia.com/gpugems/gpugems3/part-iv-image-effects/chapter-27-motion-blur-post-processing-effect>
 - Motion blur: <https://www.bryancphail.com/wp/?p=600>
 - Motion blur: https://docs.blender.org/manual/en/latest/render/cycles/render_settings/motion_blur.html
 - NPR: https://en.wikipedia.org/wiki/Non-photorealistic_rendering
 - General post-processing: <https://docs.unity3d.com/Packages/com.unity.render-pipelines.universal@7.1/manual/integration-with-post-processing.html>
- Color:
 - <https://developer.nvidia.com/sites/default/files/akamai/gameworks/hdr/UHDColorForGames.pdf>
 - <https://cinematiccolor.org/>
- Color grading:
 - <http://filmicworlds.com/blog/minimal-color-grading-tools/>
 - <https://docs.unrealengine.com/4.27/en-US/RenderingAndGraphics/PostProcessEffects/ColorGrading/>

Repository

- <https://github.com/lorentzo/IntroductionToComputerGraphics>